

# Improving native memory management & diagnostics in Node.js

Joyee Cheung

Node.js collaboration summit, April 4, London

# Agenda

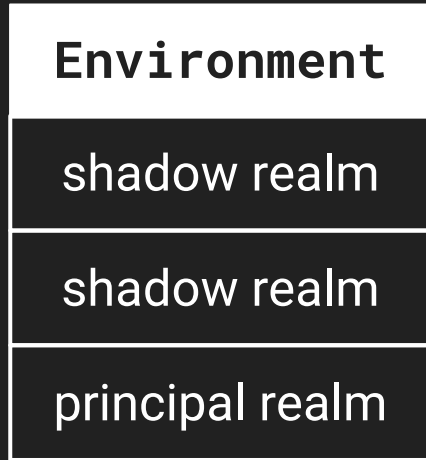
- Current infrastructure for native memory management in Node.js
- Case study of dynamic import() memory issues
  - Motivation: importModuleDynamically() broken for 2 years keeping many stuck with v16.x. Many users complained about not getting a fix prioritized. Happened to find a fix/workaround. Trying to spread the knowledge out.
- A path to Oilpan (cppgc)
- Memory debugging

# Current infrastructure for native memory management in Node.js

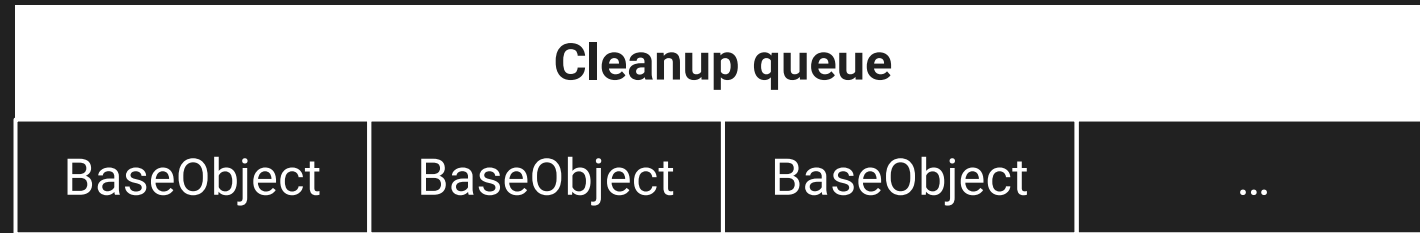
- Many internal JS land objects are created as wrappers over internal C++ objects (if you see \*Wrap in JS or C++, it's most likely to be one)
- **BaseObject**: Abstraction of (almost?) all C++ object wrappers

```
> ls src | grep wrap
async_wrap-inl.h
async_wrap.cc
async_wrap.h
cares_wrap.cc
cares_wrap.h
connect_wrap.cc
connect_wrap.h
connection_wrap.cc
connection_wrap.h
fs_event_wrap.cc
handle_wrap.cc
handle_wrap.h
js_udp_wrap.cc
module_wrap.cc
module_wrap.h
node_object_wrap.h
pipe_wrap.cc
pipe_wrap.h
process_wrap.cc
req_wrap-inl.h
req_wrap.h
```

# BaseObject tracking

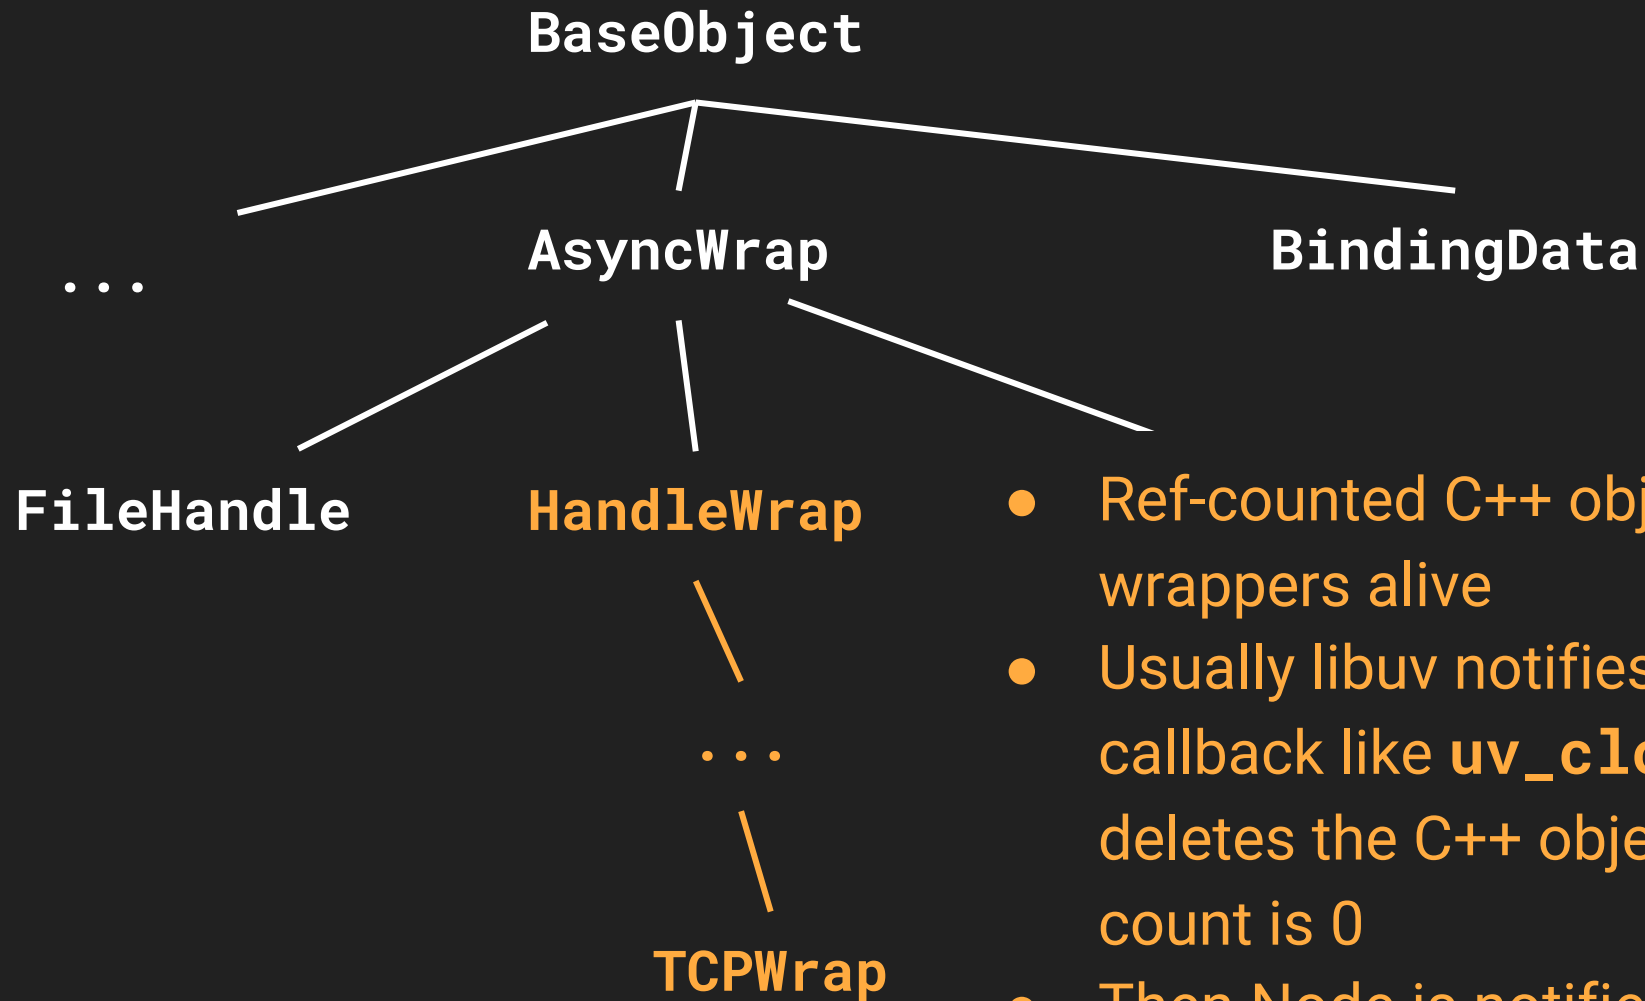


- **BaseObjects** are tracked in a per-realm cleanup queue
- At realm shutdown, **BaseObjects** that are not yet cleared would all be released through the cleanup queue.



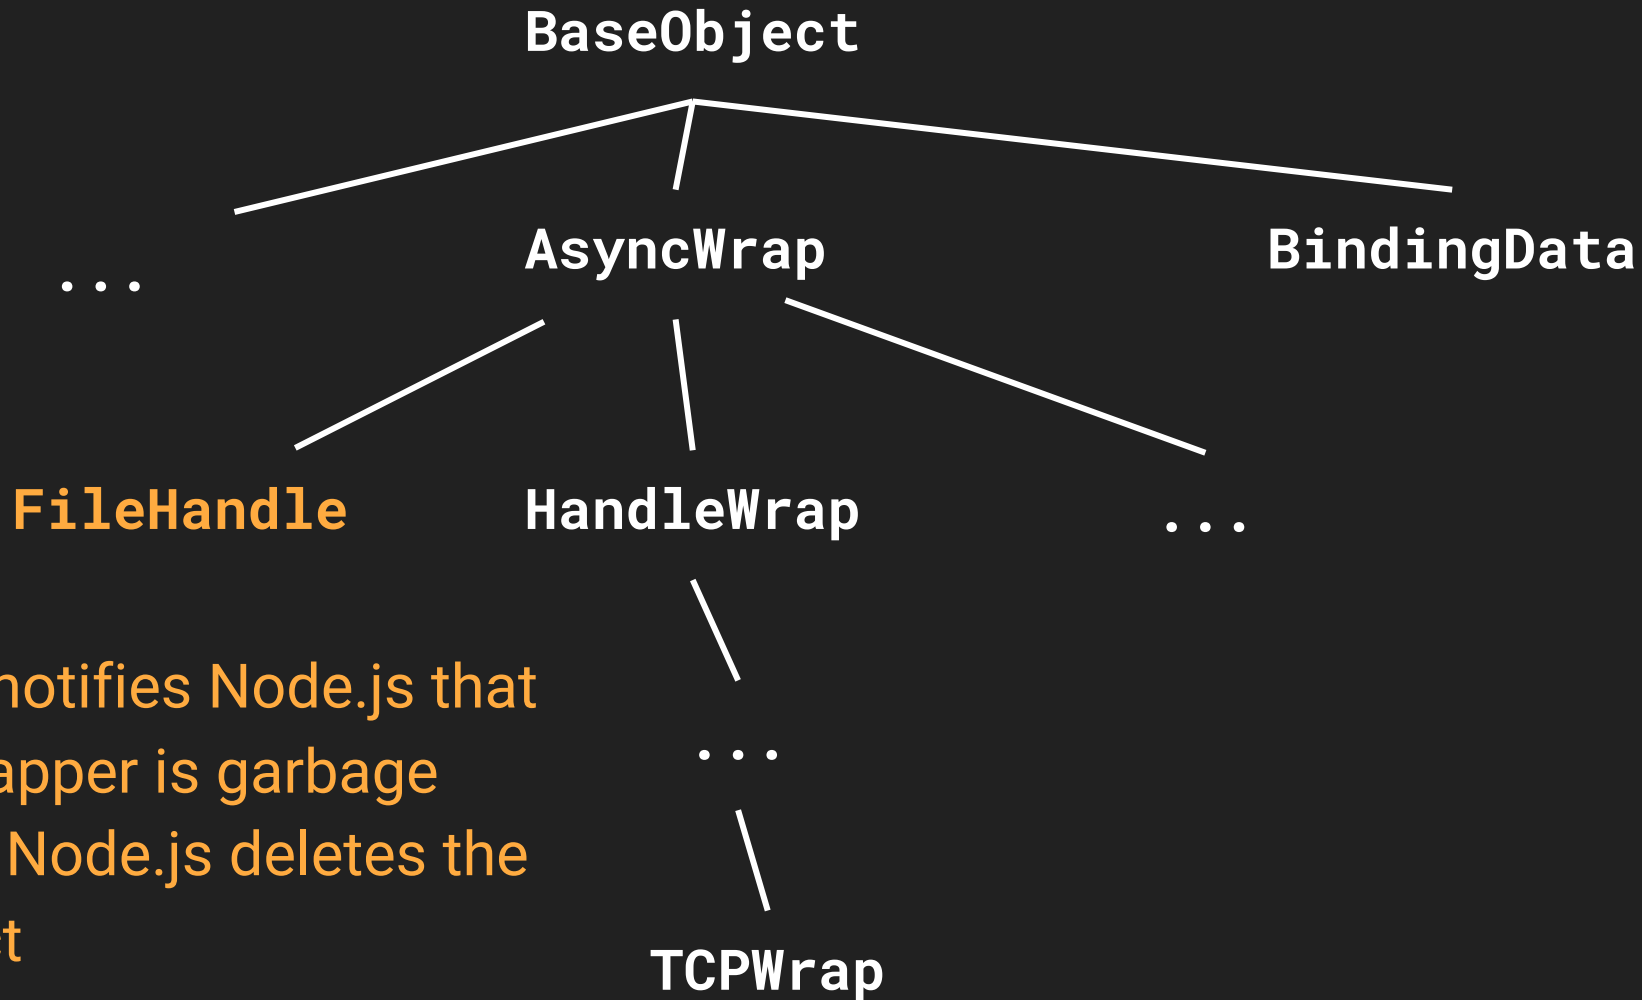
```
BaseObject::BaseObject(Realm* realm, Local<Object> object)
    : persistent_handle_(realm->isolate(), object), realm_(realm) {
    SetInternalFields(realm->isolate_data(), object, static_cast<void*>(this));
    realm->AddCleanupHook(DeleteMe, static_cast<void*>(this));
    realm->modify_base_object_count(1);
}
```

# BaseObject: the base class of (almost) all wrappers



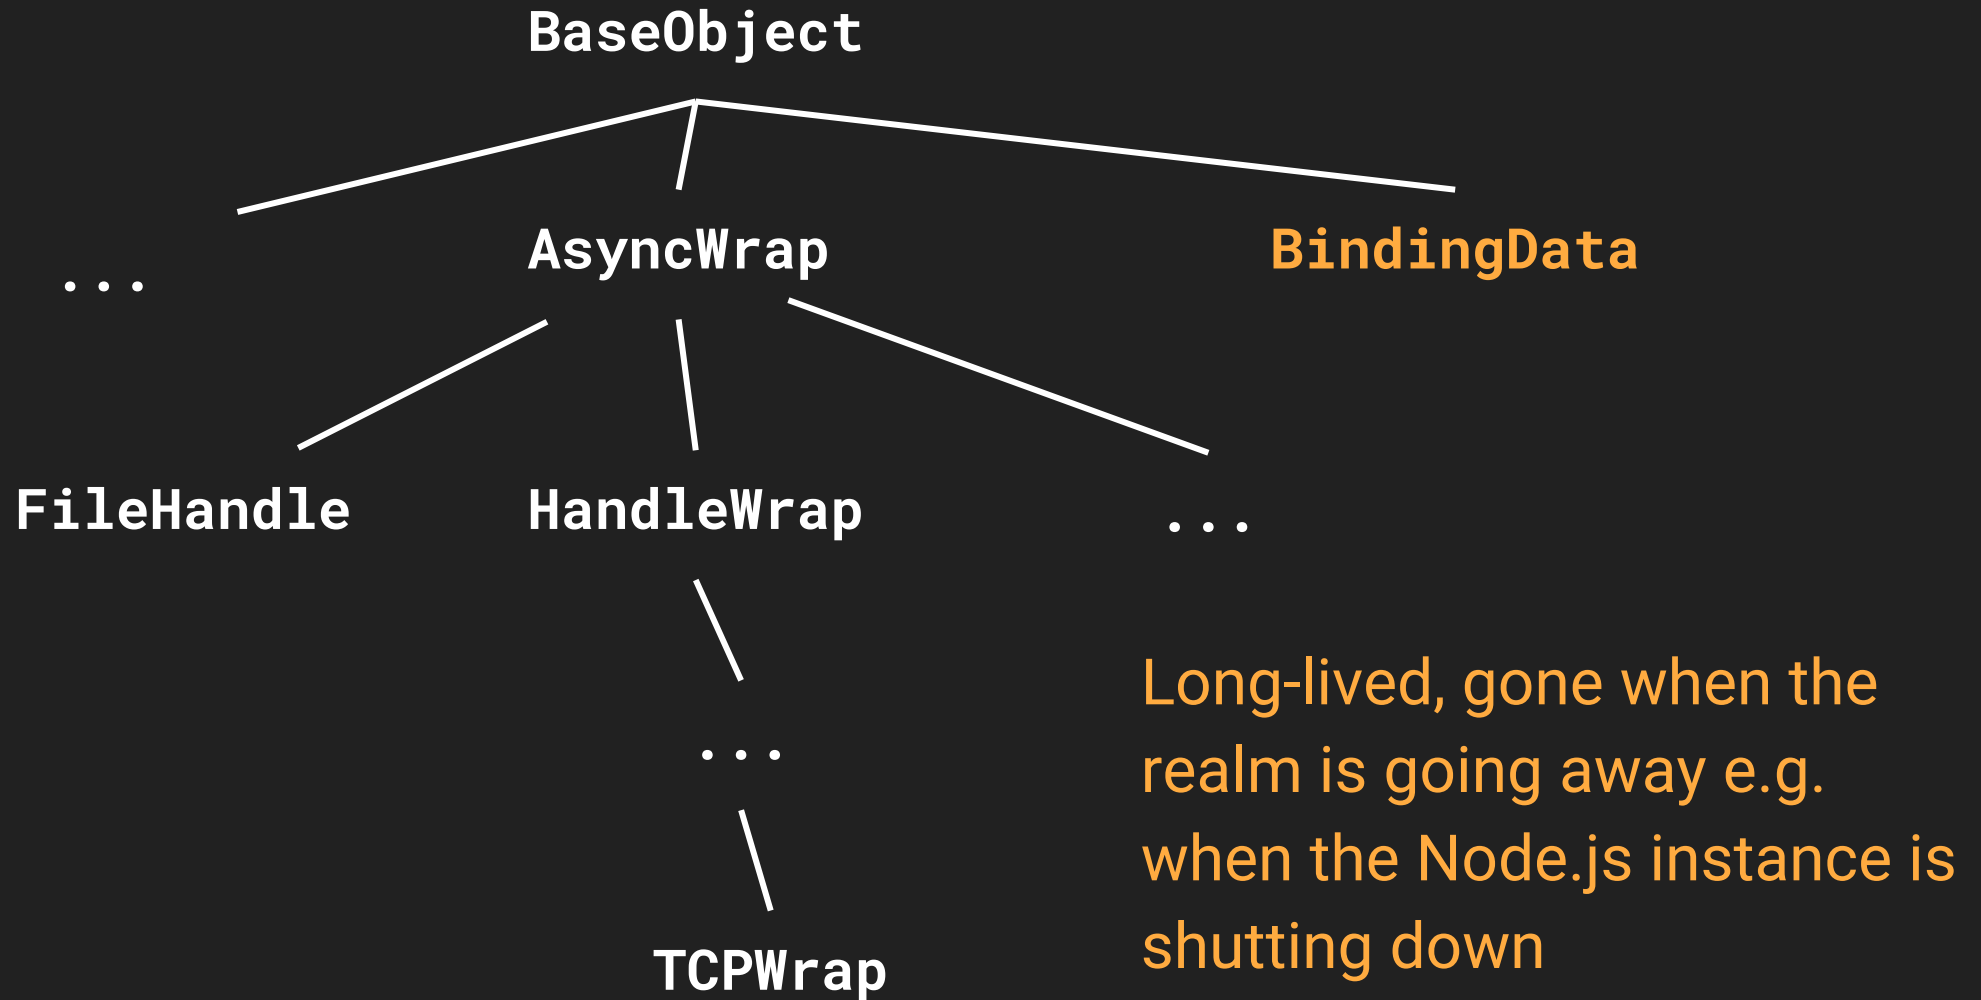
- Ref-counted C++ objects hold JS wrappers alive
- Usually libuv notifies Node.js in a callback like **uv\_close()**, Node.js deletes the C++ object when ref count is 0
- Then Node.js notifies V8 to release the JS wrapper

# BaseObject: the base class of (almost) all wrappers

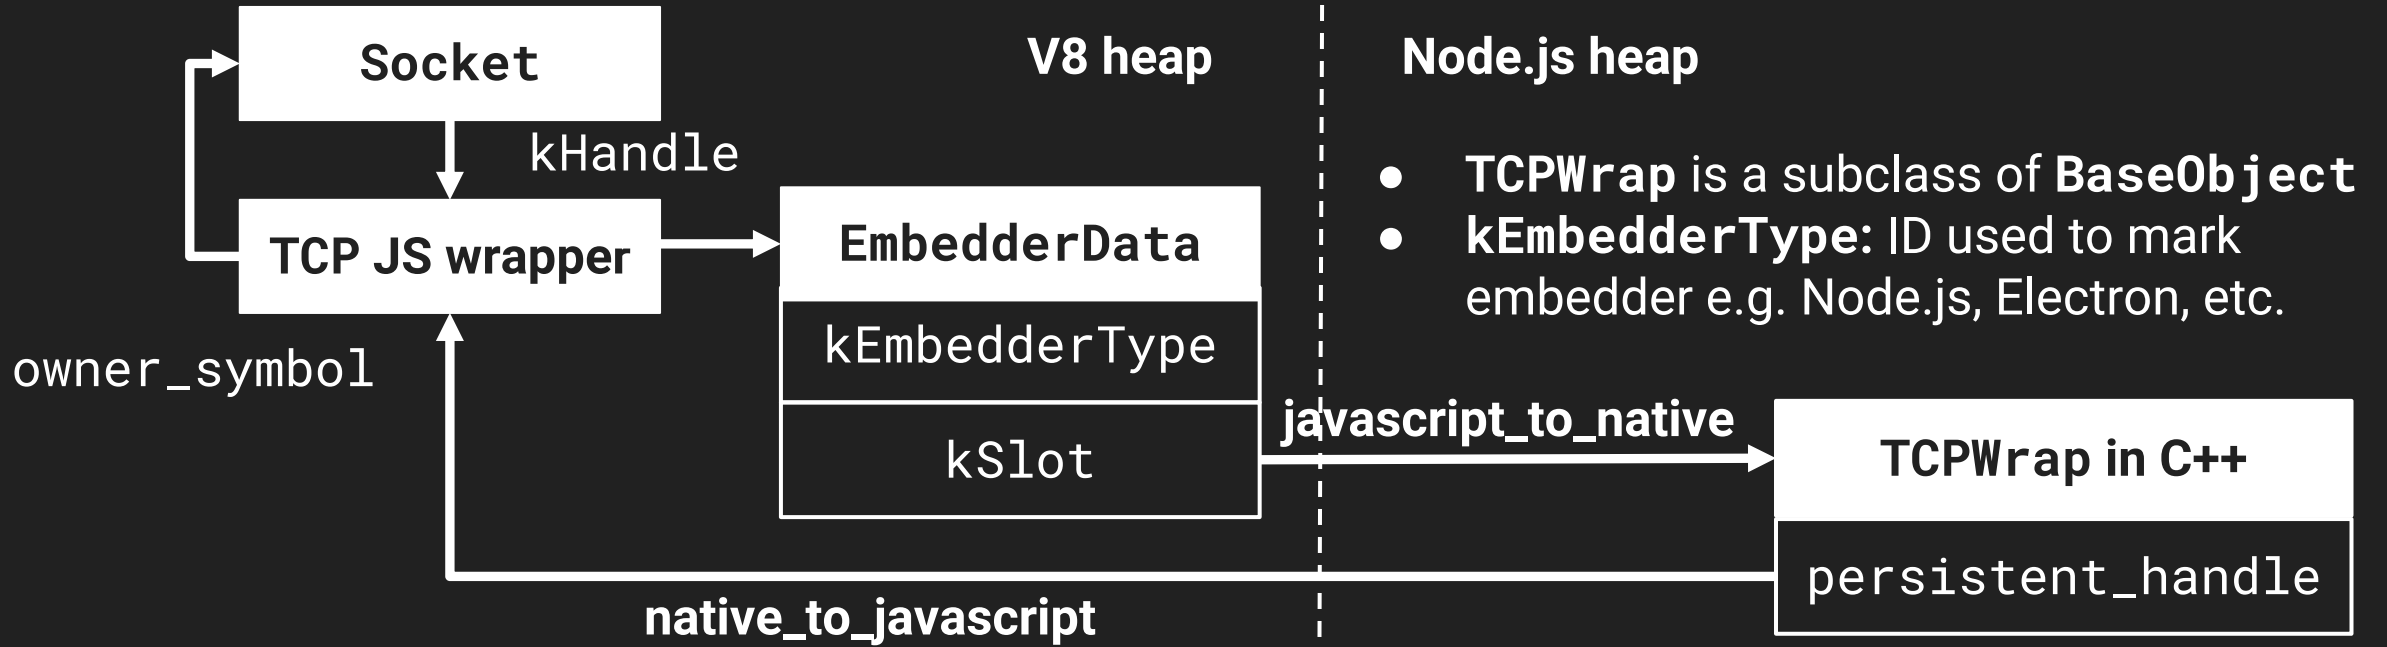


When V8 notifies Node.js that the JS wrapper is garbage collected, Node.js deletes the C++ object

# BaseObject: the base class of (almost) all wrappers



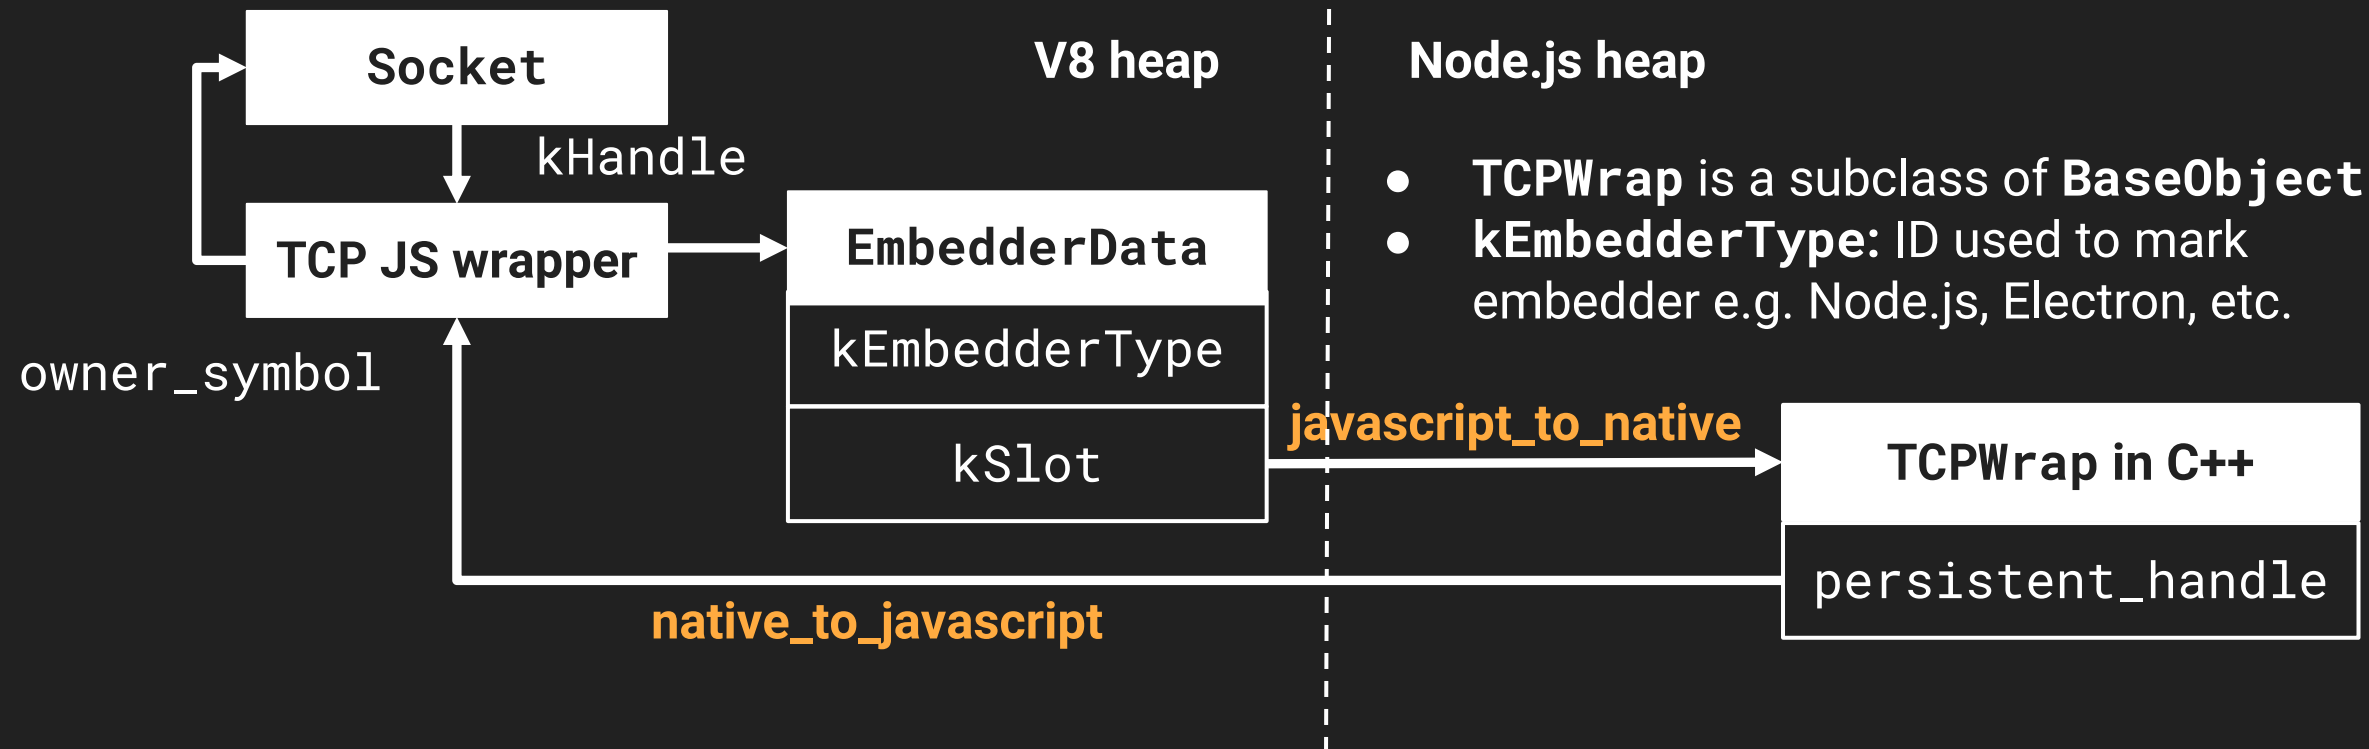
# Kept alive by C++: TCPWrap & net.Socket



```
Socket.prototype.connect = function(...args) {  
  this[kHandle] = new TCP(SOCKET);  
  this[kHandle][owner_symbol] = this;  
}
```



# Kept alive by C++: TCPWrap & net.Socket



▼ <symbol kHandle> :: TCP @6285	10	56	0 %
▶ 3 :: onStreamRead() @75257	8	64	0 %
▶ <symbol owner_symbol> :: Socket @77145	9	488	0 %
▼ javascript_to_native :: Node / TCPWrap<Socket> @50	11	424	0 %
native_to_javascript :: TCP @6285	10	56	0 %

# Kept alive by JS

```
class BaseObject : public MemoryRetainer {  
public:  
    void MakeWeak() {  
        persistent_handle_.SetWeak(this, [](const WeakCallbackInfo<BaseObject>& data) {  
            BaseObject* obj = data.GetParameter();  
            obj->persistent_handle_.Reset();  
            delete obj;  
        }, WeakCallbackType::kParameter);  
    }  
private:  
    v8::Global<v8::Object> persistent_handle_;  
};
```

When V8 notifies Node.js that the JS wrapper is garbage collected, Node.js delete the C++ object

If MakeWeak() isn't called, the BaseObject is usually kept alive by the per-realm CleanupQueue

# C++ -> JavaScript memory reference

- Many BaseObjects have additional references to JS values
- Callback-based memory management is also a hack (though it's been working)

```
/**
 * Install a finalization callback on this object.
 * NOTE: There is no guarantee as to *when* or even *if* the callback is
 * invoked. The invocation is performed solely on a best effort basis.
 * As always, GC-based finalization should *not* be relied upon for any
 * critical form of resource management!
 *
 * The callback is supposed to reset the handle. No further V8 API may be
 * called in this callback. In case additional work involving V8 needs to be
 * done, a second callback can be scheduled using
 * WeakCallbackInfo<void>::SetSecondPassCallback.
 */
template <typename P>
V8_INLINE void SetWeak(P* parameter,
                       typename WeakCallbackInfo<P>::Callback callback,
                       WeakCallbackType type);
```

# Case study of dynamic import() memory issues

```
class ContextifyScript : public BaseObject {  
  public:  
    ContextifyScript(Environment* env, Local<Object> object)  
      : BaseObject(env, object) {  
      script_.Reset(isolate, v8_script);  
      MakeWeak();  
    }  
  private:  
    v8::Global<v8::UnboundScript> script_;  
}
```

REPL input compiled with a ContextifyScript

```
> node --expose-gc  
> const fn = () => import('crypto')  
> gc()  
> fn()  
[1]      85524 segmentation fault  node --expose-gc
```

**Issue #35889**

# Case study of dynamic import() memory issues

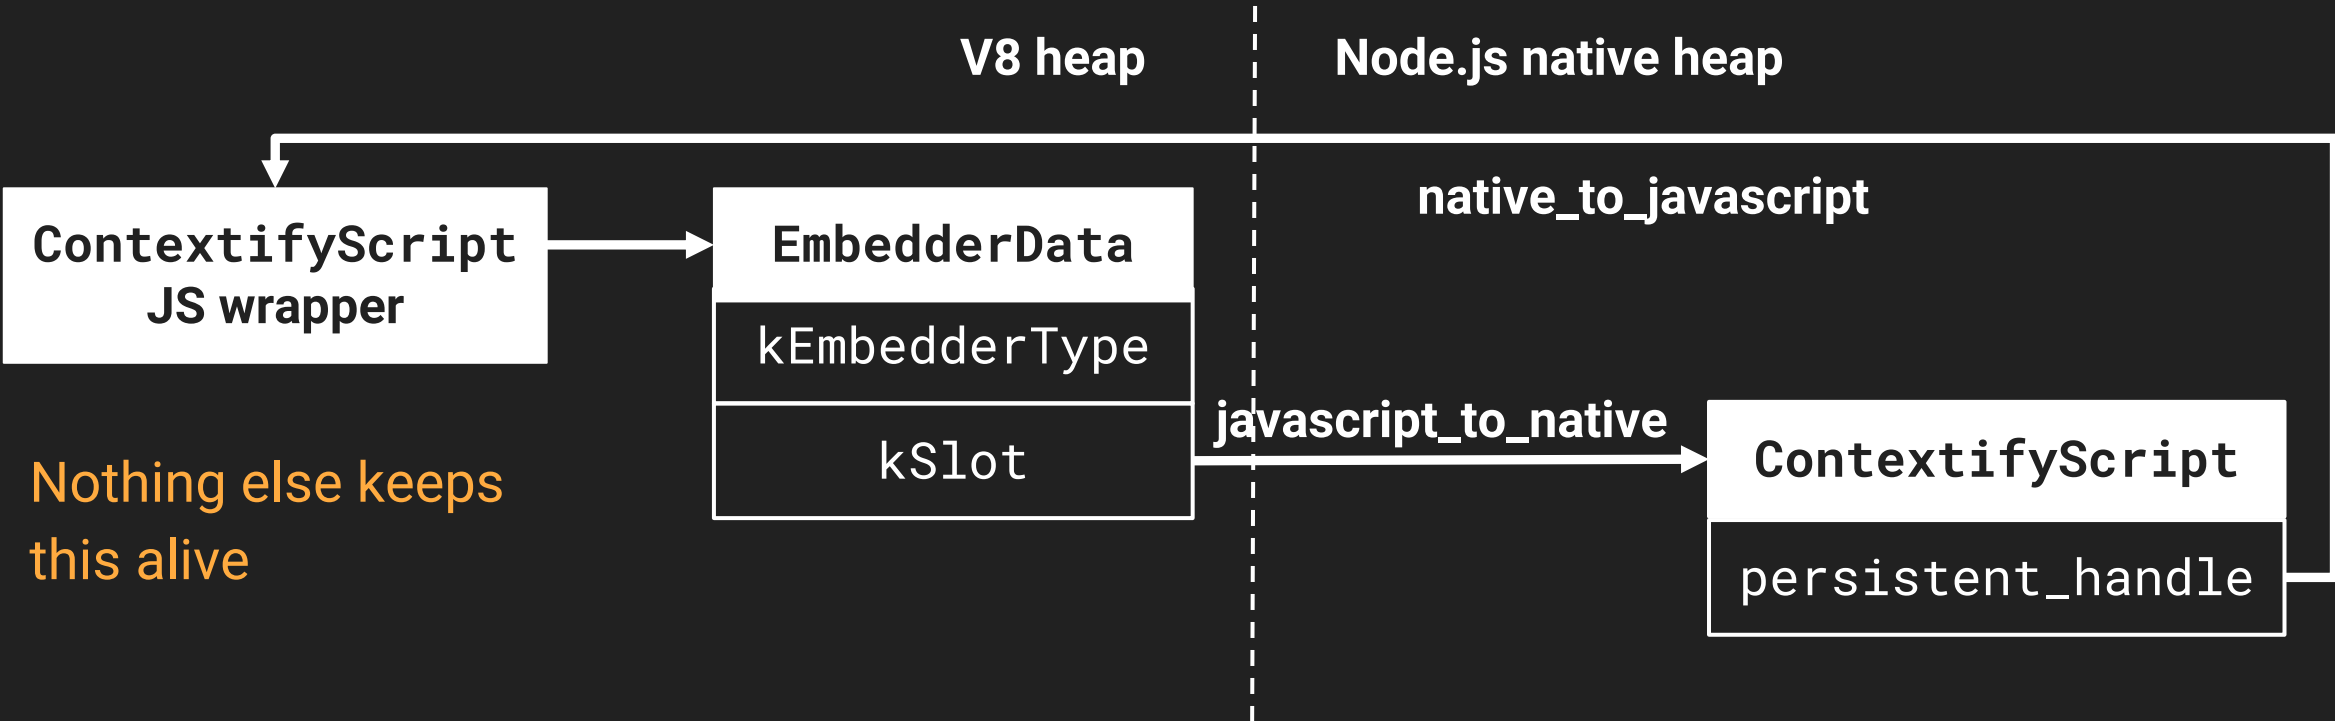
```
class ContextifyScript : public BaseObject {  
  public:  
    ContextifyScript(Environment* env, Local<Object> object)  
      : BaseObject(env, object) {  
      script_.Reset(isolate, v8_script);  
      MakeWeak();  
    }  
  private:  
    v8::Global<v8::UnboundScript> script_;  
}
```

REPL input compiled with a ContextifyScript

```
> node --expose-gc  
> const fn = () => import('crypto')  
> gc()  
> fn()  
[1]      85524 segmentation fault  node --expose-gc
```

JS land can lose reference to the JS object, and it goes away too soon and crashes due to use-after-free

# Case study of dynamic import() memory issues



▼ Detached Node / ContextifyScript @16619648 ✂	–	48	0 %	48	0 %
▼ native_to_javascript :: Script @6187	–	48	0 %	48	0 %
▶ __proto__ :: ContextifyScript @44437	12	24	0 %	1 096	0 %
▶ map :: system / Map @44435	–	72	0 %	120	0 %
▶ sourceMapURL :: system / Oddball @69 ☐	2	48	0 %	120	0 %
▶ javascript_to_native :: Detached Node / ContextifyScript @16619648 ✂	–	48	0 %	48	0 %
▶ 1 / part of key (Script @6187) -> value (Object @68131) pair in WeakMap (table	–	32	0 %	32	0 %

# Case study of dynamic import() memory issues

```
class ContextifyScript : public BaseObject {
```

```
public:
```

```
ContextifyScript(Environment* env, Local<Object> object)
```

```
    : BaseObject(env, object) {
```

```
    script_.Reset(isolate, v8_script);
```

```
    MakeWeak();
```

```
}
```

```
private:
```

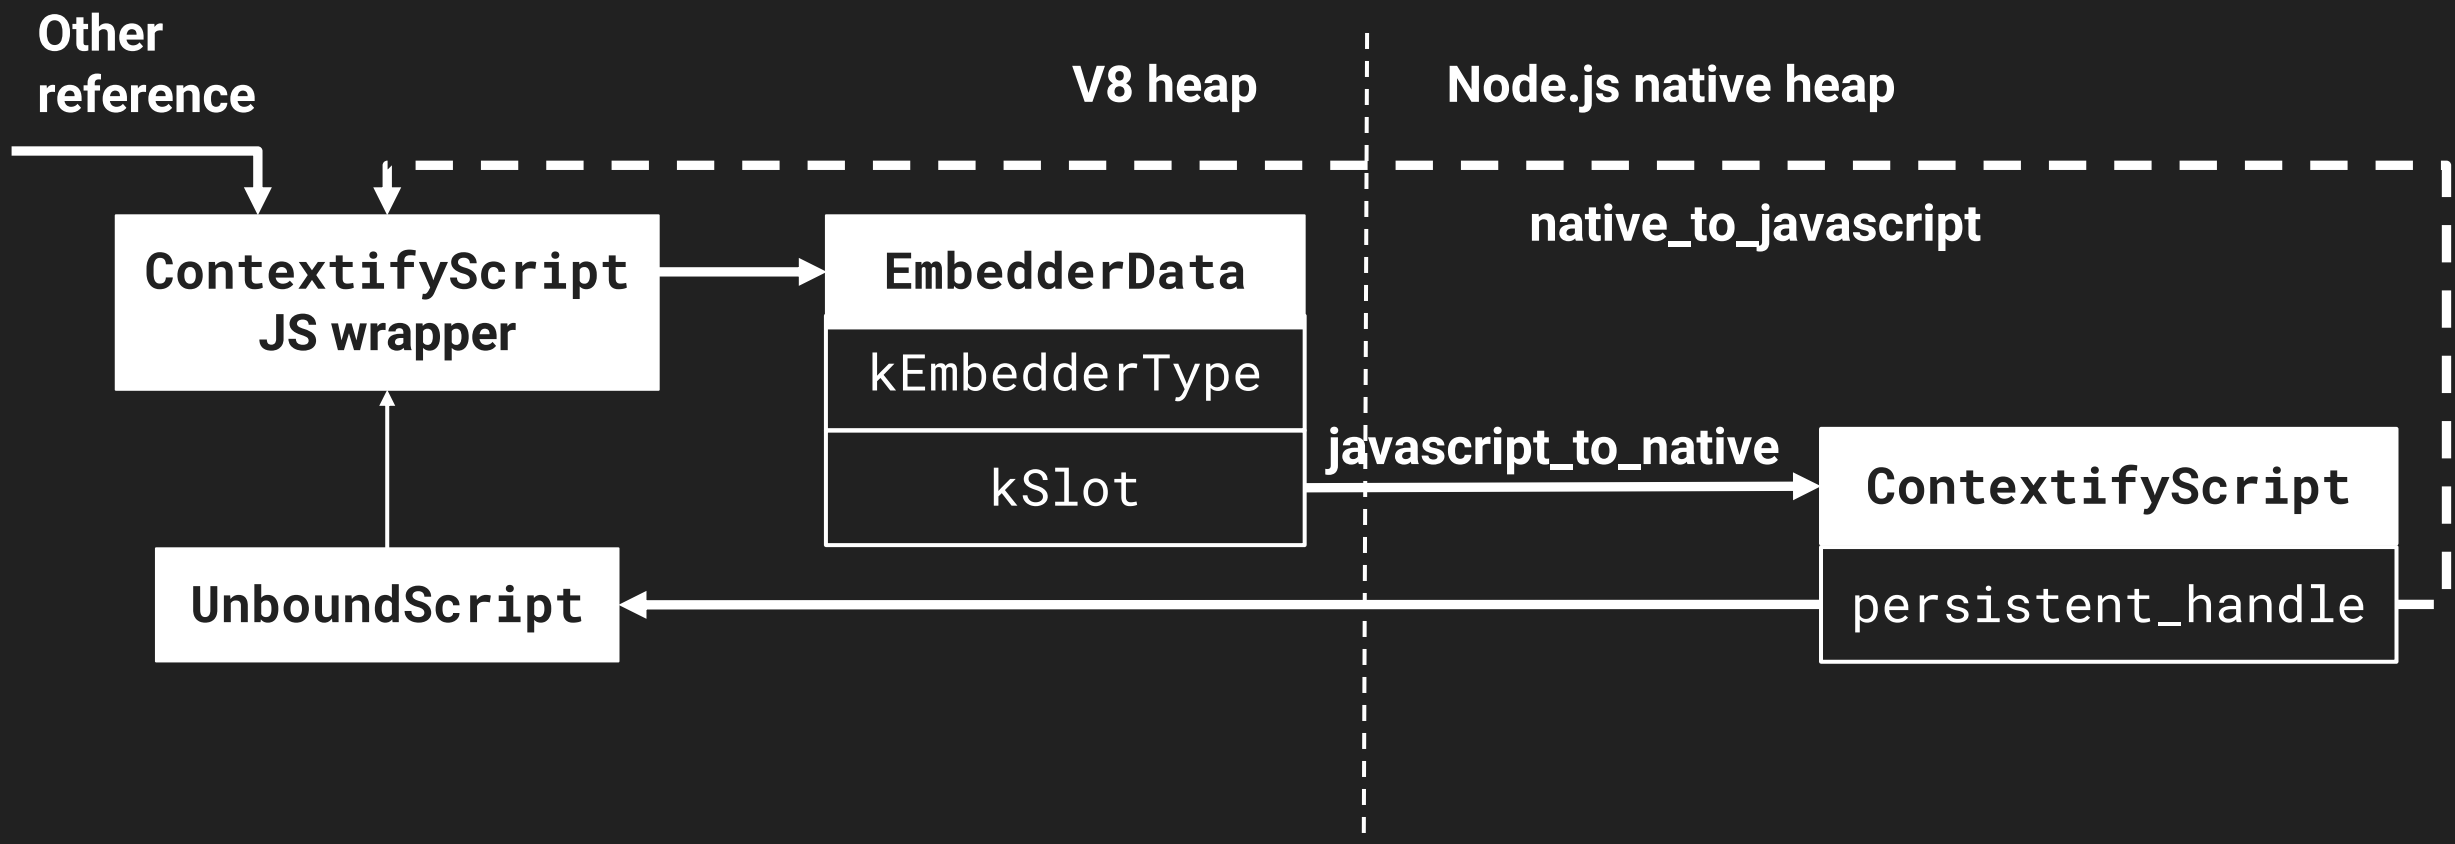
```
    v8::Global<v8::UnboundScript> script_;
```

```
}
```

After fixing JS land reference to ensure ContextifyScript is kept alive while import() can still be called from the Script...

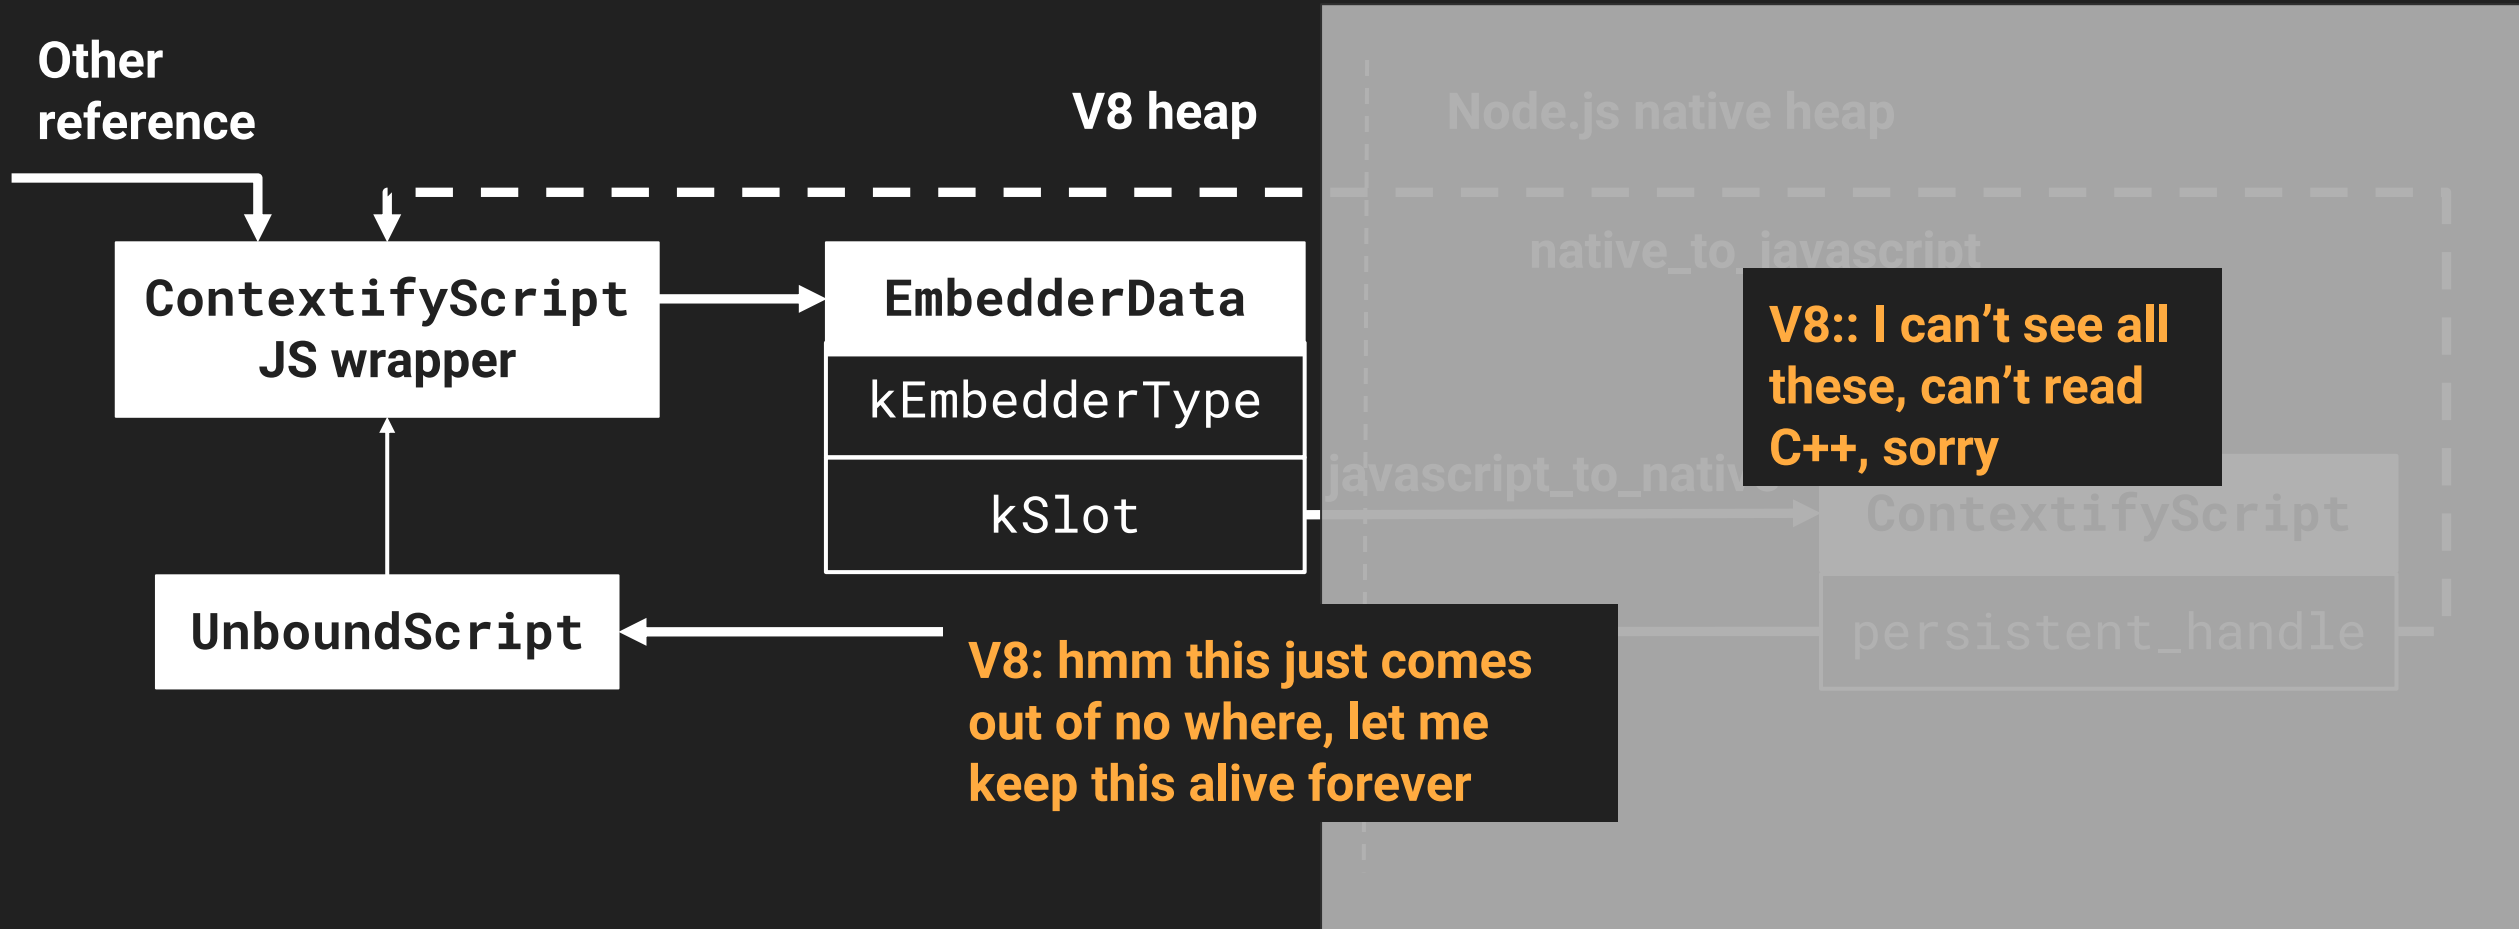
😓 This started to leak because the UnboundScript kept ContextifyScript alive in a cycle (similar to #42080)

# Case study of dynamic import() memory issues





# Case study of dynamic import() memory issues



# Case study of dynamic import() memory issues

```
class ContextifyScript : public BaseObject {
```

```
public:
```

```
ContextifyScript(Environment* env, Local<Object> object)
```

```
    : BaseObject(env, object) {
```

```
    script_.Reset(isolate, v8_script);
```

```
    MakeWeak();
```

```
    script_.SetWeak();
```

```
    object->SetInternalField(kUnboundScriptSlot, v8_script);
```

```
}
```

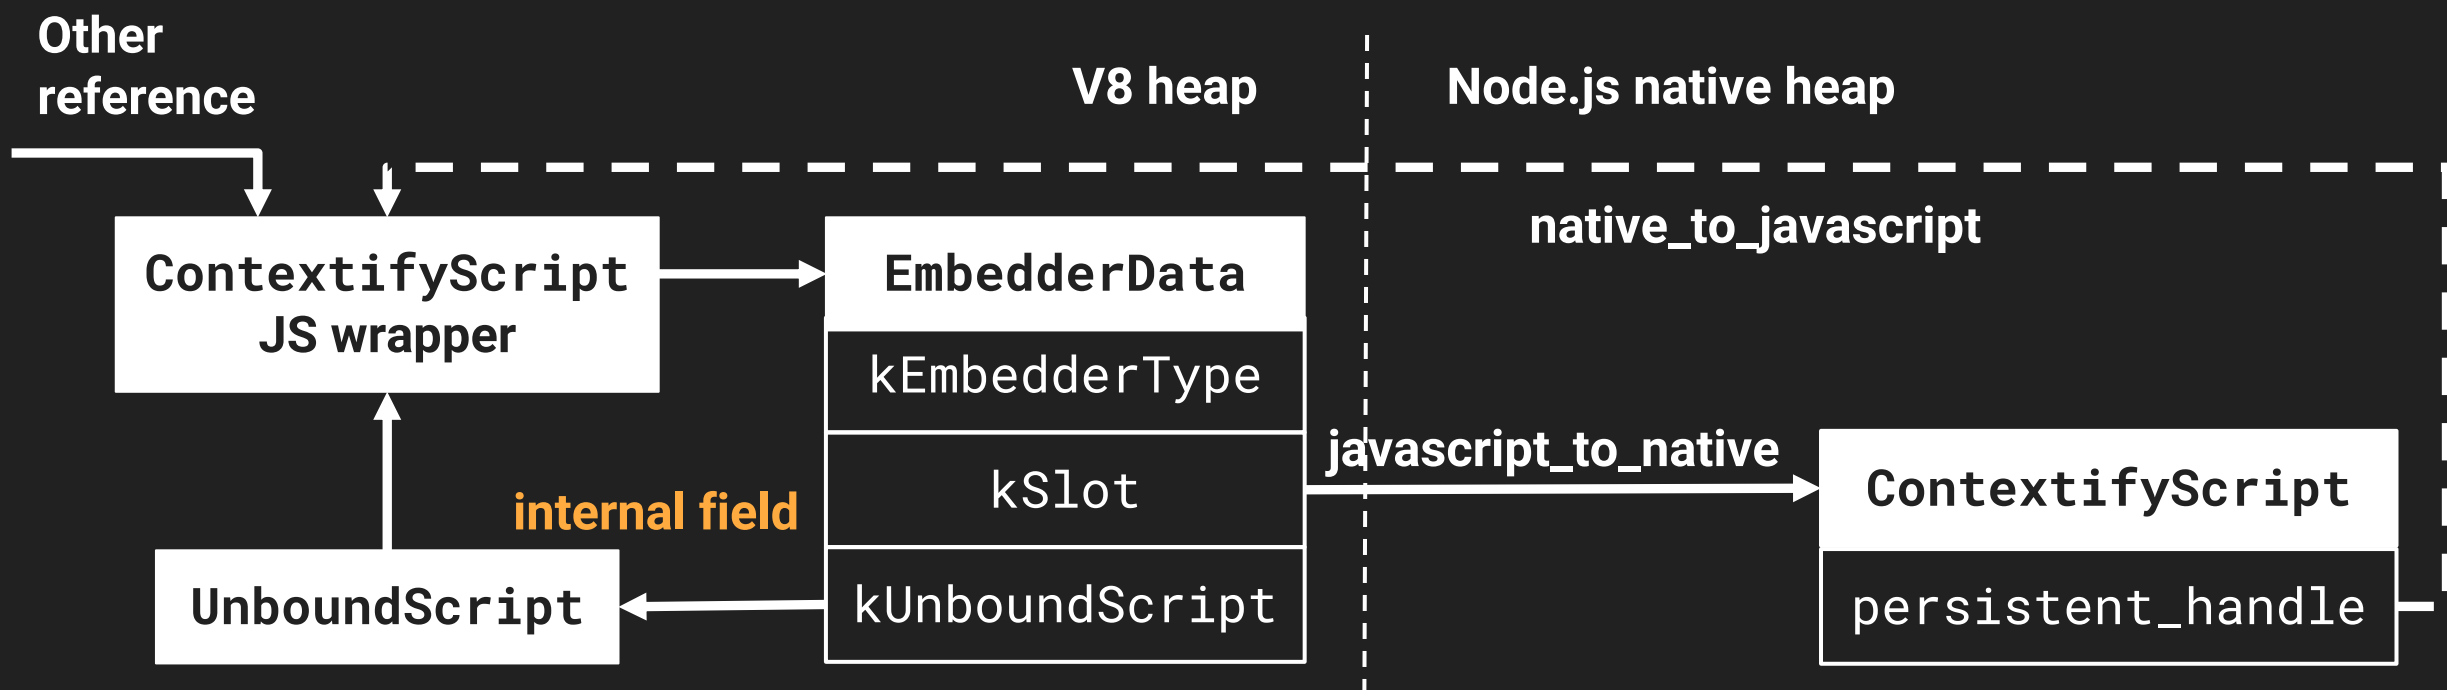
Patched V8 a bit to make this possible to fix the leak

```
private:
```

```
    v8::Global<v8::UnboundScript> script_;
```

```
}
```

# Case study of dynamic import() memory issues

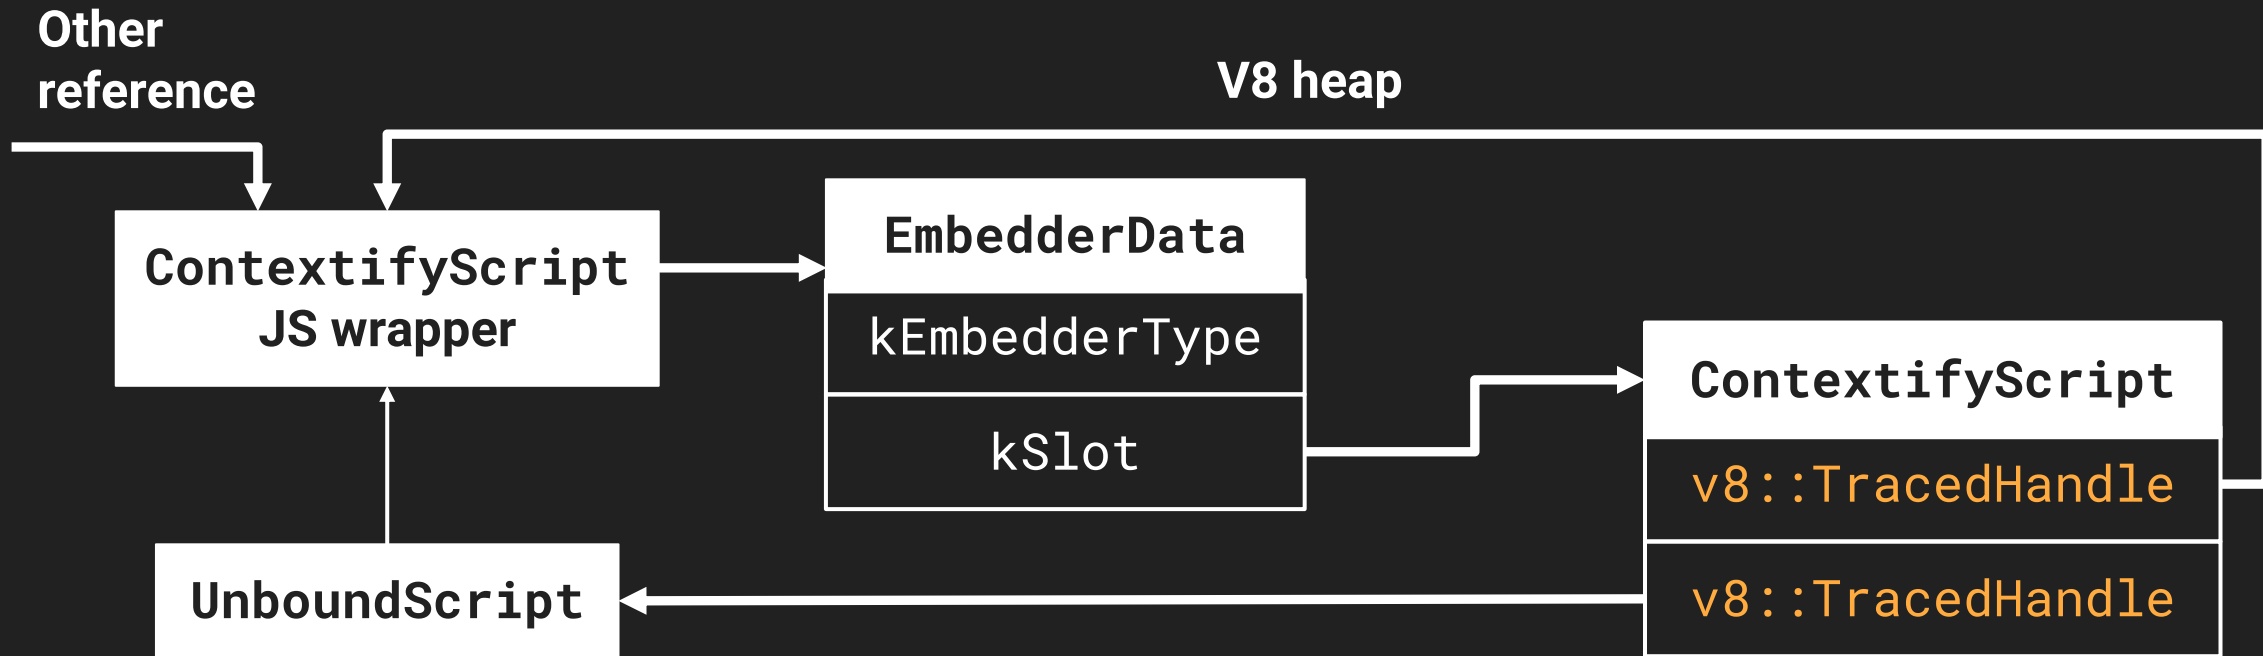


▼ Detached Node / ContextifyScript @50 ✂	–	40	0 %	40	0 %
▼ native_to_javascript :: Script @6393	–	64	0 %	9 904	0 %
▶ 2 :: (shared function info) @28793	–	64	0 %	9 840	0 %
▶ __proto__ :: ContextifyScript @78335	12	24	0 %	1 120	0 %
▶ map :: system / Map @78365	–	72	0 %	216	0 %
▶ sourceMappingURL :: system / Oddball @67 ☐	2	48	0 %	120	0 %
▶ javascript_to_native :: Detached Node / ContextifyScript @50 ✂	–	40	0 %	40	0 %
▶ <symbol node:host defined option symbol> :: symbol @6559	8	24	0 %	24	0 %

# A path to Oilpan

- V8's C++ GC library
- Can be used to create C++ <-> JS references that V8 knows how to track to avoid issues caused by cycles
- Bootstrapped in core with some cctests
- Currently working on migration plan ([design doc](#))

# A path to Oilpan (WIP PR: #52295)



▼ Node / ContextifyScript @6393

▶ [5] :: Node / ContextifyScript @6393

▶ [6] :: (shared function info) @28793

▶ \_\_proto\_\_ :: ContextifyScript @51179

▶ map :: system / Map @67851

▶ sourceMapURL :: system / Oddball @67

▶ <symbol node:host\_defined\_option\_symbol> :: symbol @6559

traced handles

–	104	0 %	10 456	0 %
–	104	0 %	10 456	0 %
–	64	0 %	10 352	0 %
12	24	0 %	1 120	0 %
–	72	0 %	216	0 %
2	48	0 %	120	0 %
8	24	0 %	24	0 %

# A path to Oilpan (WIP PR: #52295)

- `crypto.createHash()` can be ~27% faster after migration. The whole hashing operation can be ~10% faster. [#51017](#)
- Blocker: external memory tracking in heap snapshots. Working on a new V8 API.
- Blocker: saw minor regressions in some APIs (V8 serdes), investigating impact & why.
- Blocker: Need to be able to trace `v8::Data`. Working on an upstream patch [crrev:5403888](#)

# How to analyze memory issues

- Good old heap snapshots: `v8.writeHeapSnapshot()` or `--heapsnapshot-near-heap-limit`
- `--heap-profiler-show-hidden-objects`
- `--heap-prof`
- `vm.measureMemory()` etc. can provide hints
- Sometimes `lnode` helps too
- Discuss: what else can we provide? What are some common memory issues that can use better tooling support from Node.js?