

December 11-12, 2019
Montreal, Canada


How Node.js Bootstraps Itself: 2019 Edition

Joyee Cheung, Igalia



December 11-12, 2019
Montreal, Canada

About me

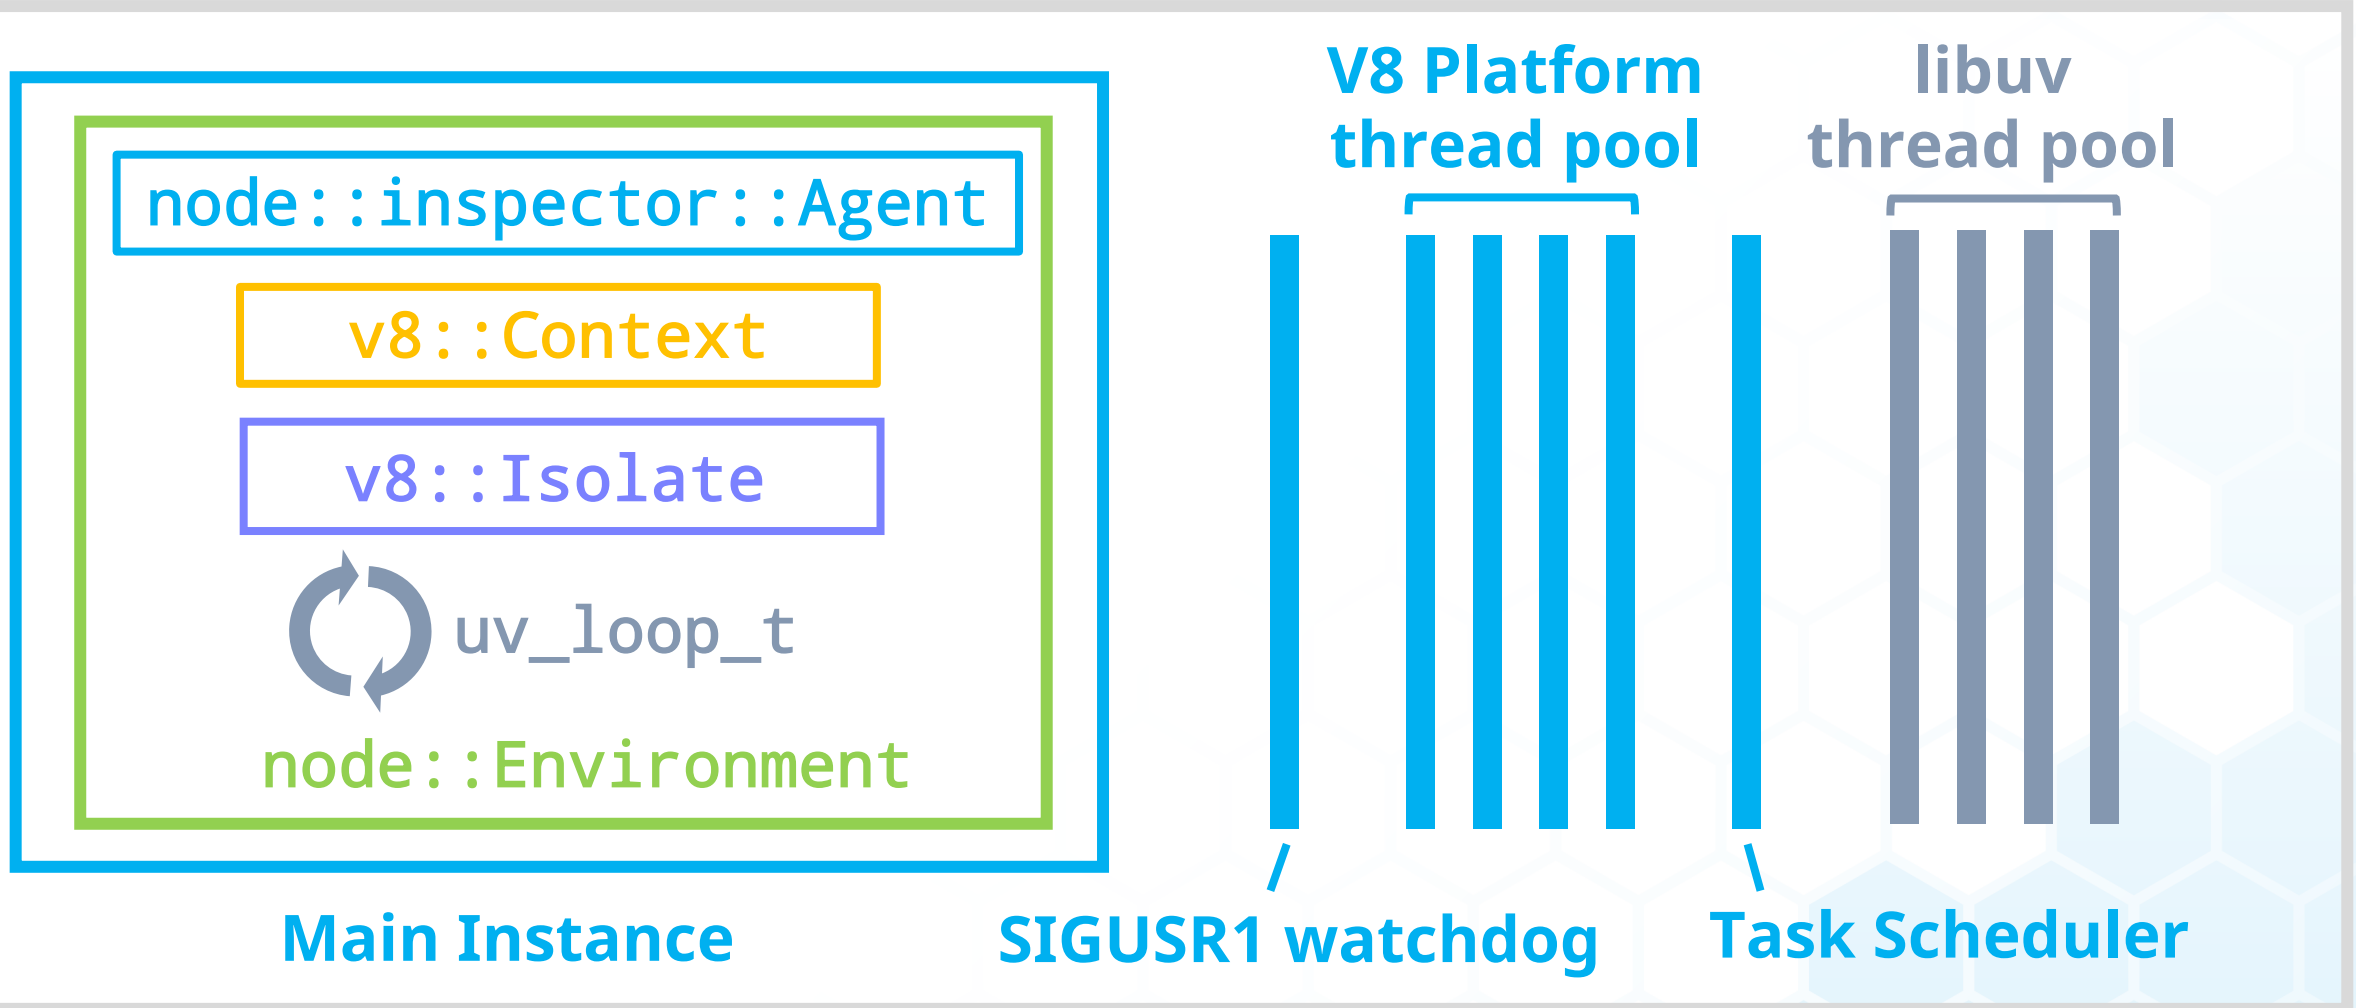


- Joyee Cheung (Cantonese) / Qiuyi Zhang (Mandarin)
- Compilers @ Igalia
- Node.js TSC & V8 Committer
- Champion of the startup performance initiative in Node.js
 - <https://github.com/nodejs/TSC/blob/master/Strategic-Initiatives.md>
- @joyeecheung on GitHub & Twitter

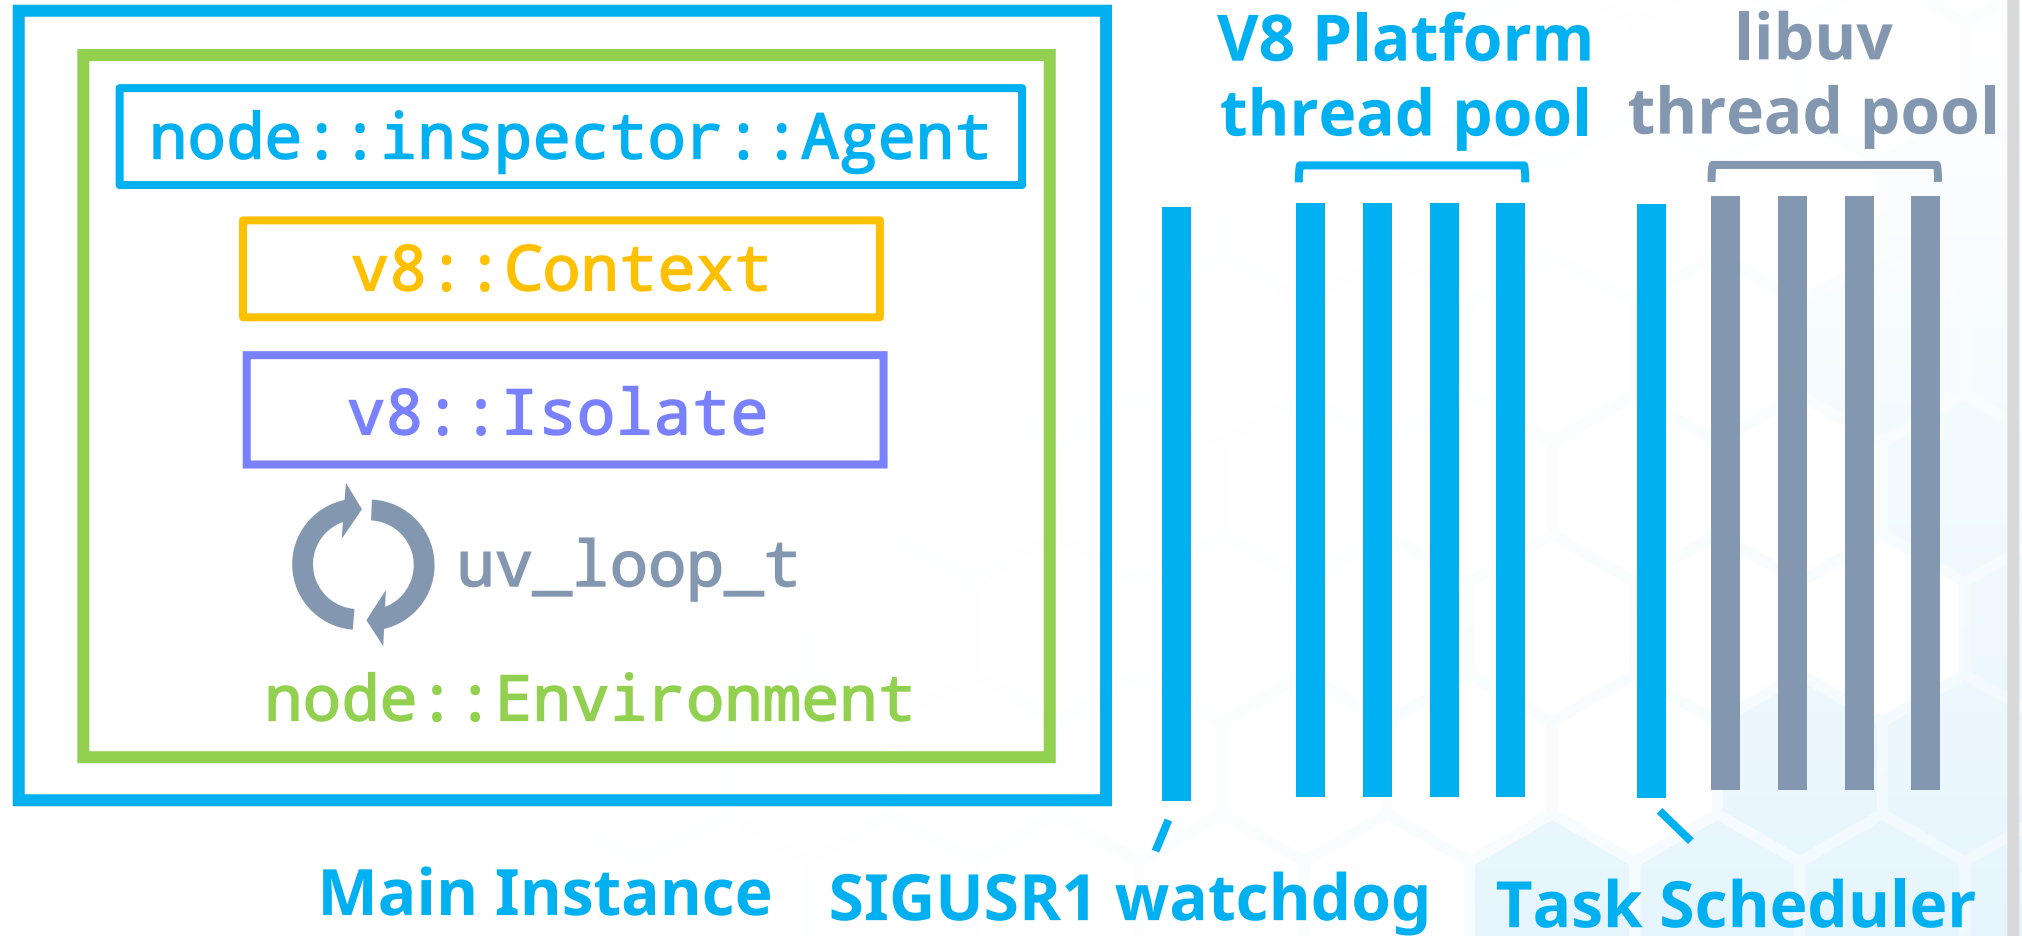
Slides of this talk:

https://github.com/joyeecheung/talks/blob/master/node_js_interactive_2019/how-node-js-bootstraps-itself-2019-edition.pdf


Overview of a Node.js process



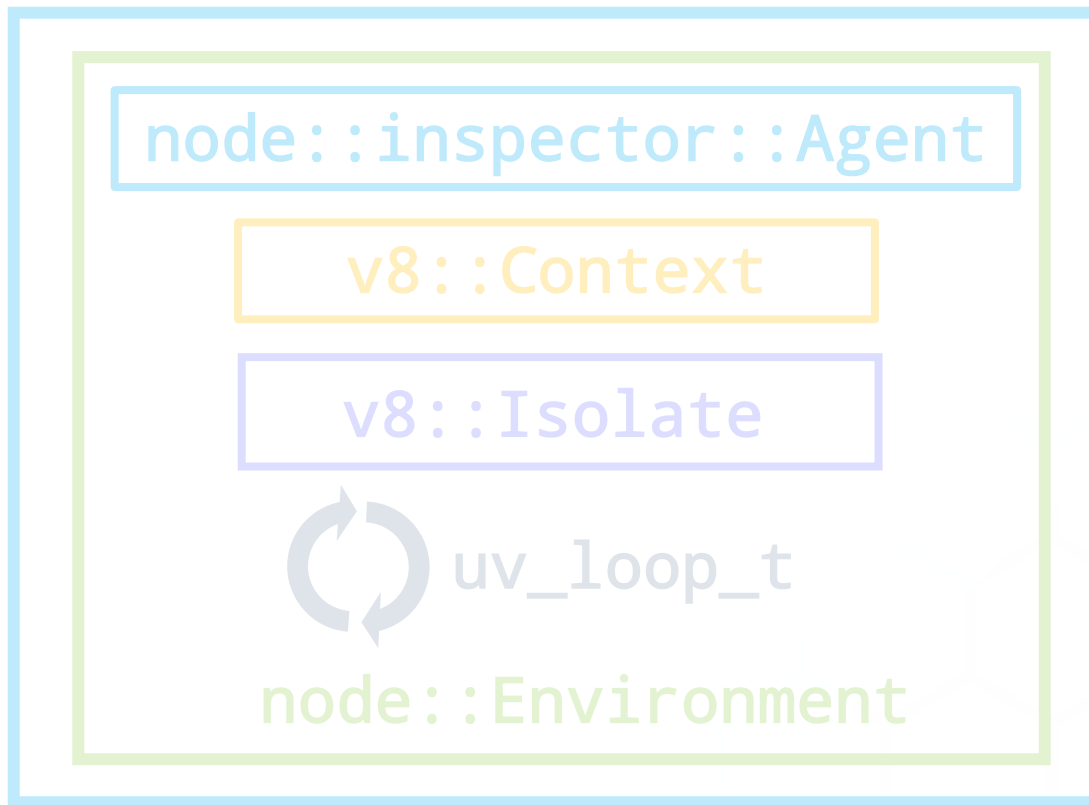
Creating a Worker



Creating a Worker




Worker



Main Instance


V8 Platform
thread pool

libuv
thread pool




SIGUSR1 watchdog

Task Scheduler





December 11-12, 2019
Montreal, Canada

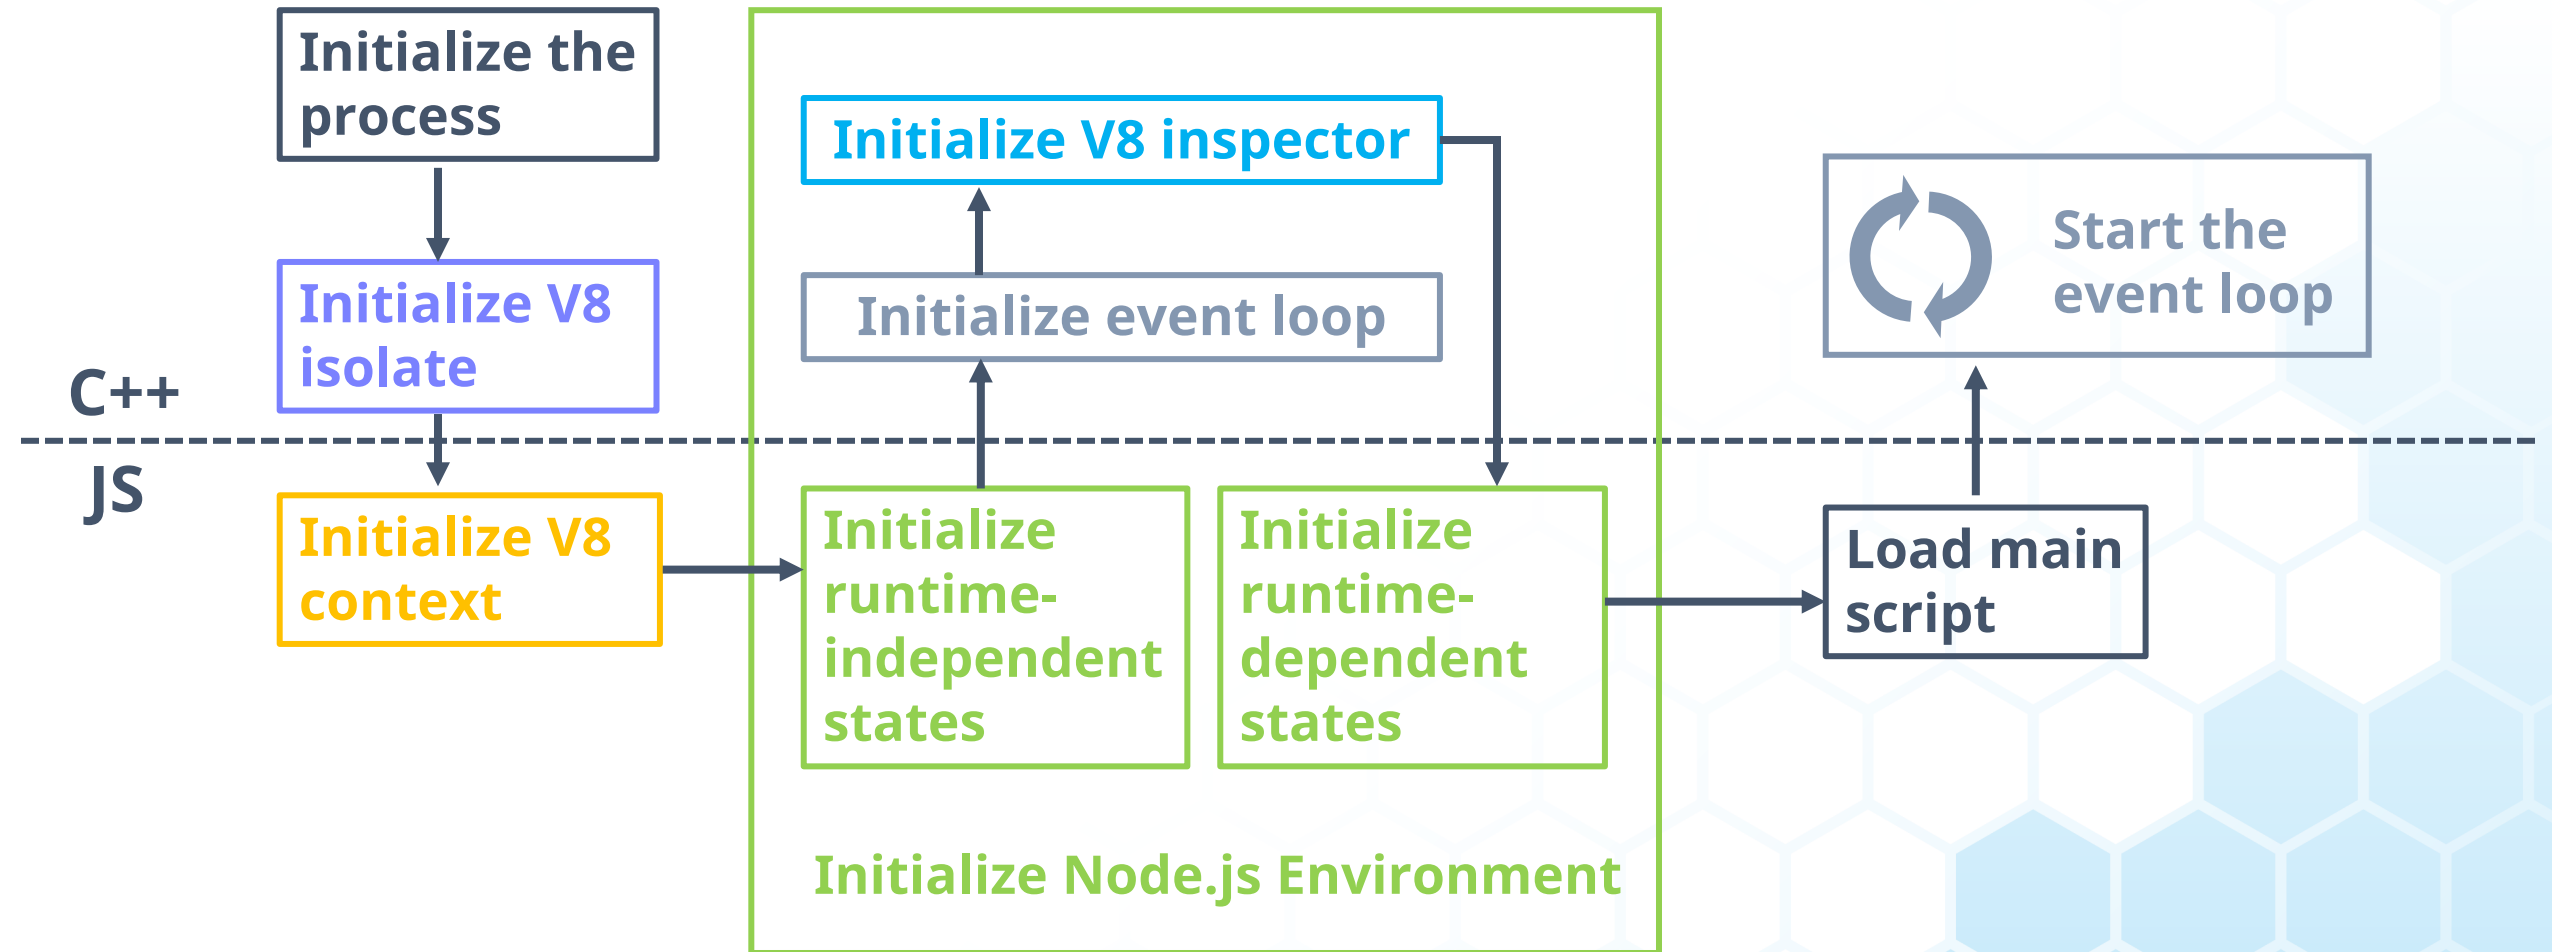
Creating a child process



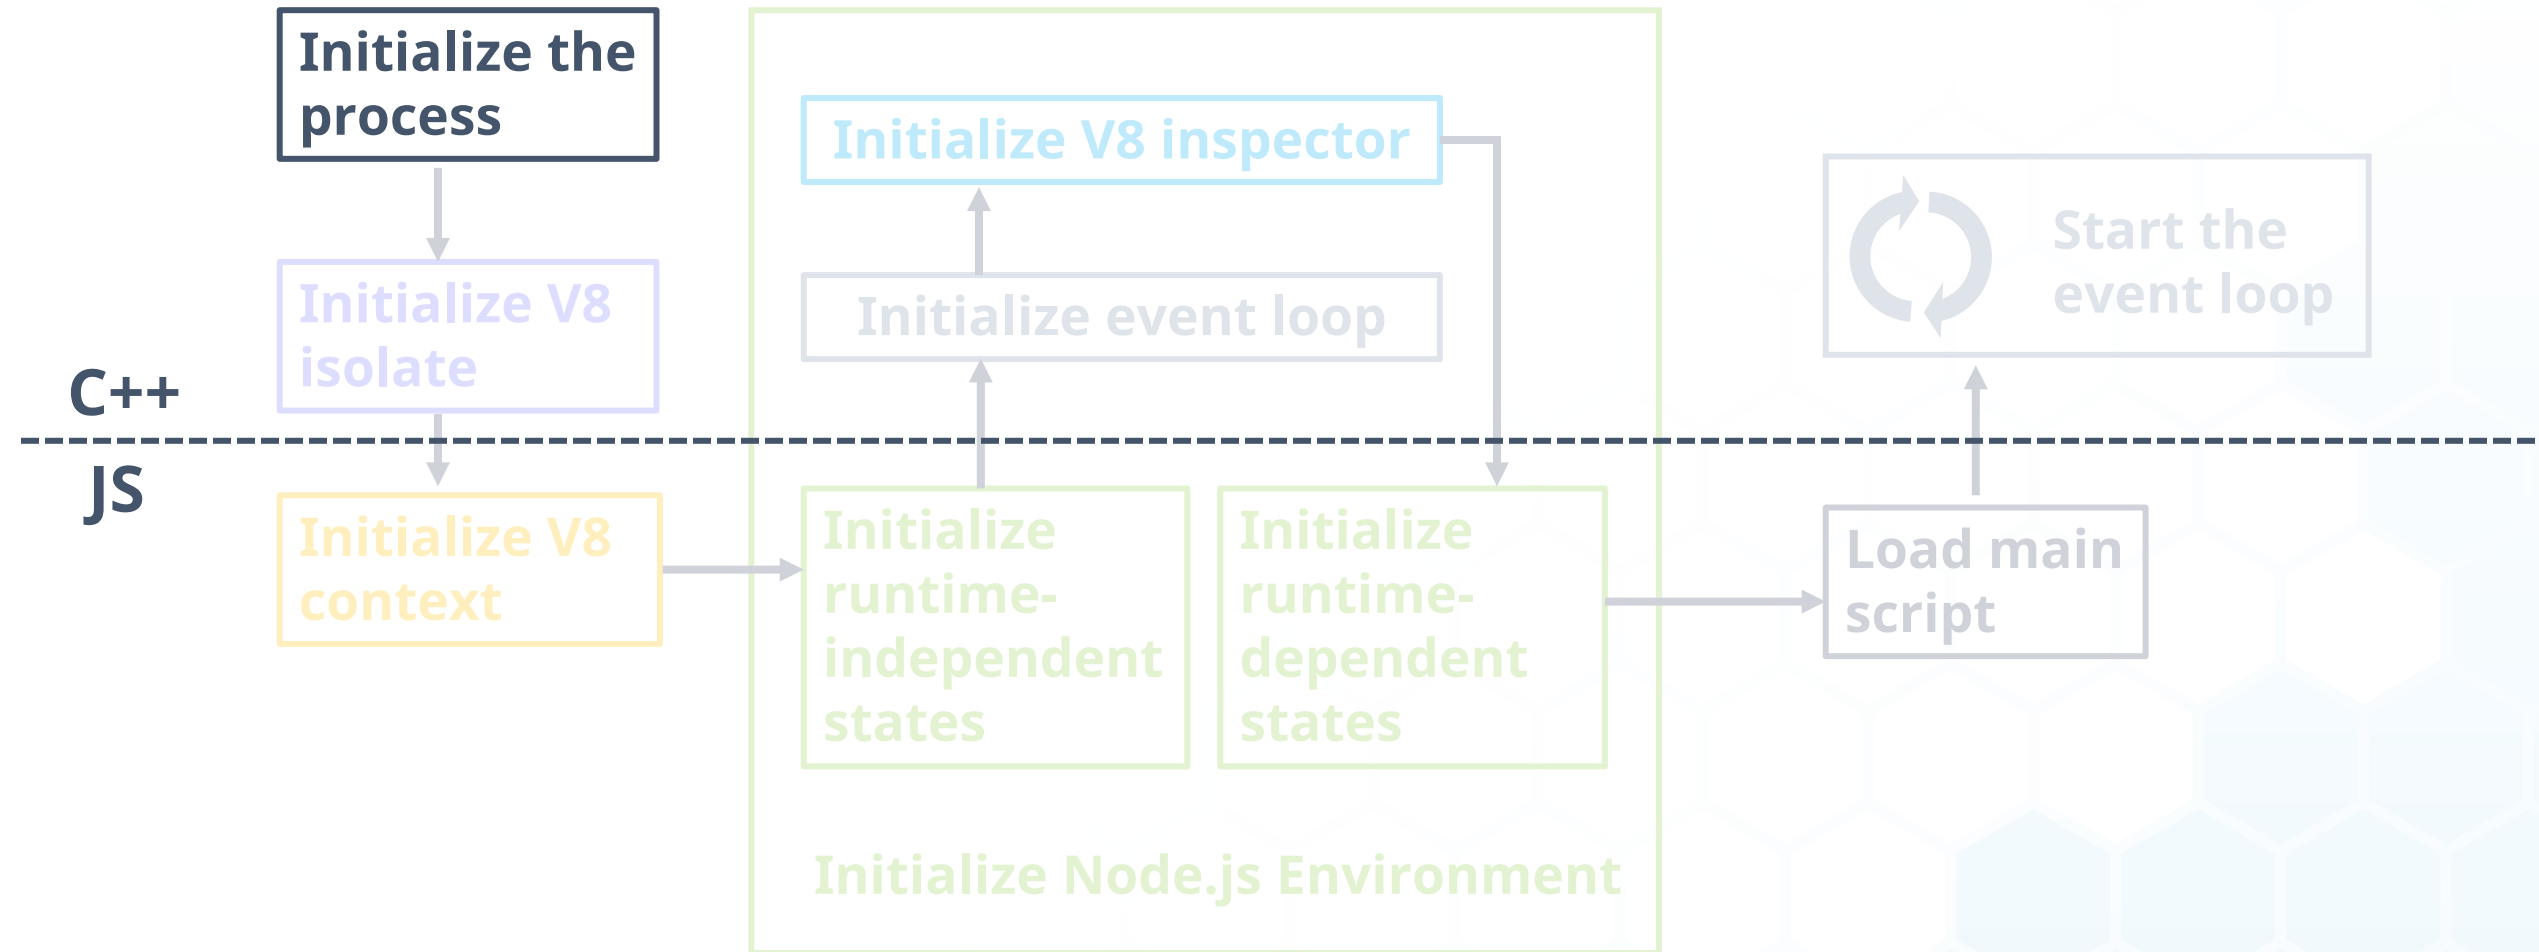
Creating a child process



Overview of the bootstrap



Setting up the process





December 11-12, 2019
Montreal, Canada

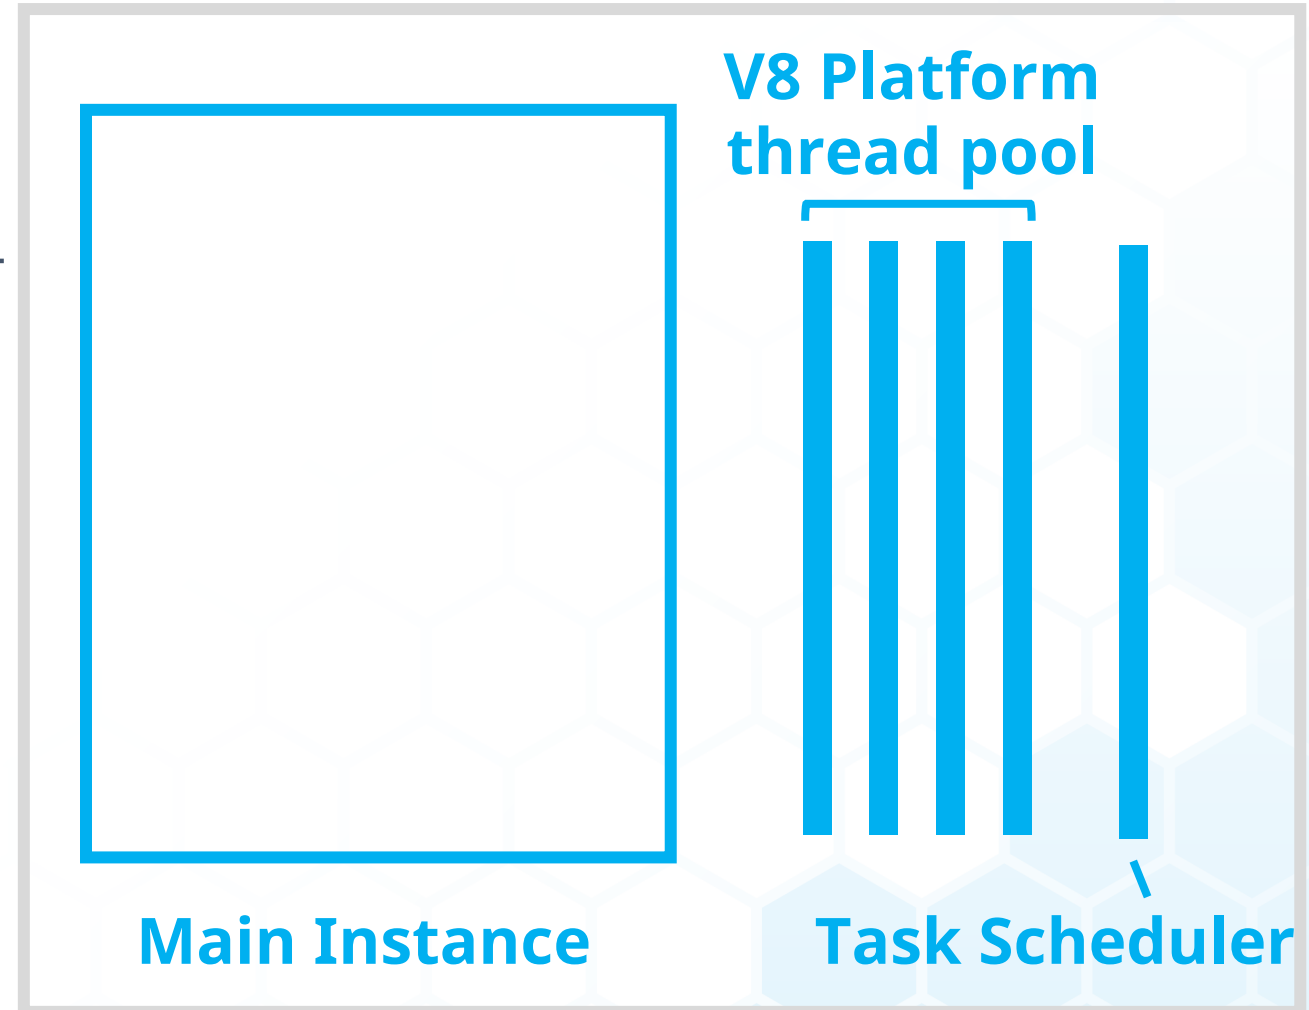
Setting up the process

- Setup signal handlers
- Parse CLI arguments (strings to C++ structures)
- Initialize ICU, OpenSSL



Setting up the process

- Setup signal handlers
- Parse CLI arguments (strings to C++ structures)
- Initialize ICU, OpenSSL
- **Initialize the V8 platform**

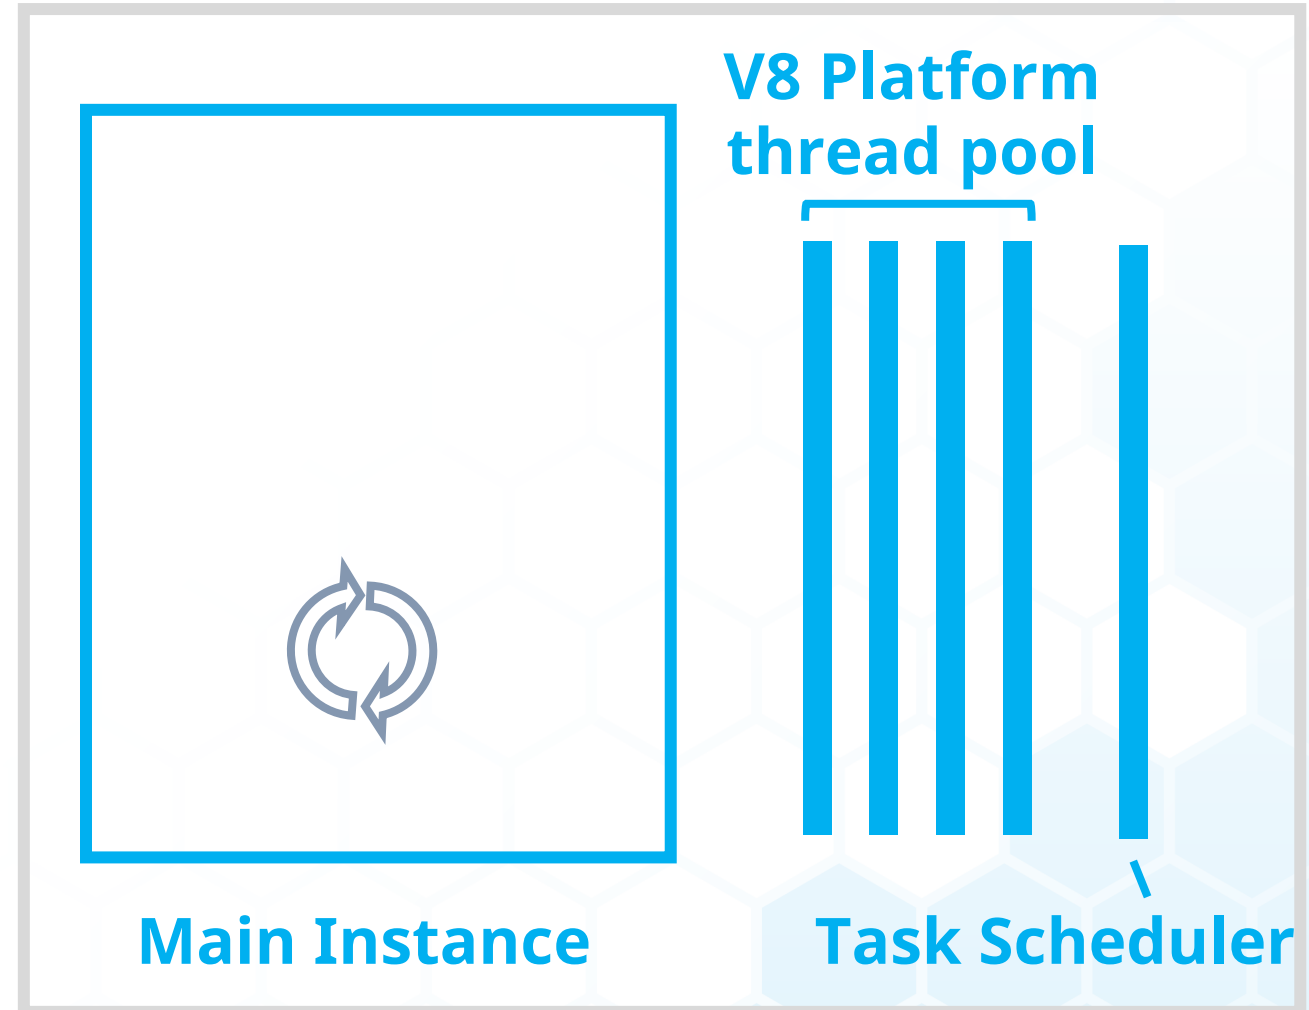


Setting up the process

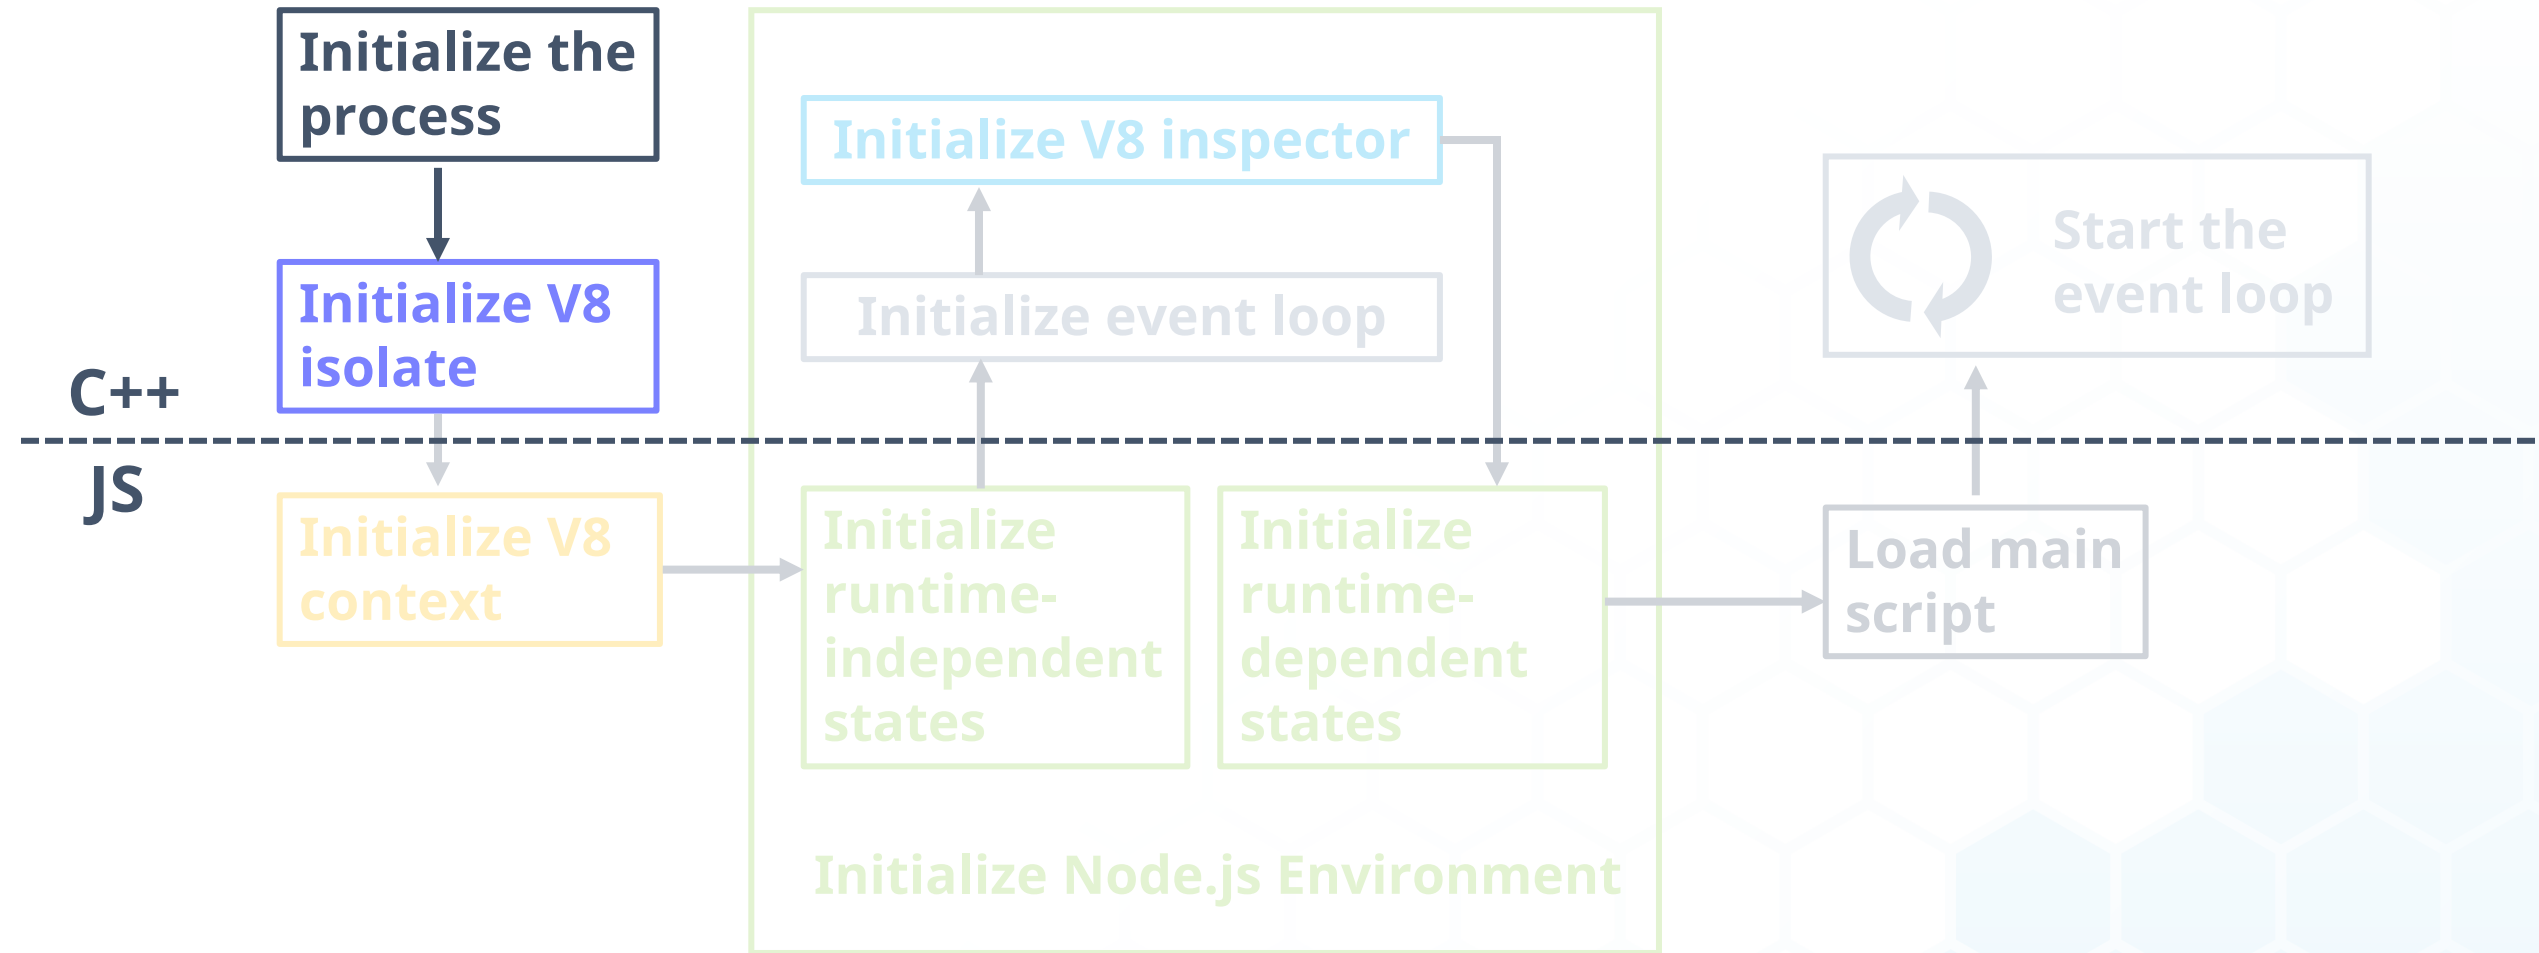
- Setup signal handlers
- Parse CLI arguments (string to C++ structures)
- Initialize ICU, OpenSSL
- Initialize the V8 platform

For the main instance

- Initialize the **default libuv event loop**



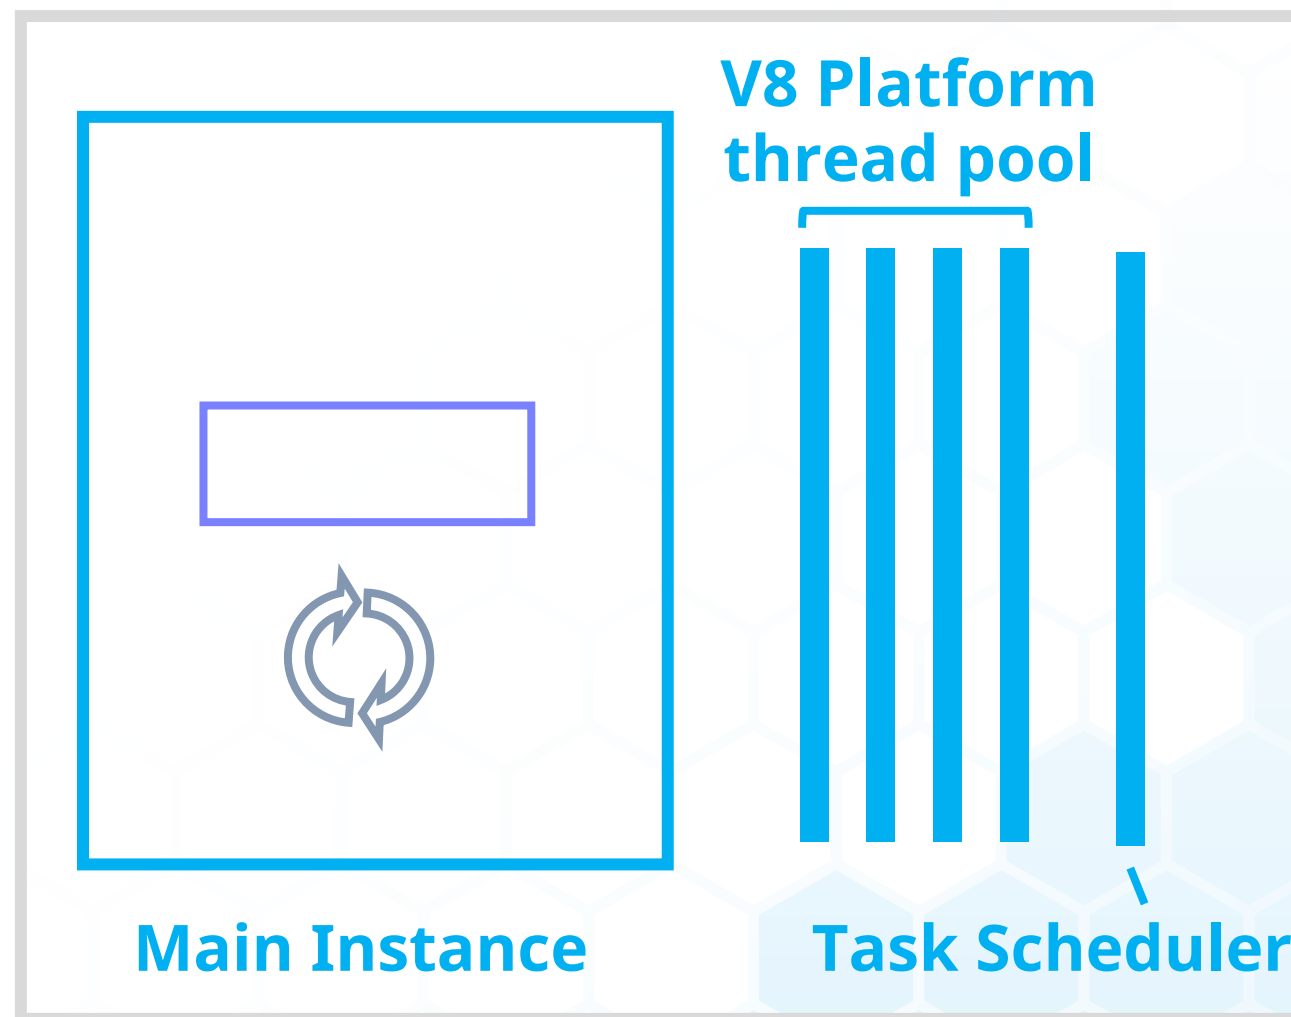
Setting up the V8 isolate



Setting up the V8 isolate

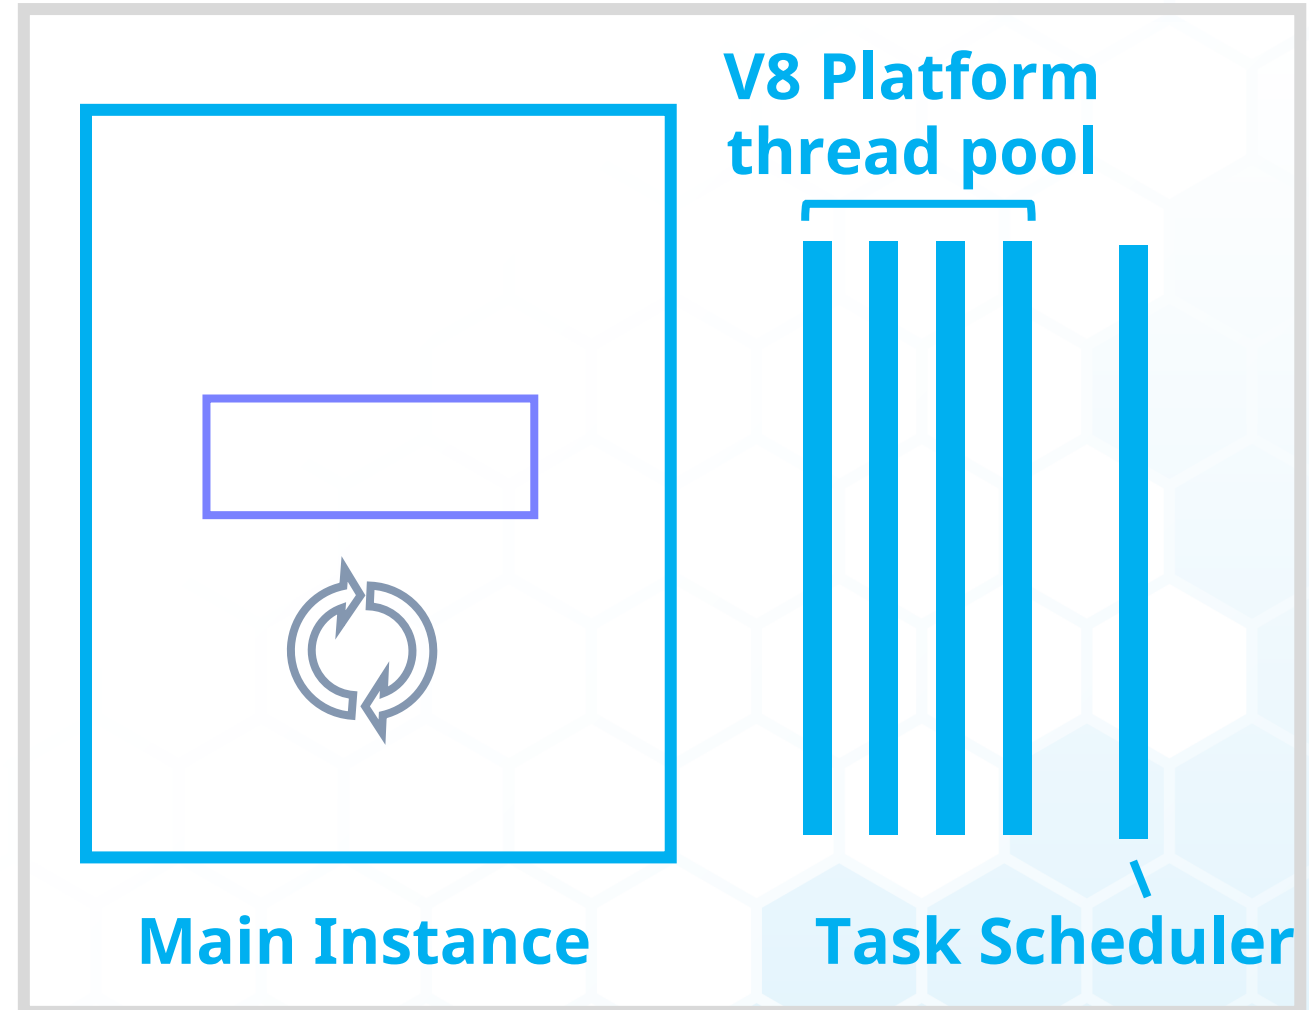
What's a V8 isolate?

- `v8::Isolate` is the instance of the v8 JavaScript engine
- Encapsulates the JS heap, microtask queue, pending exceptions...



Setting up the V8 isolate

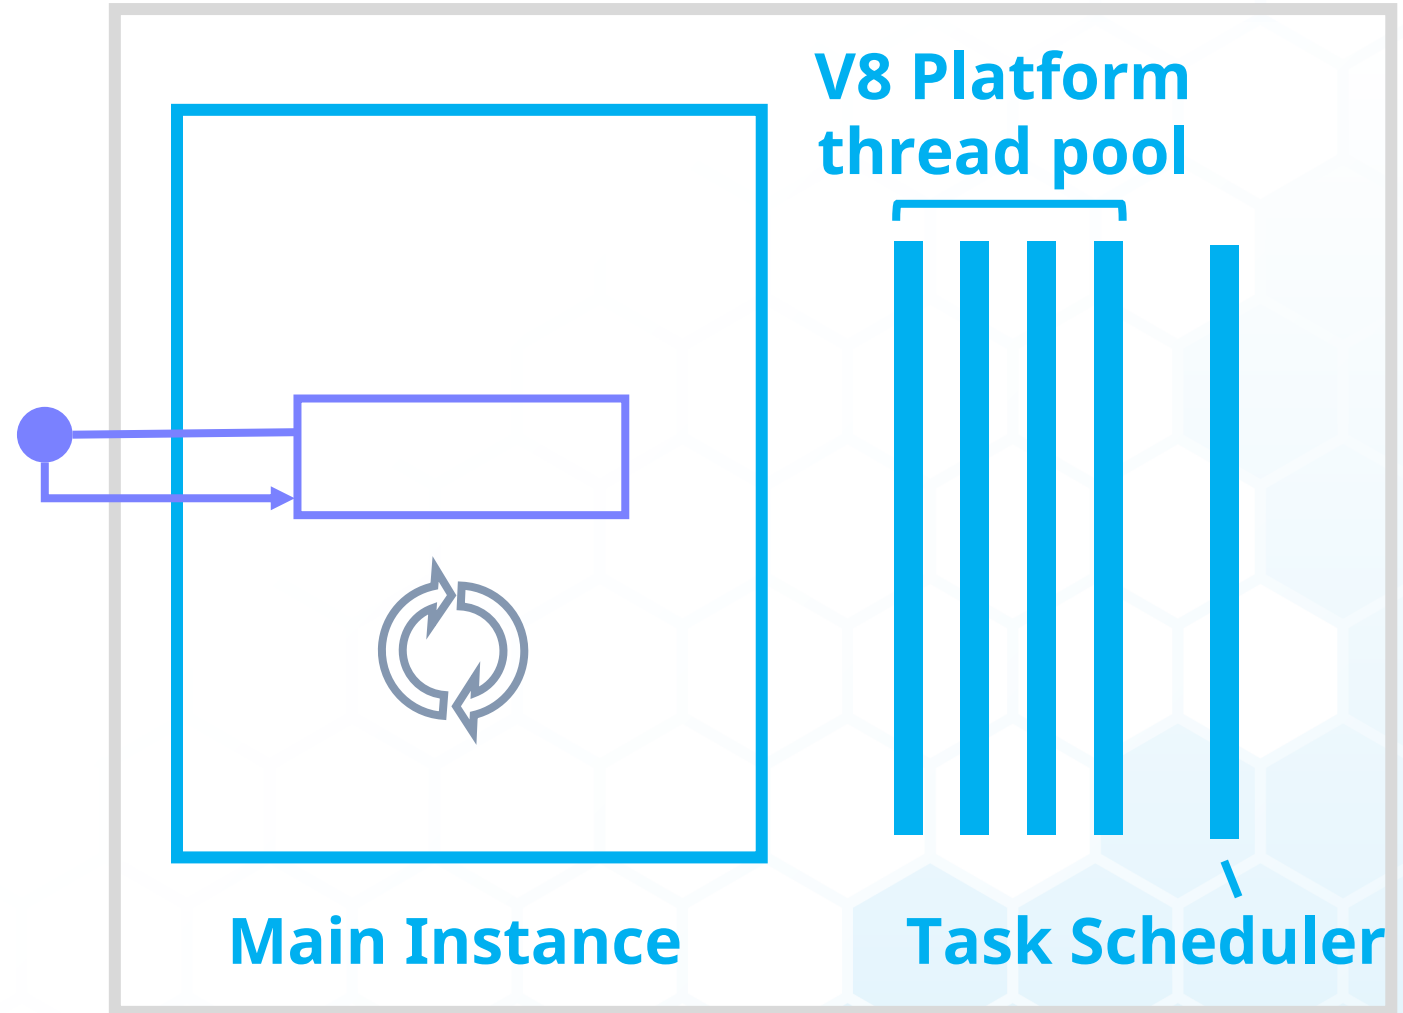
- Configure resource constraints (e.g. memory)
- Create array buffer allocator
- Deserialize from the V8 isolate snapshot



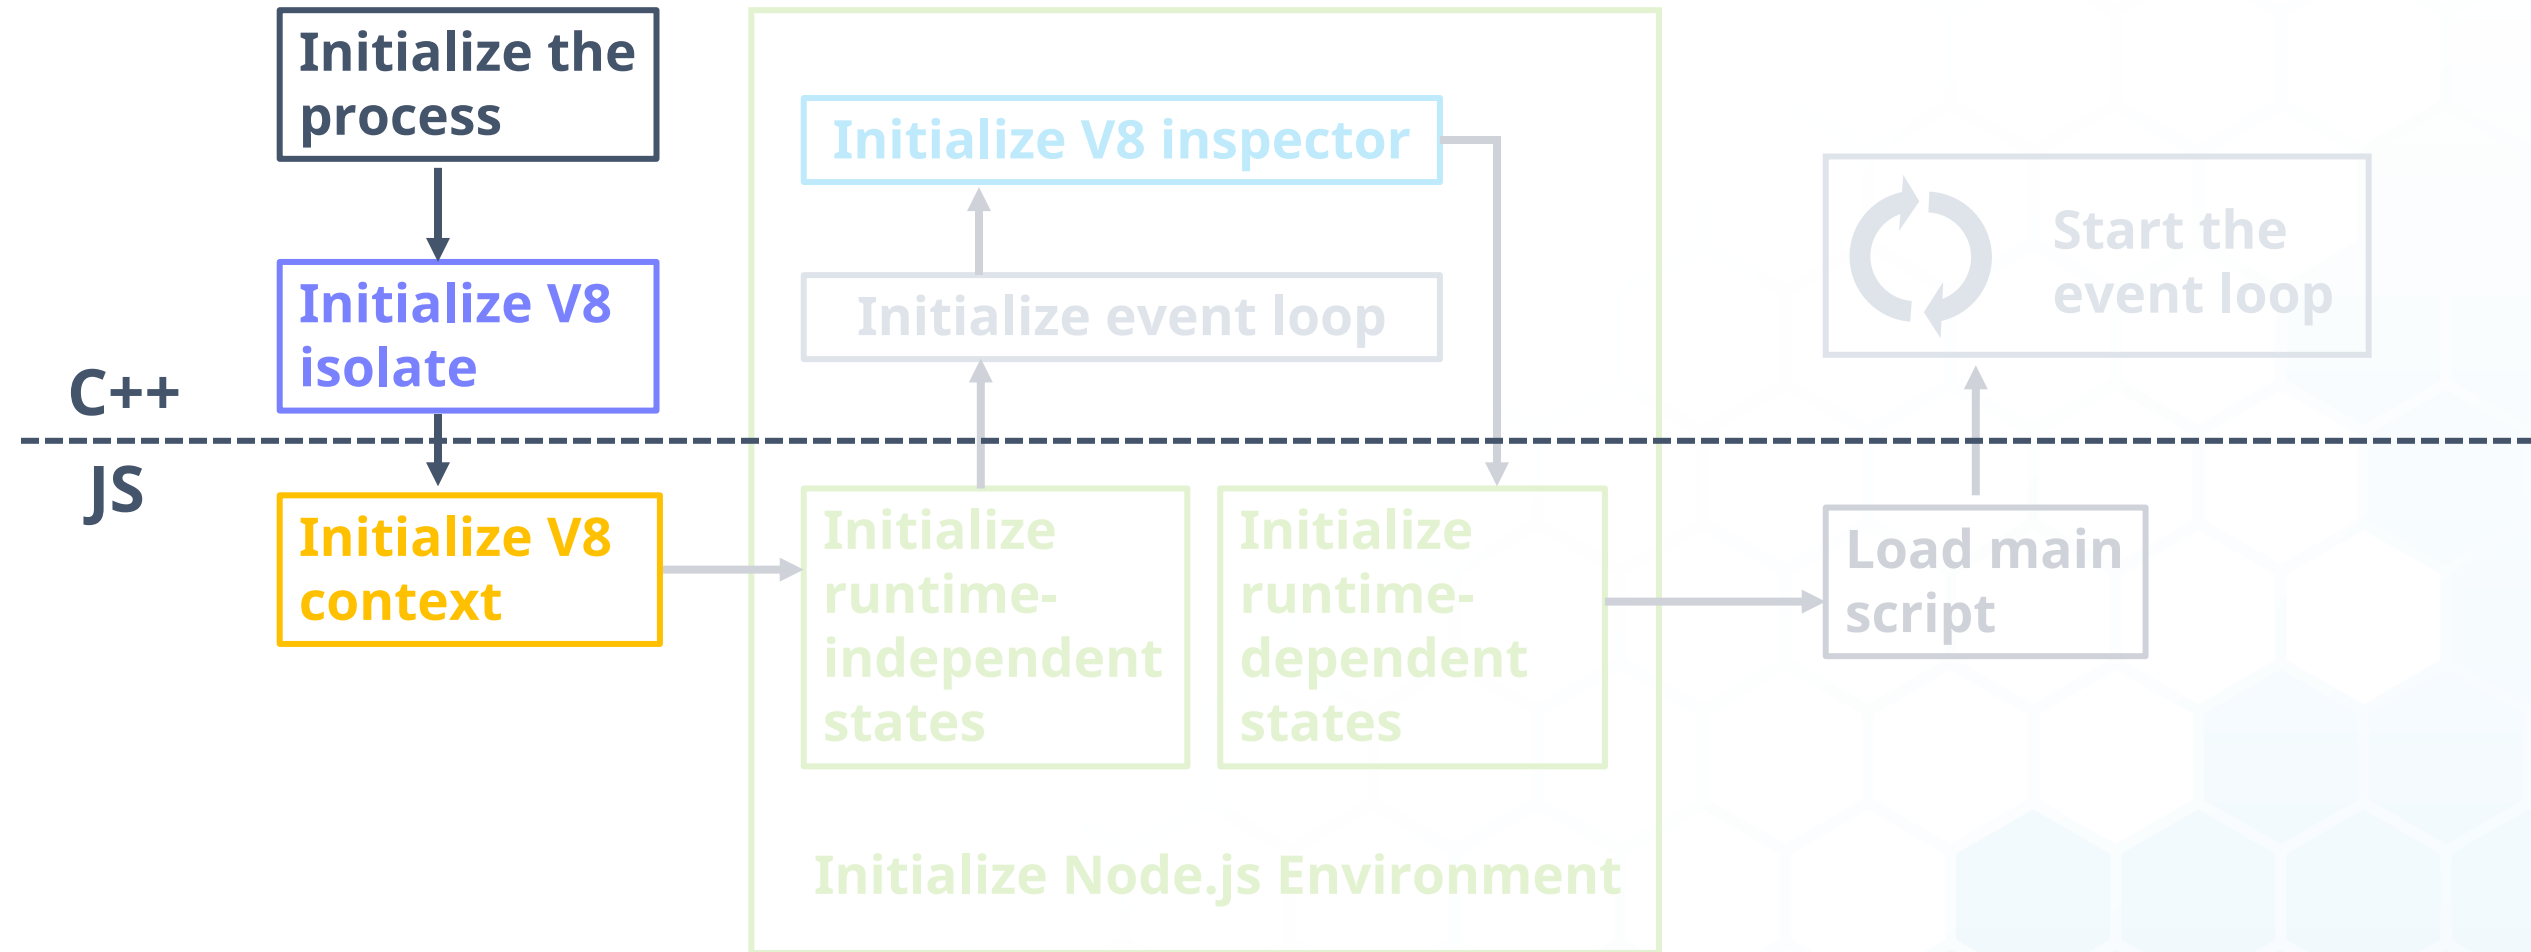
Setting up the V8 isolate

Setup various per-isolate callbacks in C++

- GC callbacks
- Uncaught exception listeners
- Promise rejection callbacks
- etc.



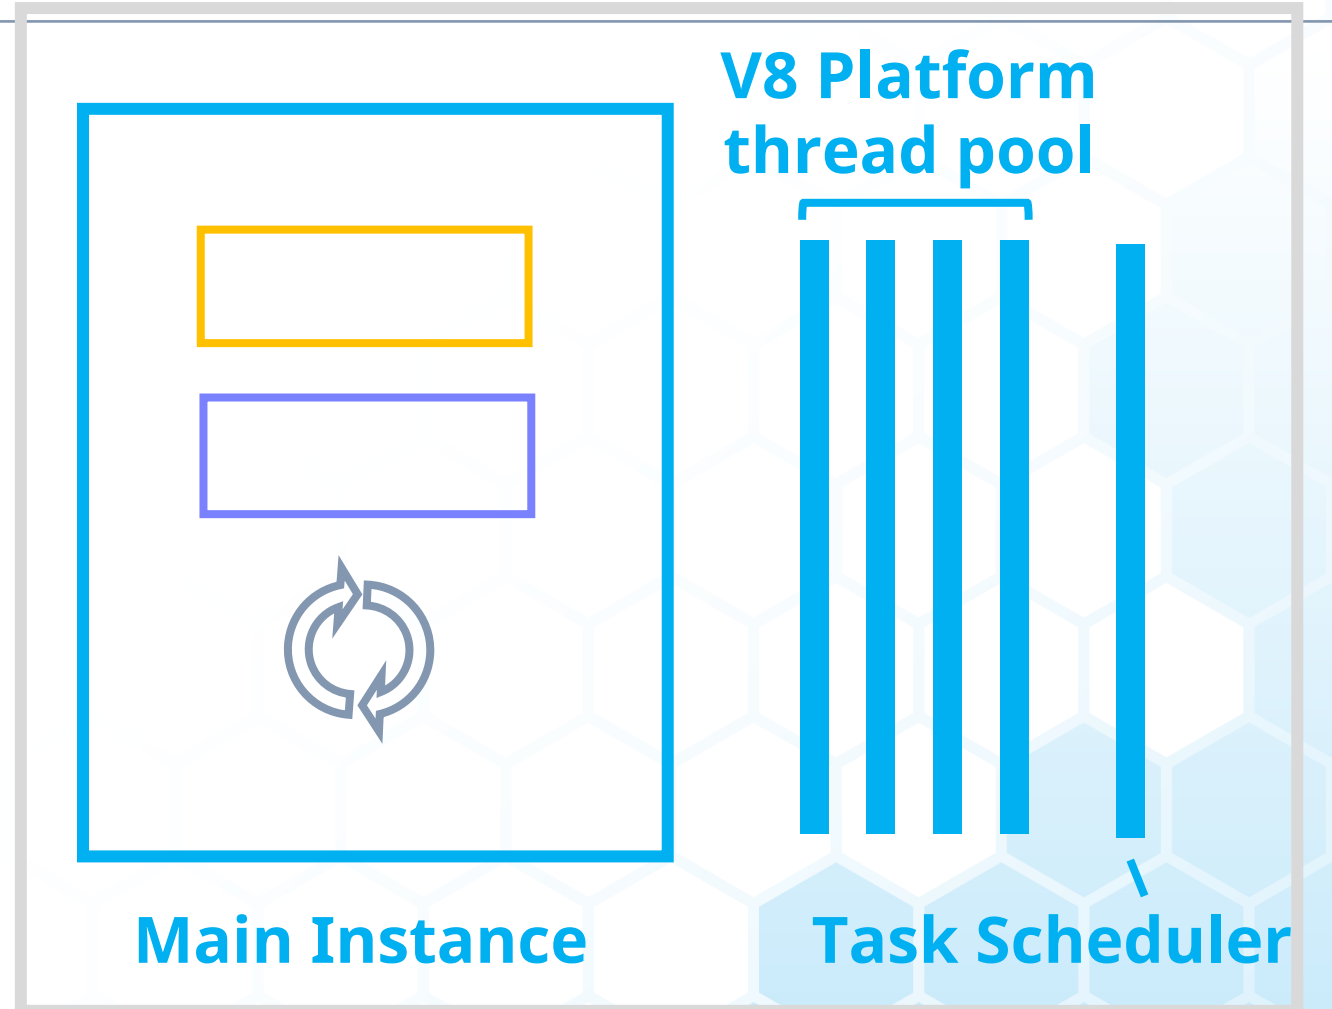
Setting up the V8 context



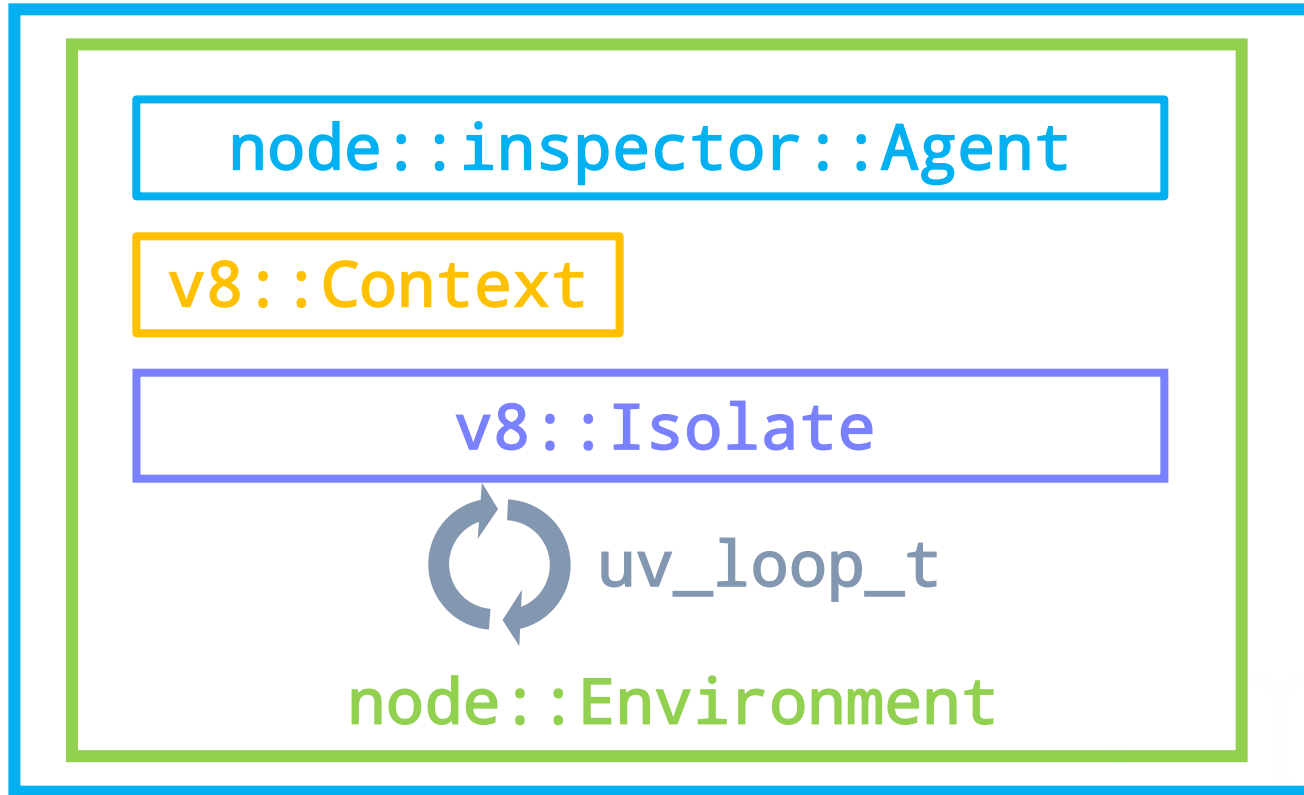
Setting up the V8 context

What's a V8 context?

- A sandboxed execution context
- Encapsulates JavaScript builtins (primordials) e.g. `globalThis`, `Array`, `Object`...
- What's inside the returned result of `vm.createContext()`



Creating a vm Context



Main Instance

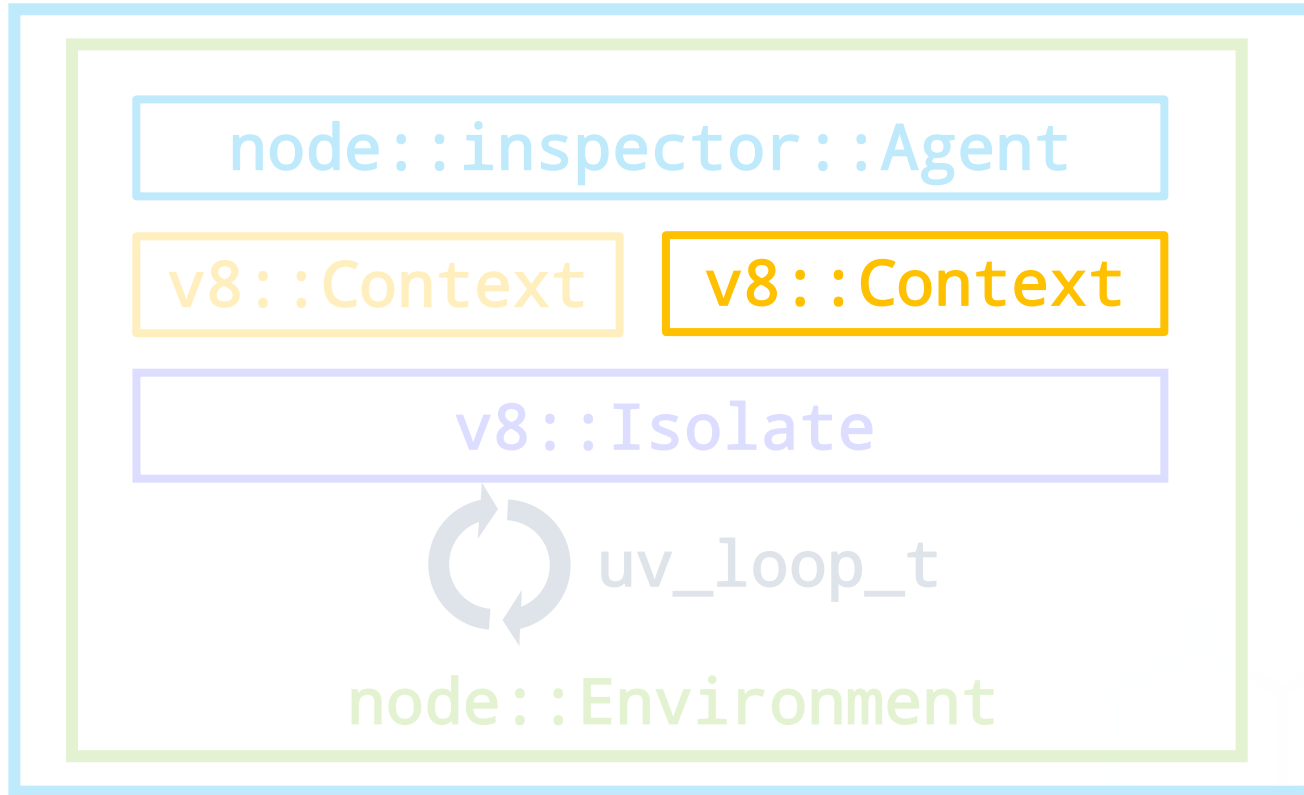
SIGUSR1 watchdog

V8 Platform
thread pool

libuv
thread pool

Task Scheduler

Creating a vm Context



Main Instance

SIGUSR1 watchdog

V8 Platform
thread pool

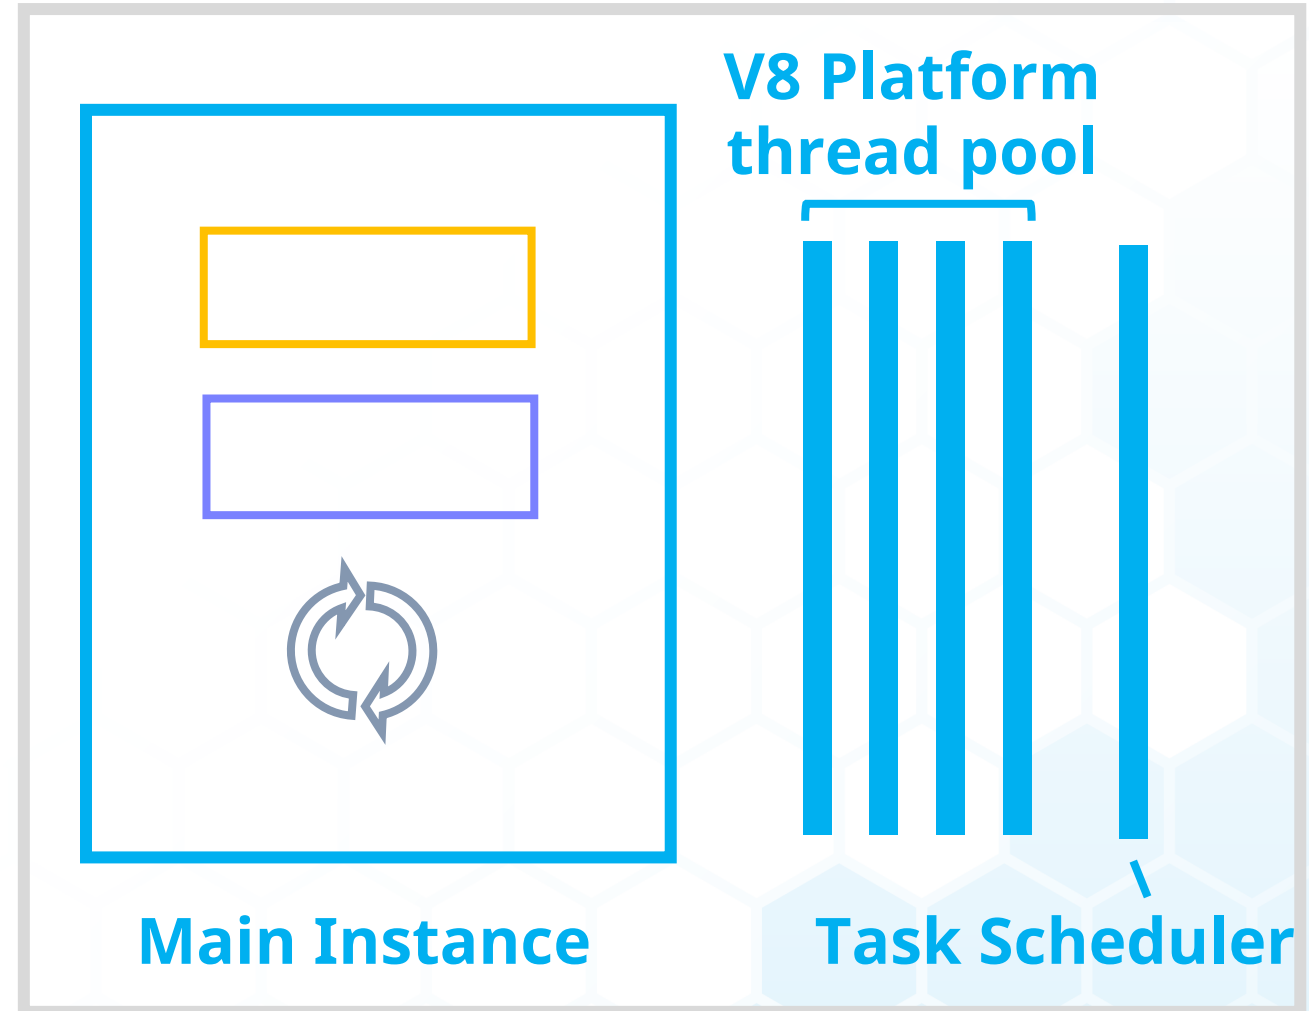
libuv
thread pool

Task Scheduler

Setting up the V8 context

What's inside a Node.js's V8 context?

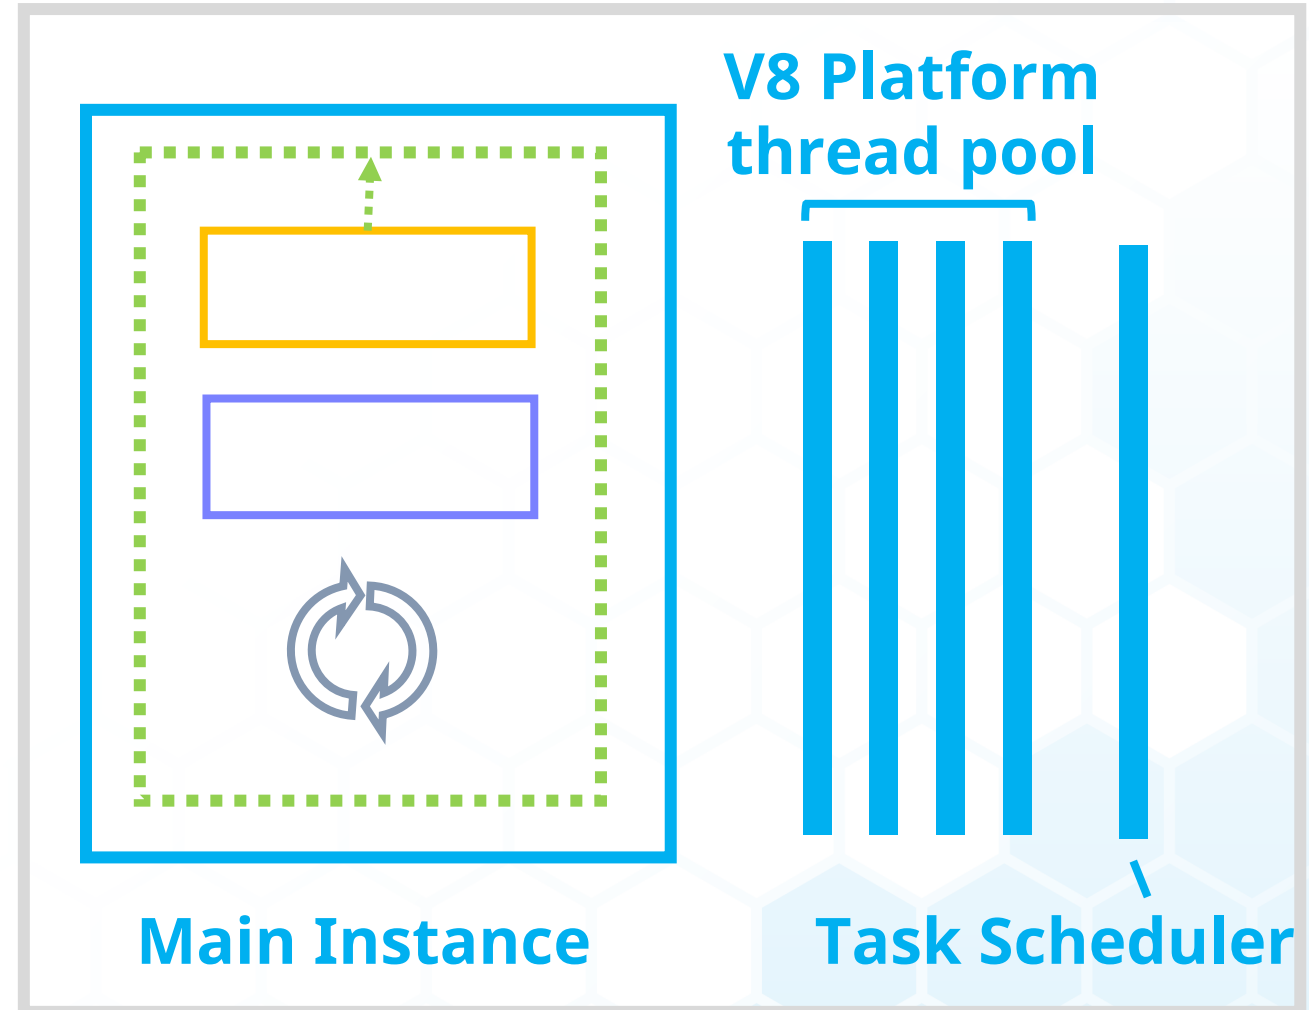
- Immutable copy of **primordials**
- **DOMException** for Web APIs (surprise!)



Setting up the V8 context

Main contexts & vm contexts

- Each Node.js instance has a main context (and potentially contexts created with vm).
- Main contexts contain a pointer to its associated Node.js **Environment** (not yet created)
- VM contexts do not have that pointer and are not bootstrapped further





NODE+JS INTERACTIVE

December 11-12, 2019
Montreal, Canada

V8 snapshot: before

Array **Object** **String**
...

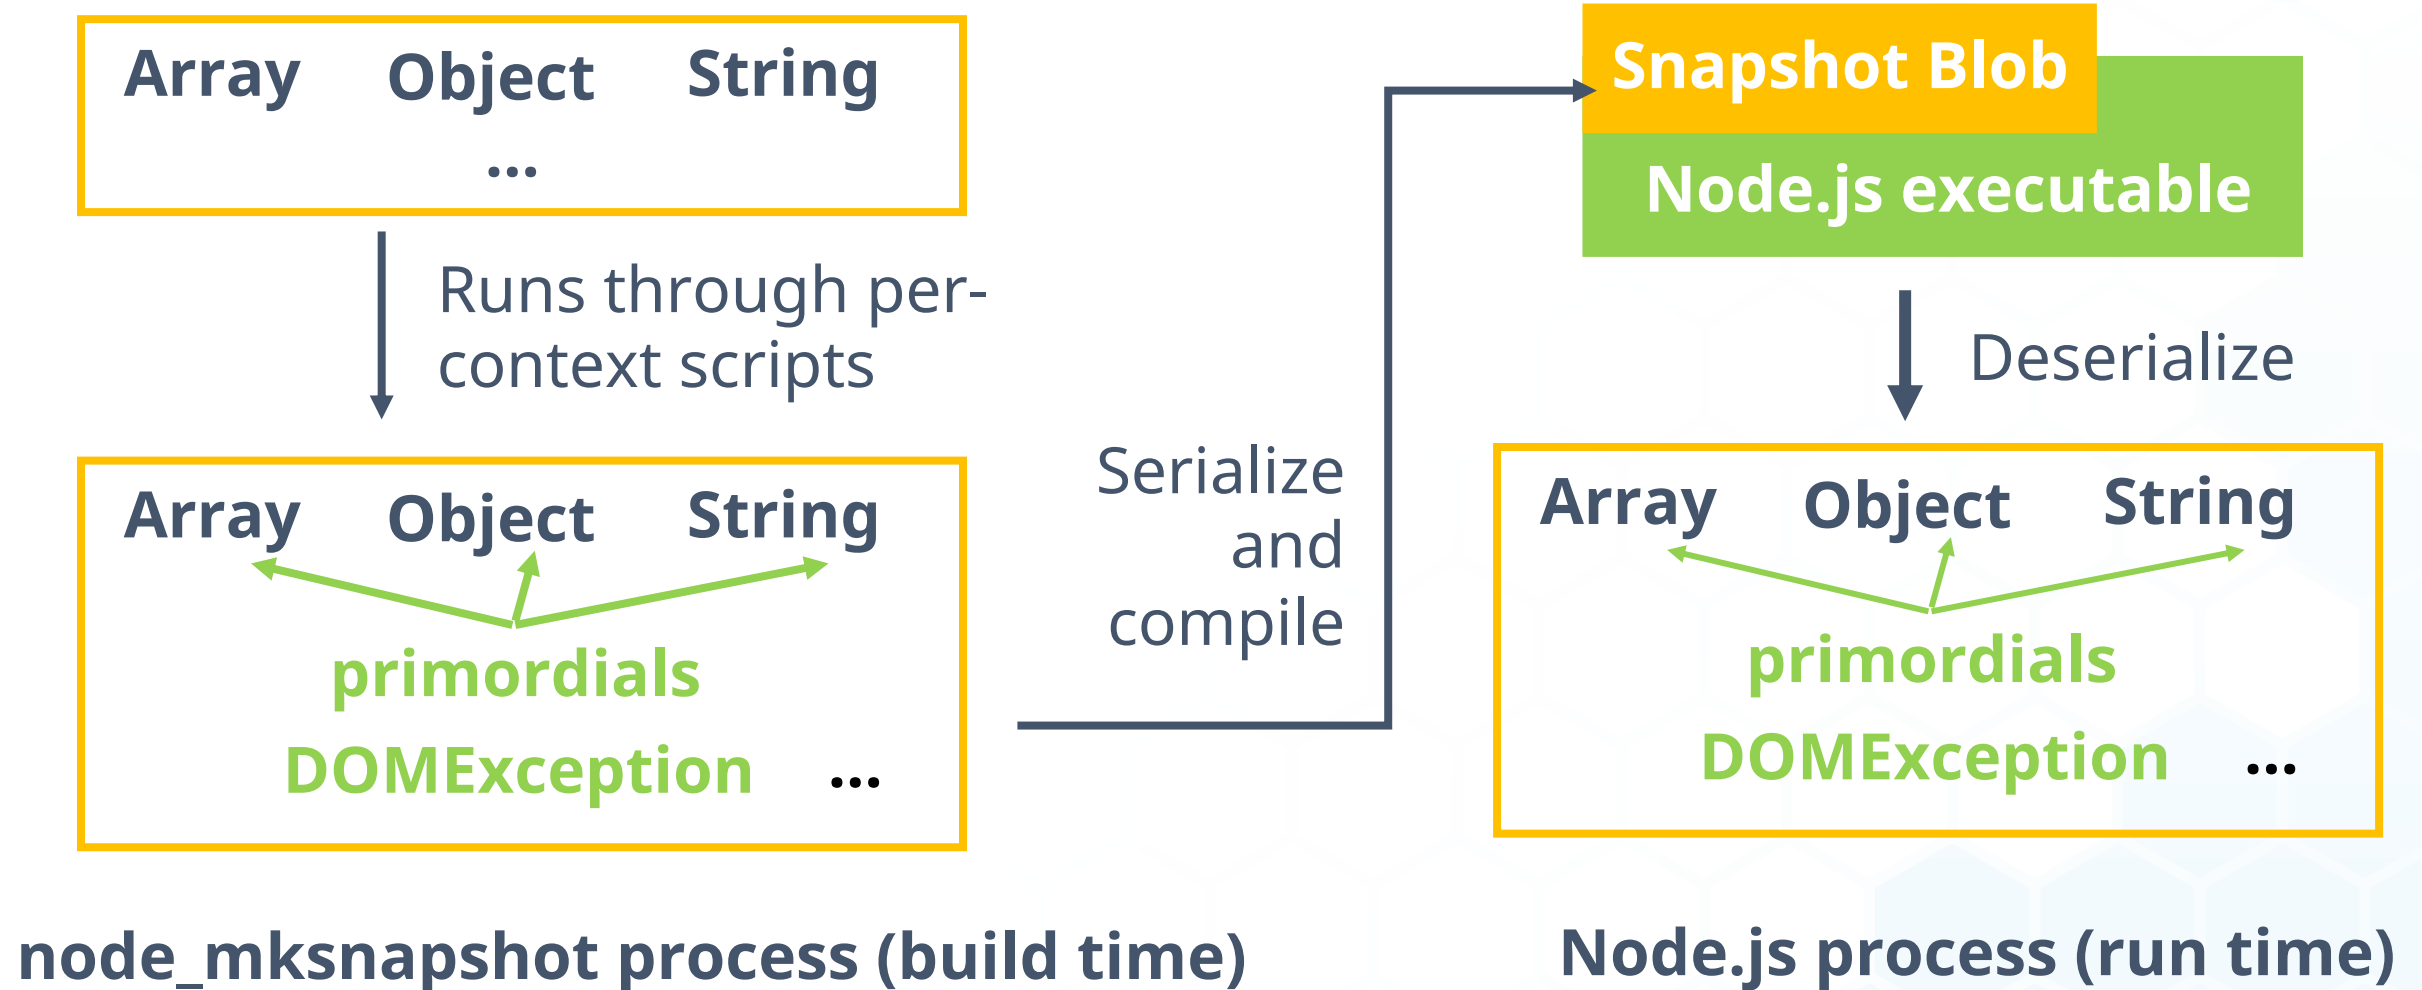


Runs through per-
context scripts

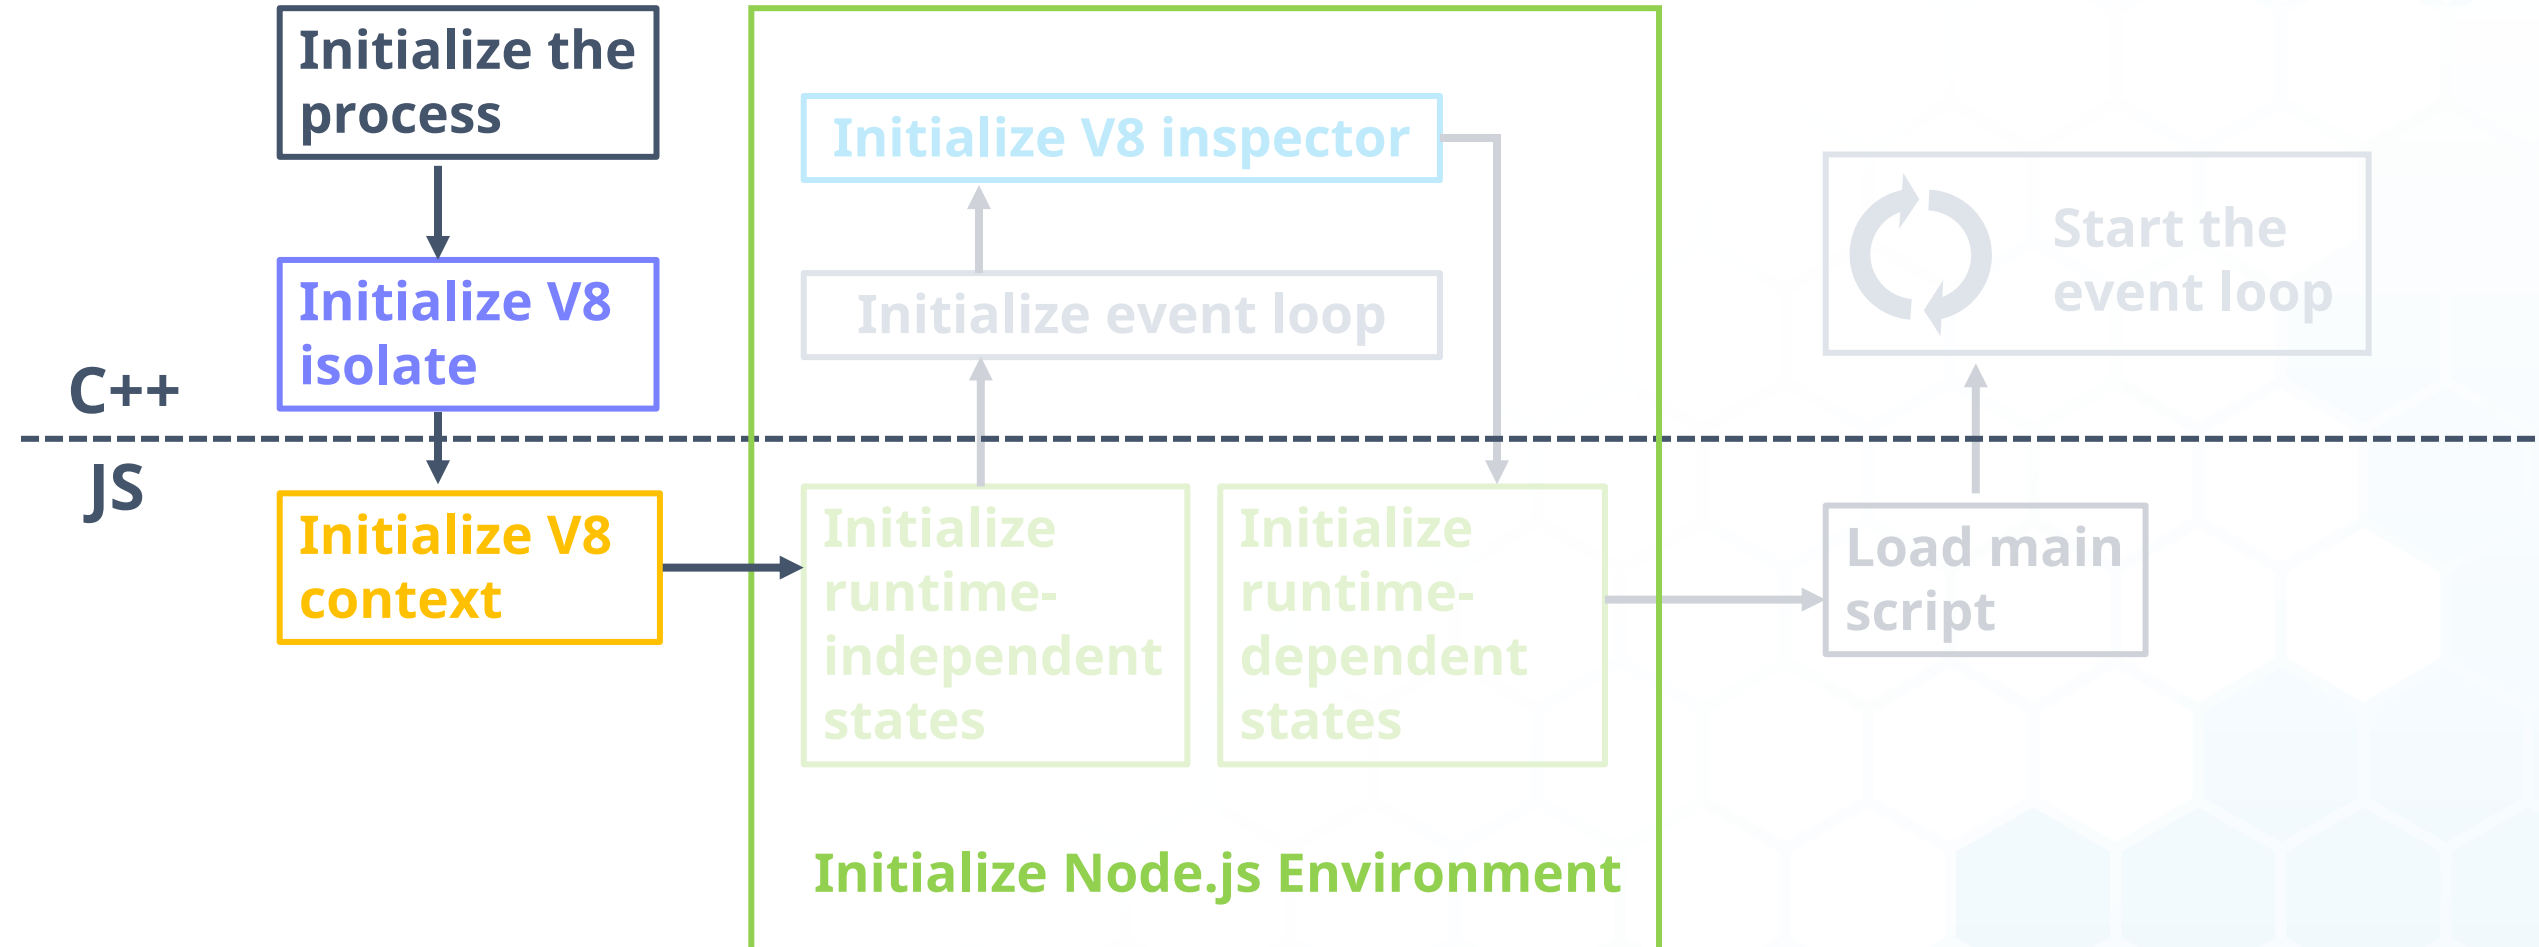
Array **Object** **String**
← ← →
primordials
DOMException ...

Node.js process

V8 snapshot: after



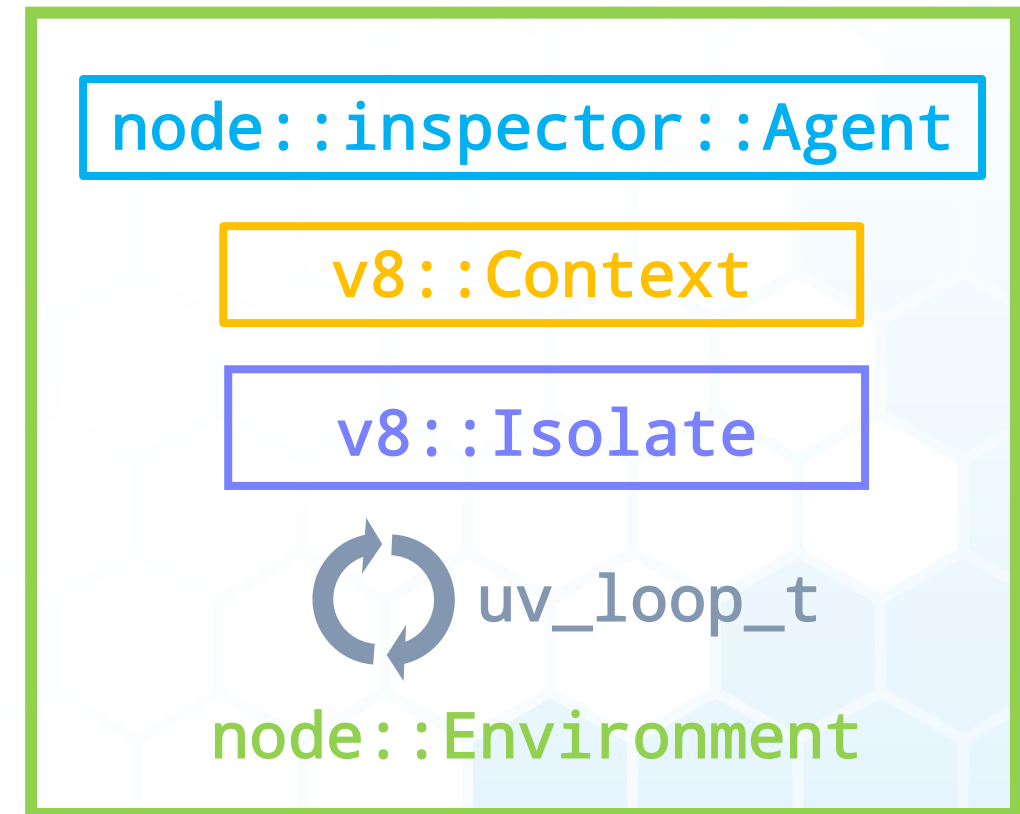
Setting up the Environment



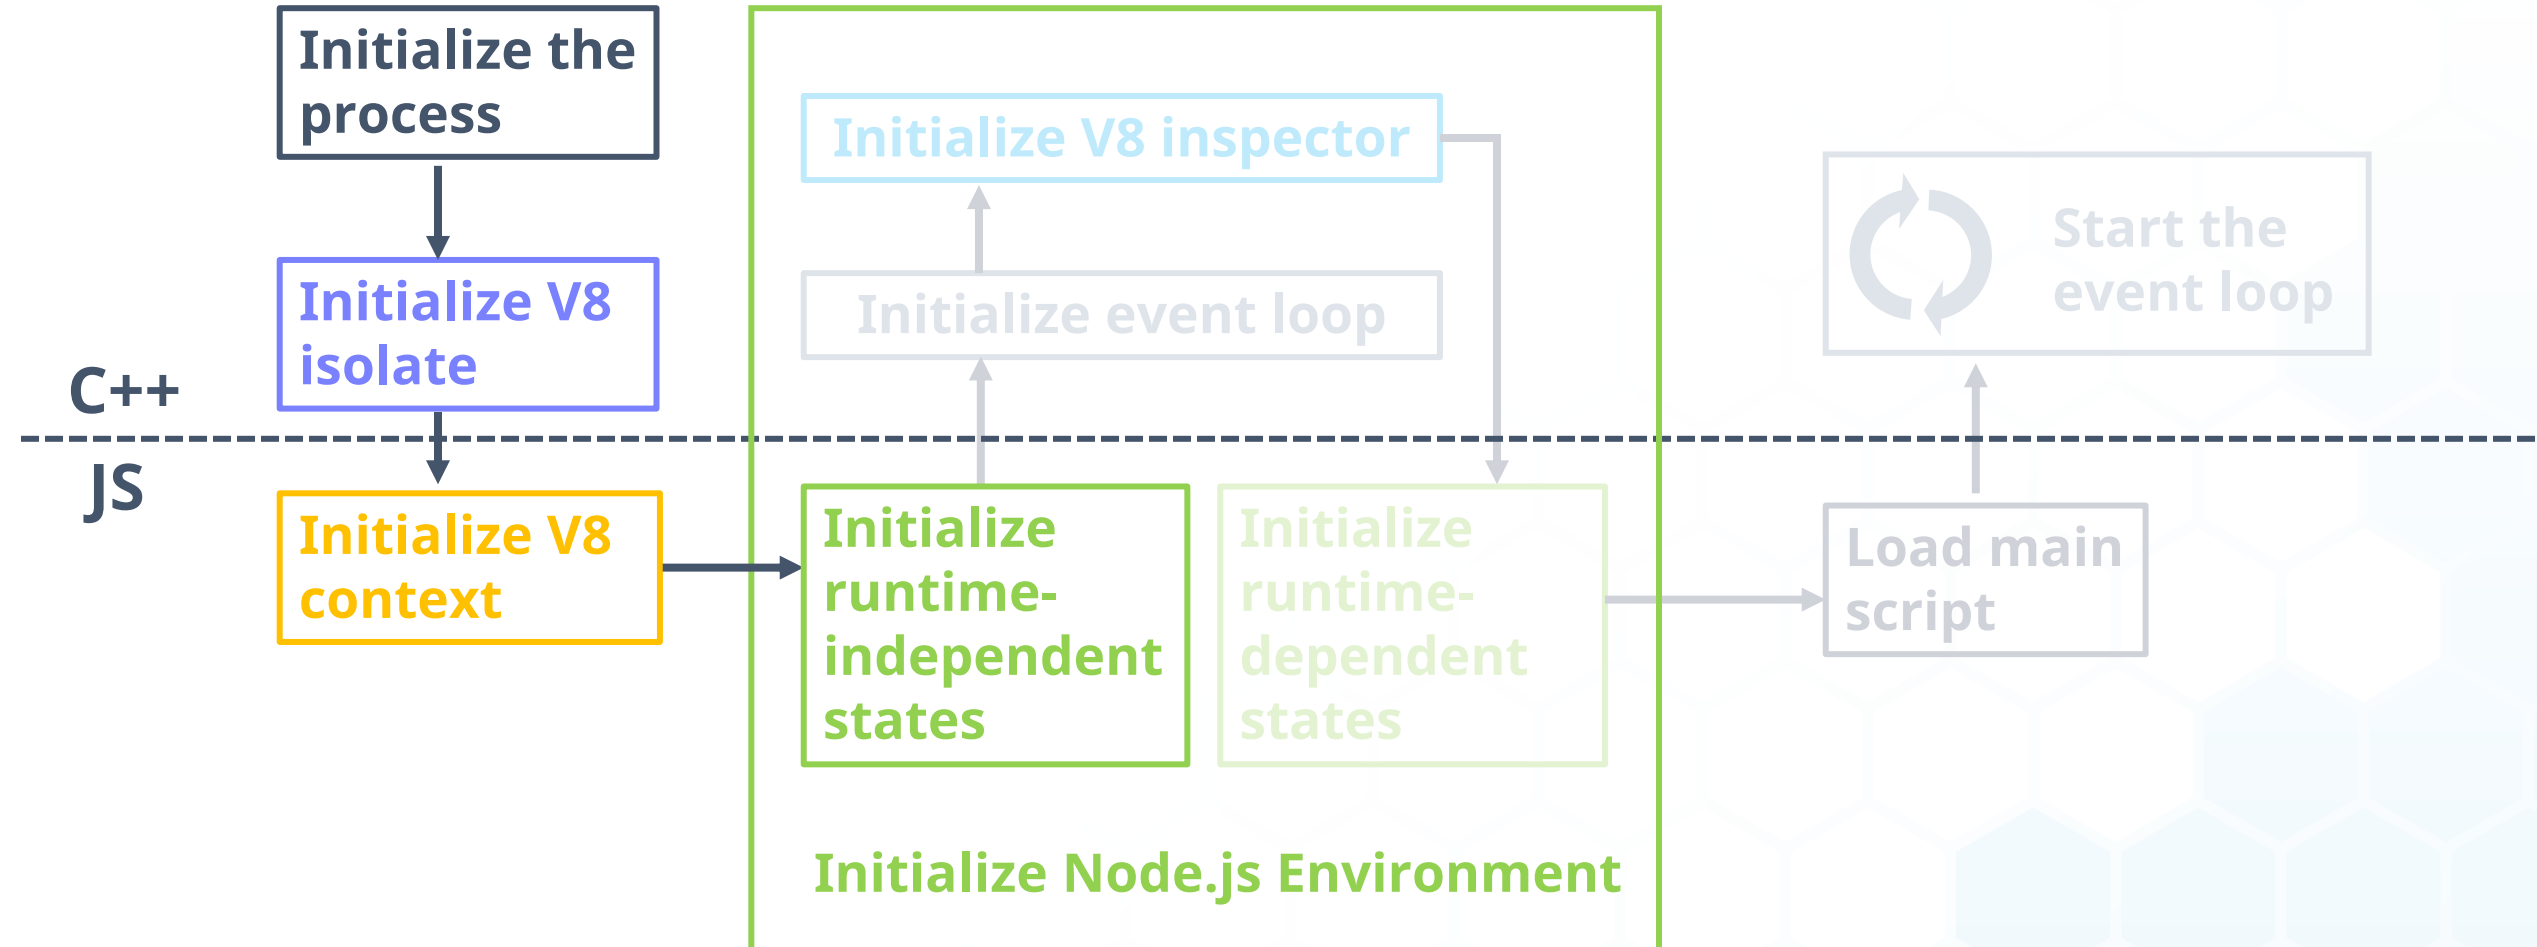
Setting up the Environment

What's a Node.js Environment?

- Encapsulation of the Node.js instance
- Associated with
 - One V8 inspector agent (for JS debugging)
 - One main V8 context
 - One V8 isolate
 - One libuv event loop



Setting up the Environment





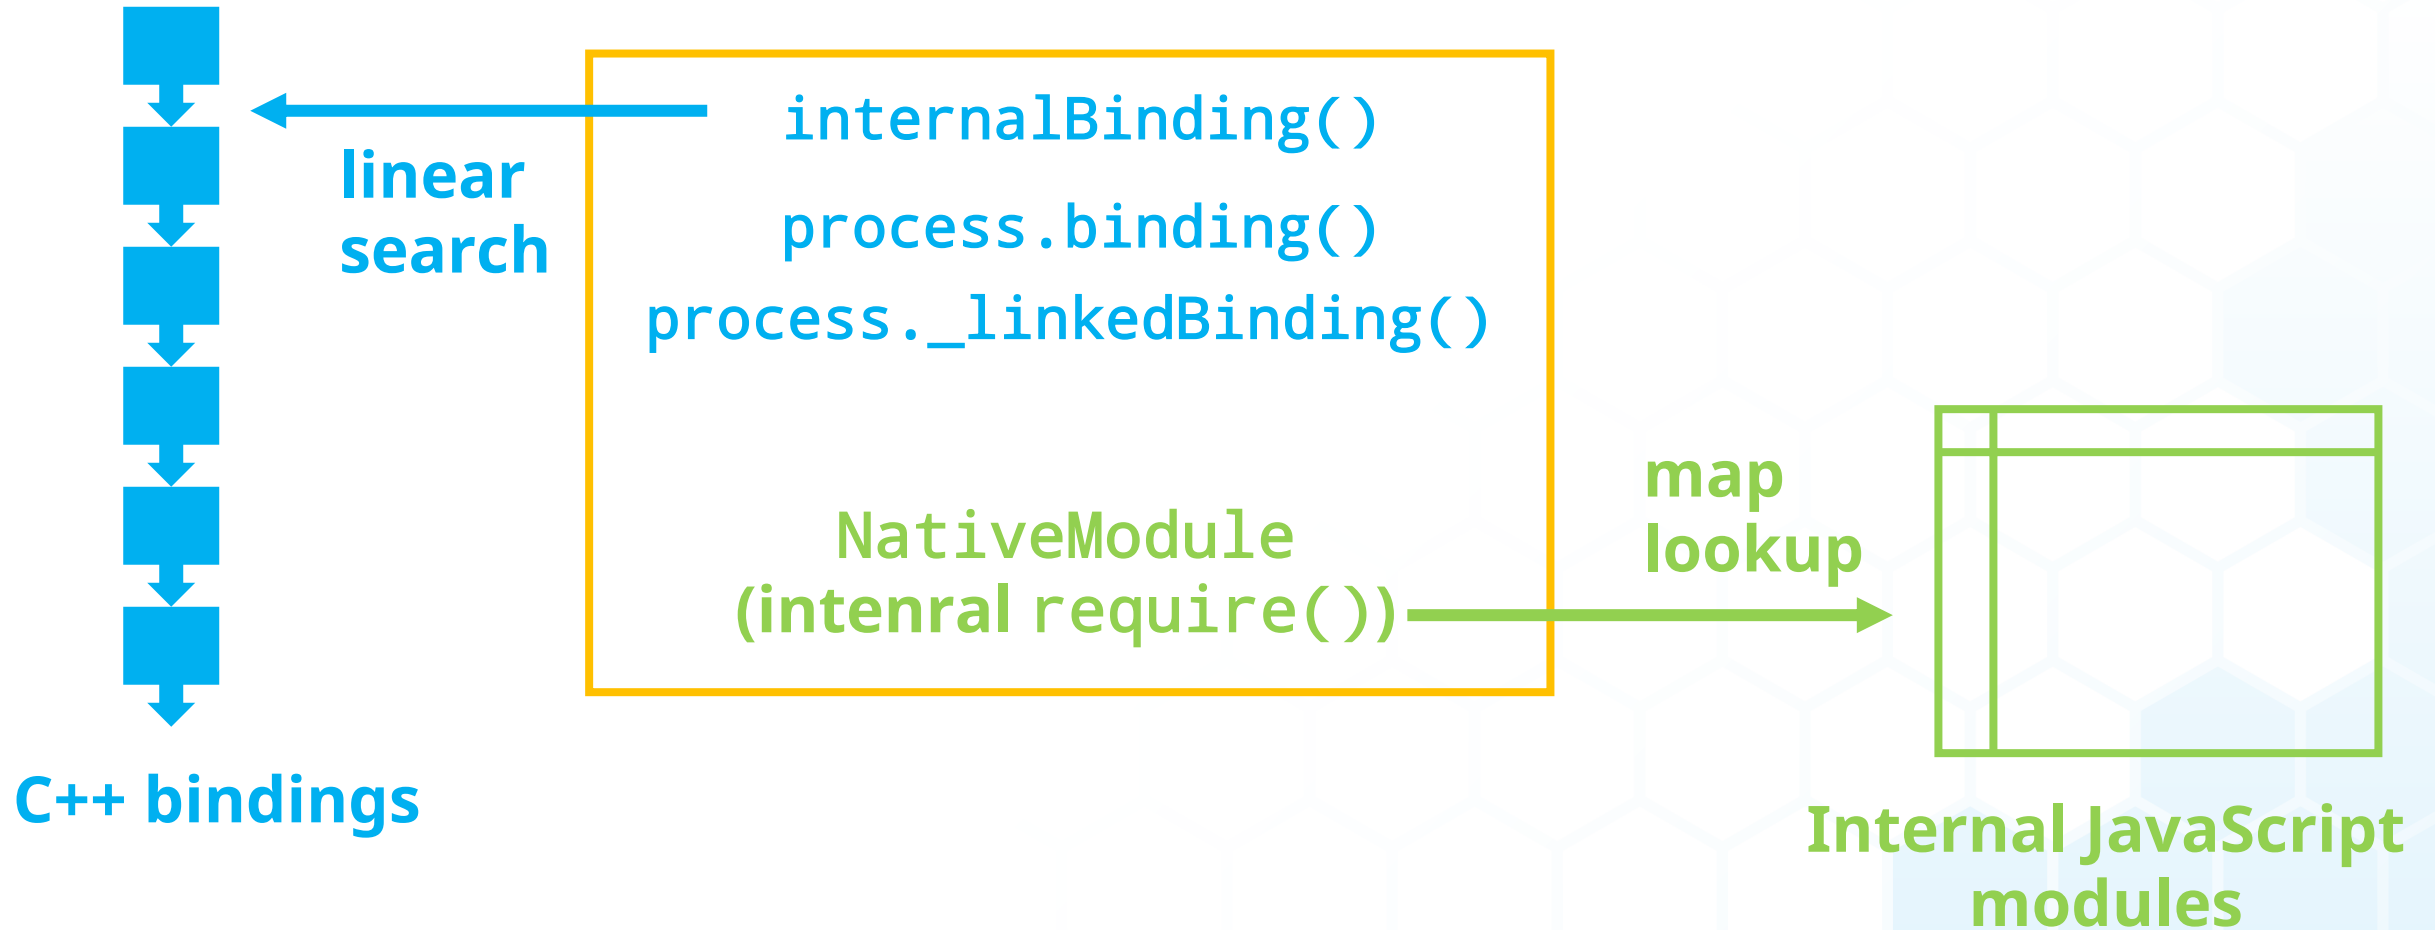
December 11-12, 2019
Montreal, Canada

Setting up the Environment


What needs to be initialized?

- **Components independent of runtime states**
 - Internal JavaScript module and C++ binding loaders
 - The process object and other globals
 - JavaScript callbacks that C++ hooks invoke

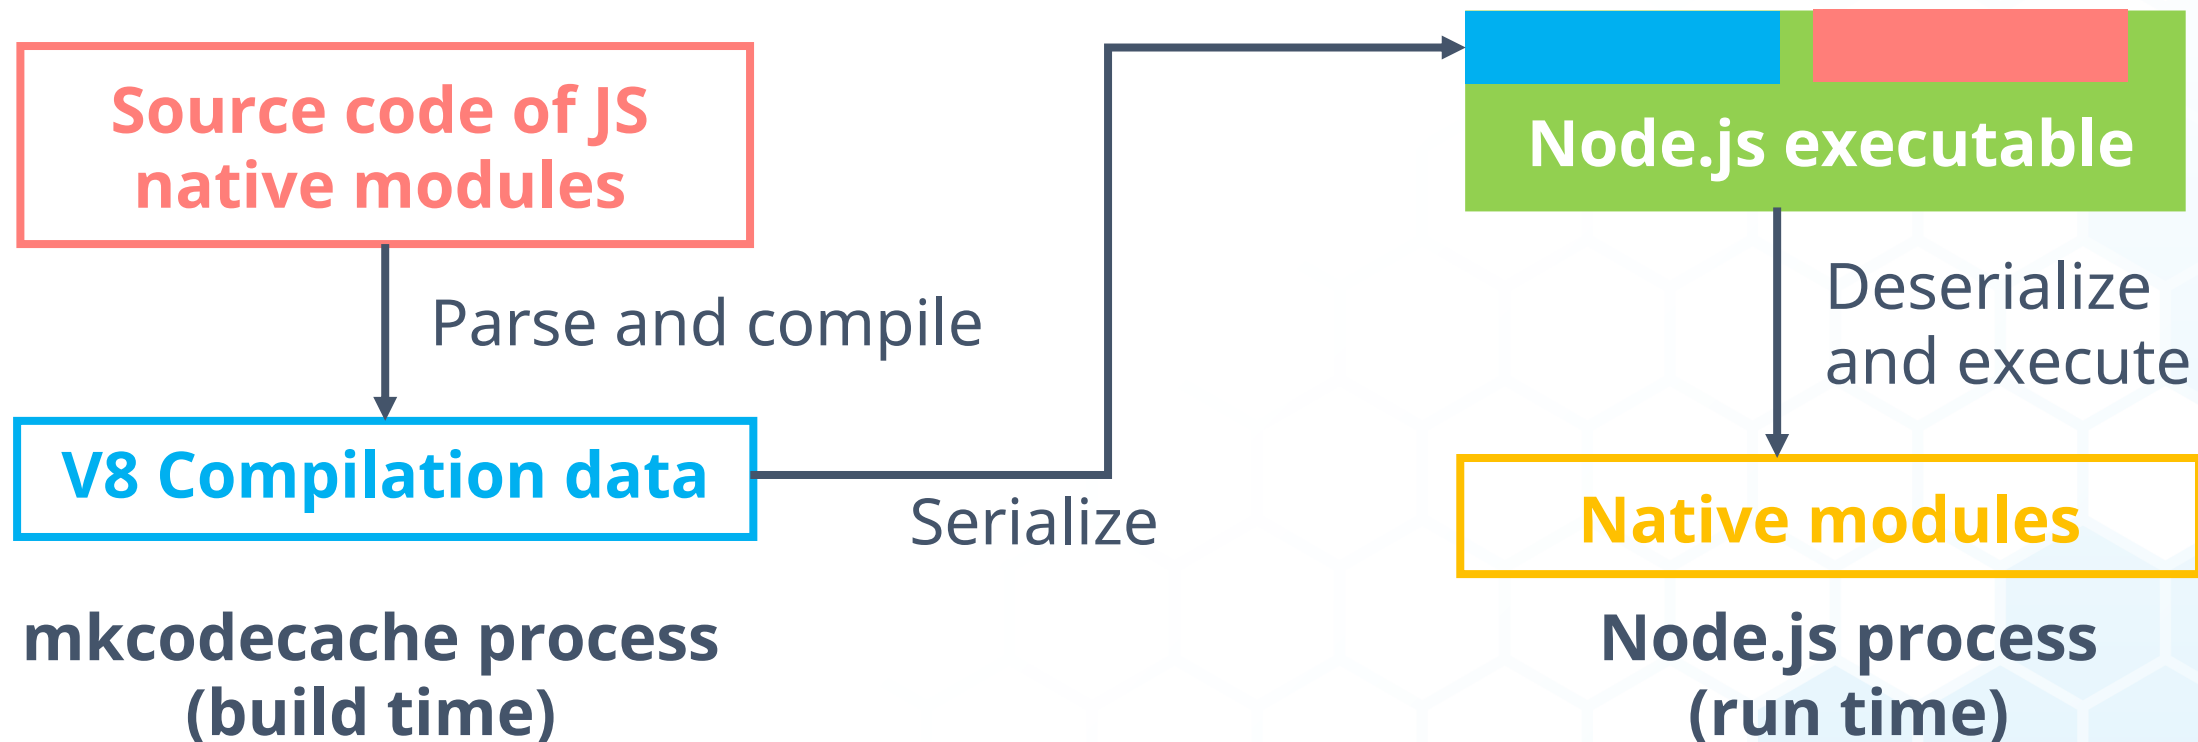
Environment: internal loaders



V8 code cache: before



V8 code cache: after





December 11-12, 2019
Montreal, Canada

Environment: globals

- Most globals are implemented in internal JavaScript with access to internal C++ bindings (glue to dependencies)
- Created and attached to `global` or `process` during bootstrap
- `global` is a legacy alias to the ECMAScript stage 4 `globalThis`



December 11-12, 2019
Montreal, Canada

Environment: globals

```
process.nextTick =  
  require('internal/process/task_queues').setupTaskQueue().nextTick;  
...  
  
Object.defineProperty(global, 'process', { value: process, ...});  
Object.defineProperty(global, 'global', { value: global, ...});  
Object.defineProperty(global, 'setTimeout', {  
  value: require('timers').setTimeout,  
  ...  
});
```



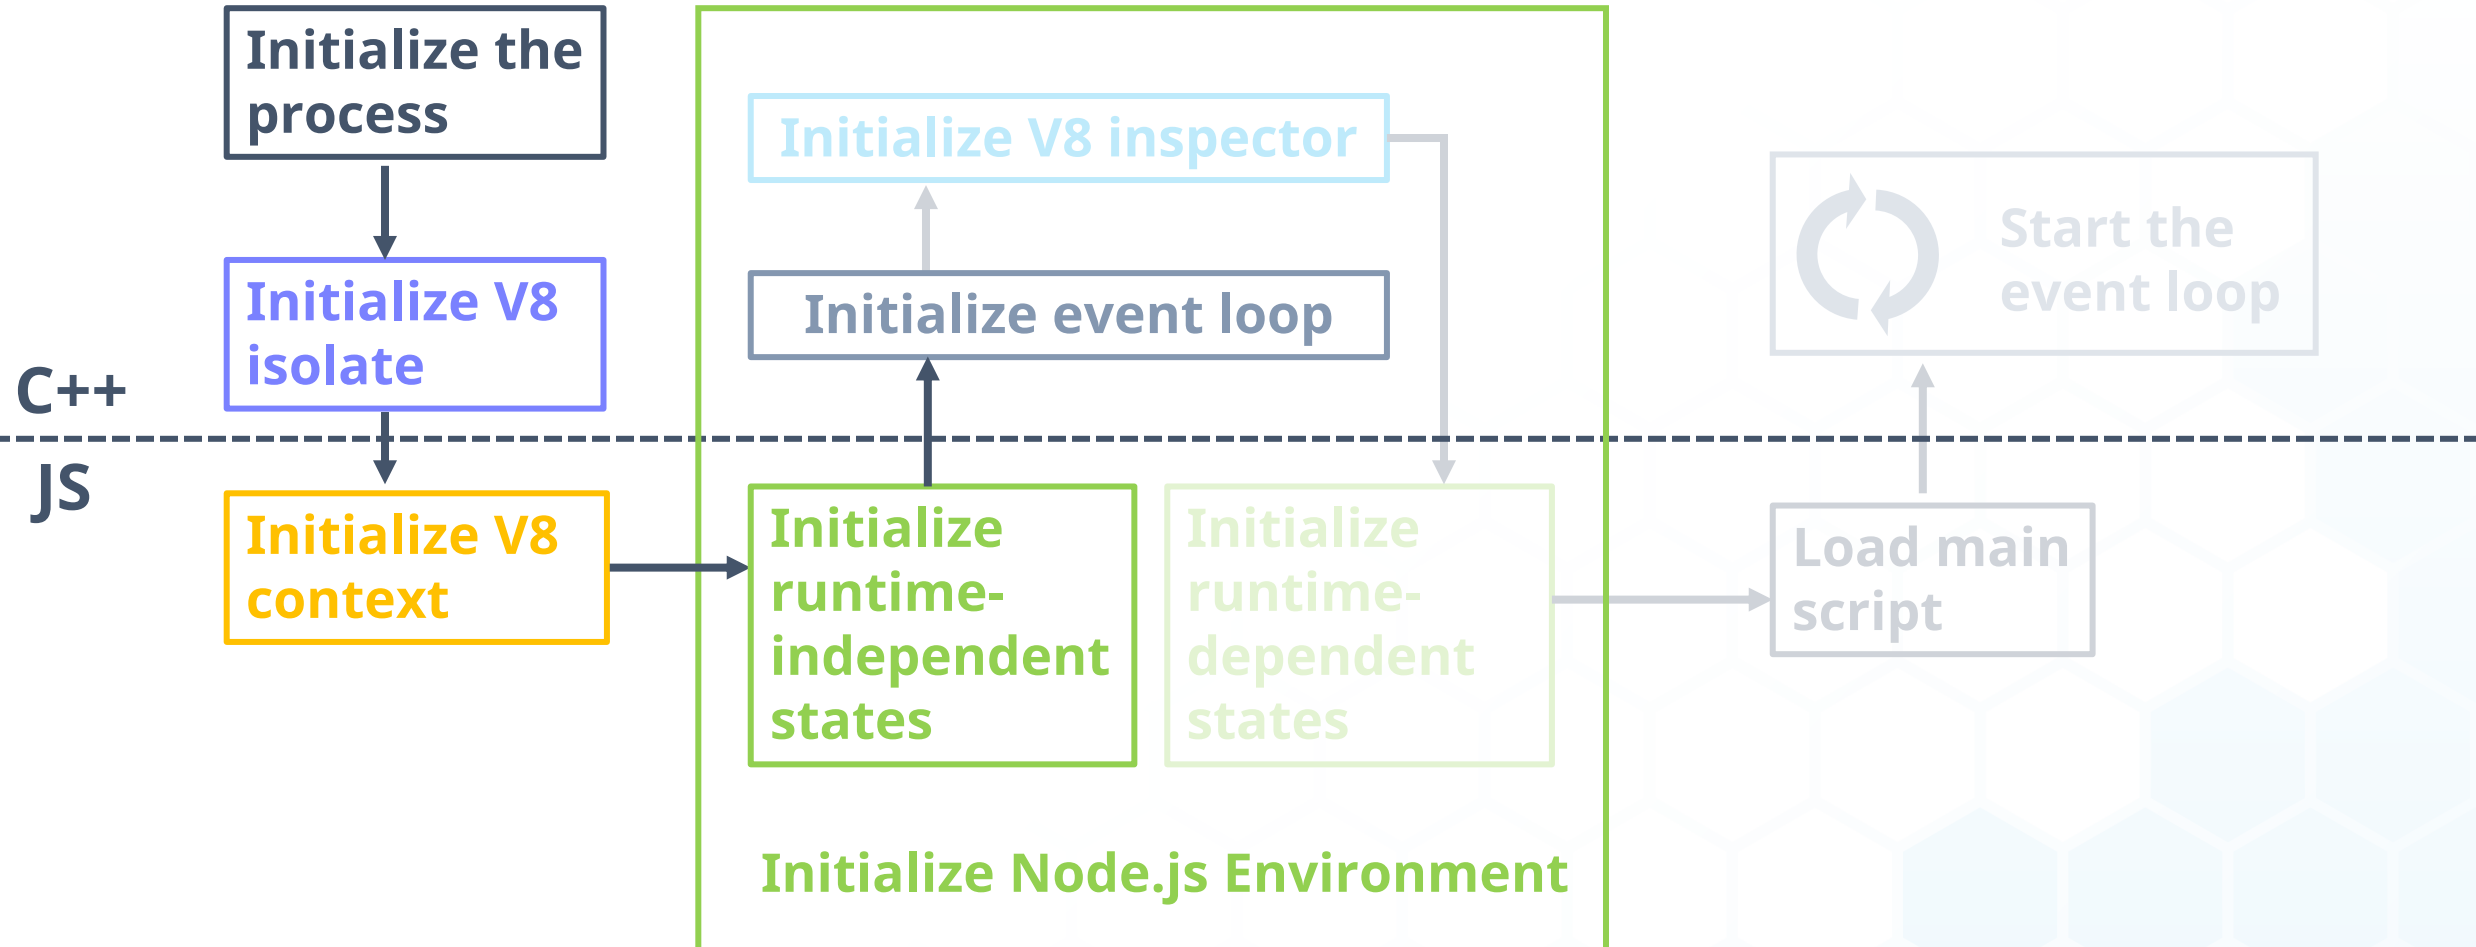
December 11-12, 2019
Montreal, Canada

Environment: initialize hooks

Example of hooks


- **Async hooks (per-Environment)**
 - User callbacks need to be invoked at different stages of async operations
- **`process.nextTick` (per-Environment)**
 - Queued user callbacks need to be invoked when async operations are done
- **`Error.prepareStackTrace` (per-Isolate)**
 - User hook needs to be invoked when `error.stack` is accessed
- **`process.on('uncaughtException')` / `process.on('unhandledRejection')` (per-Isolate)**
 - Need to be invoked when there are uncaught exceptions/unhandled rejections

Setting up the Environment

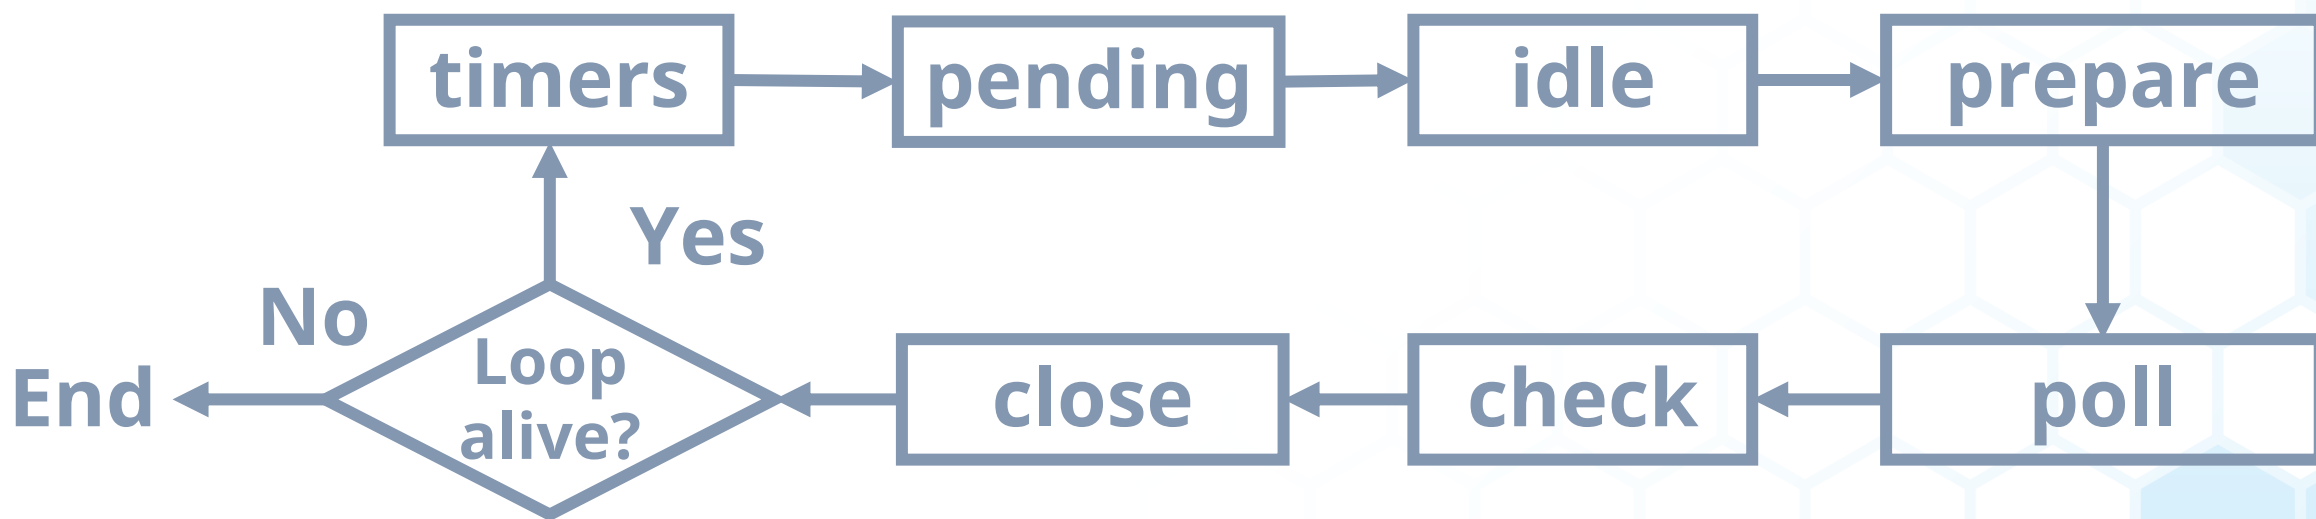


Environment: initialize the event loop

- Some handles are initialized at bootstrap
- Some are activated immediately, some are activated on demand
- More handles (e.g. poll handles) can be added on-demand (e.g. for I/O)

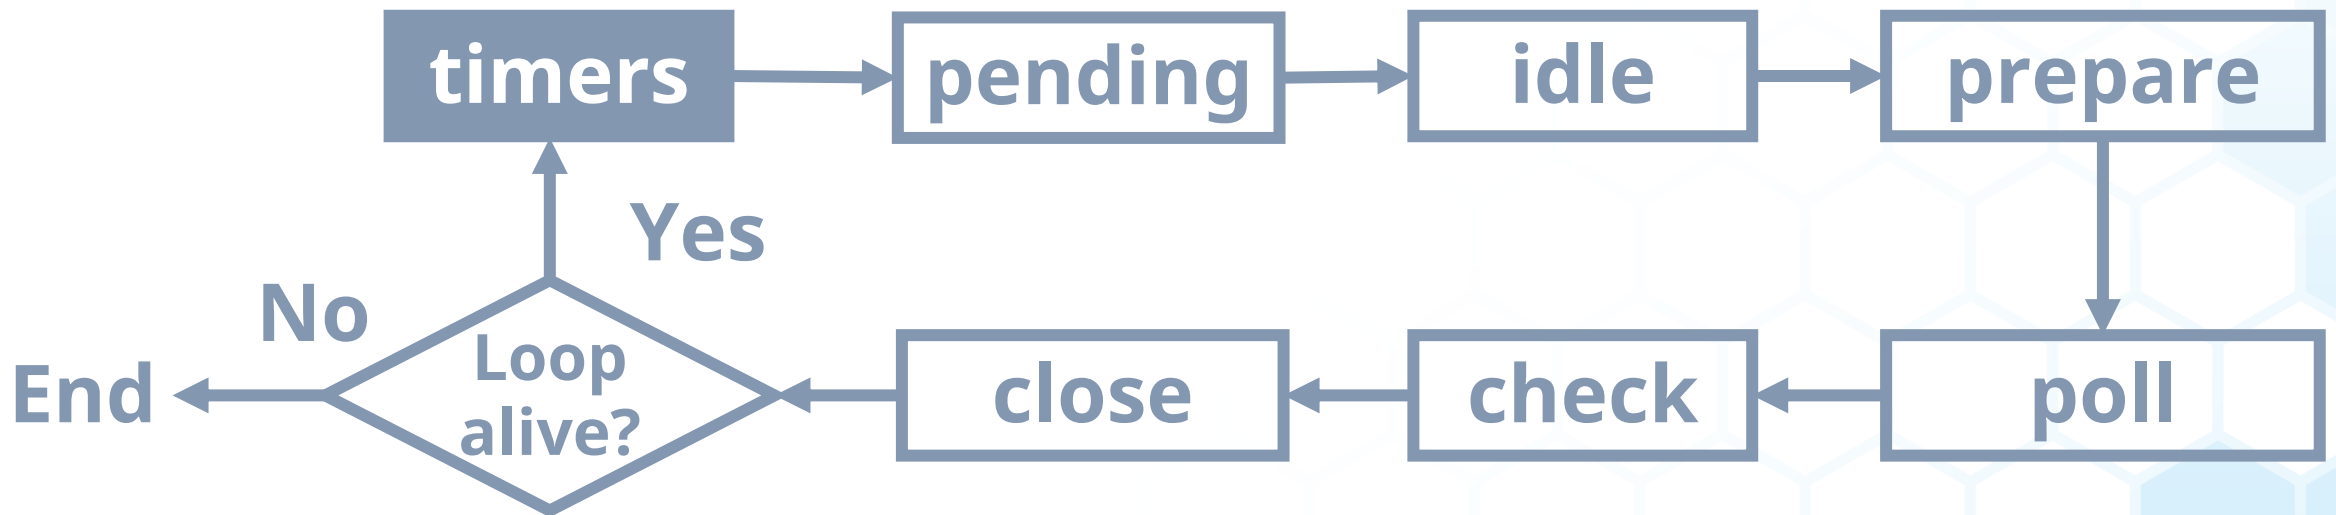


Environment: initialize the event loop

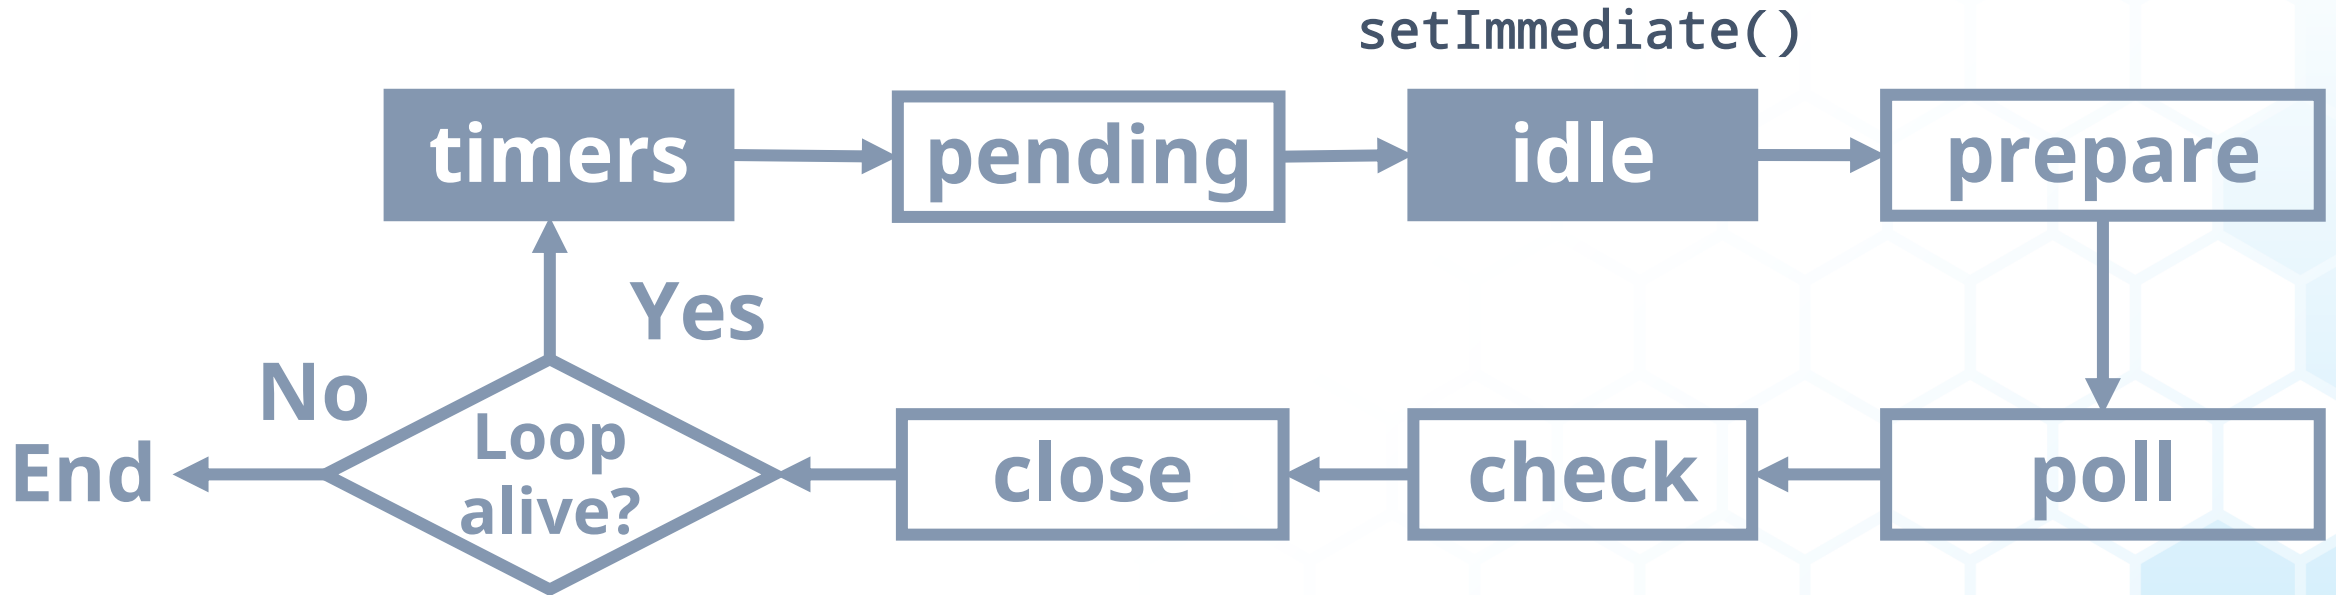


Environment: initialize the event loop

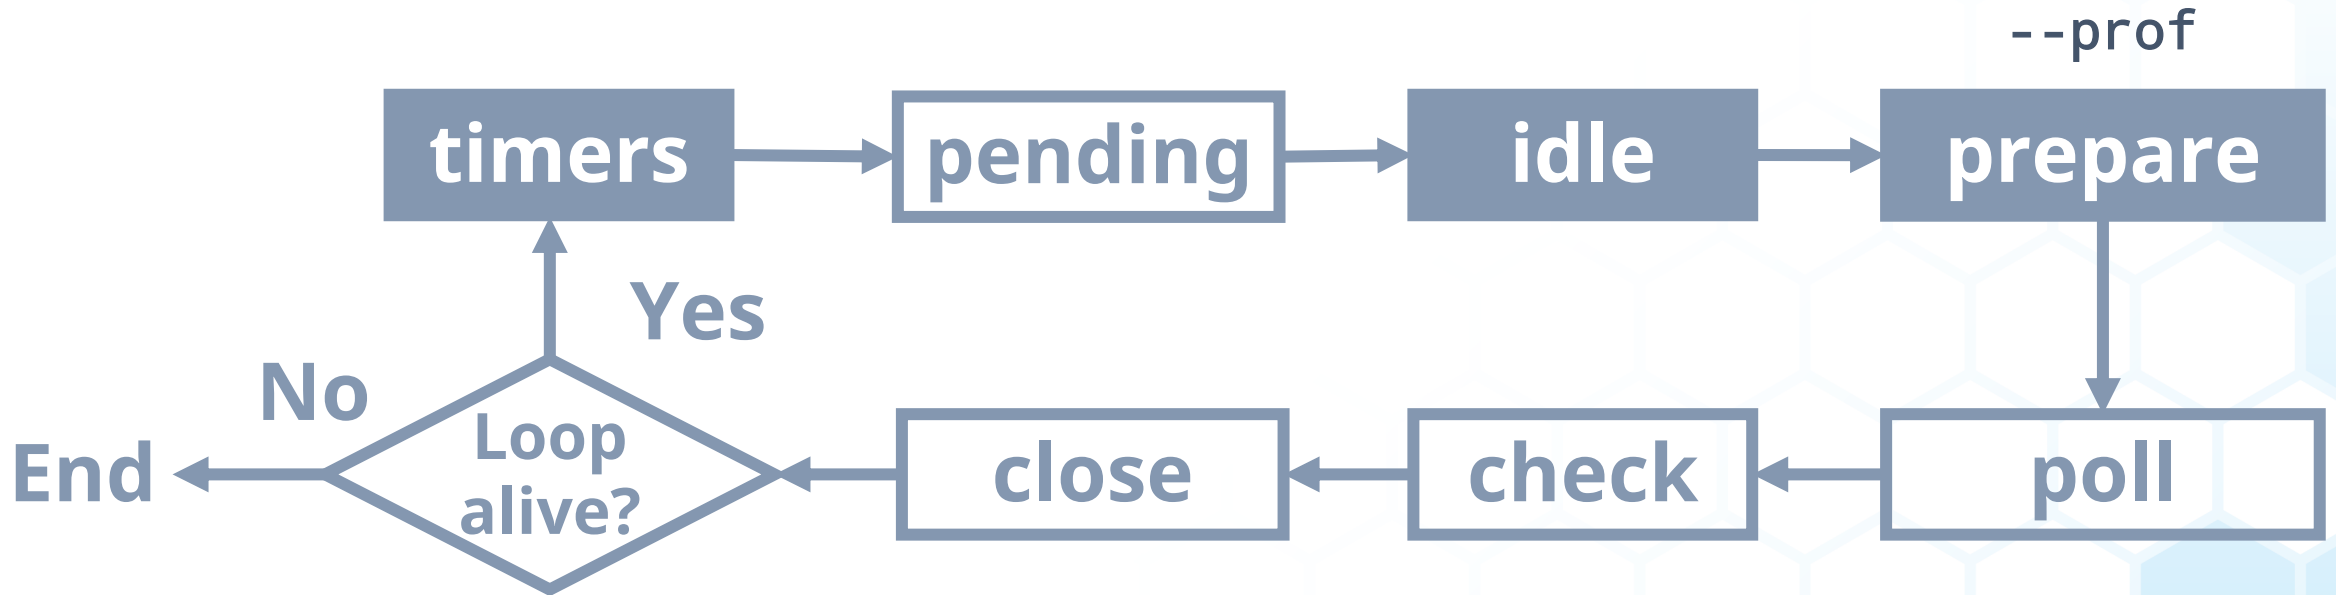
`setTimeout()` / `setInterval()`



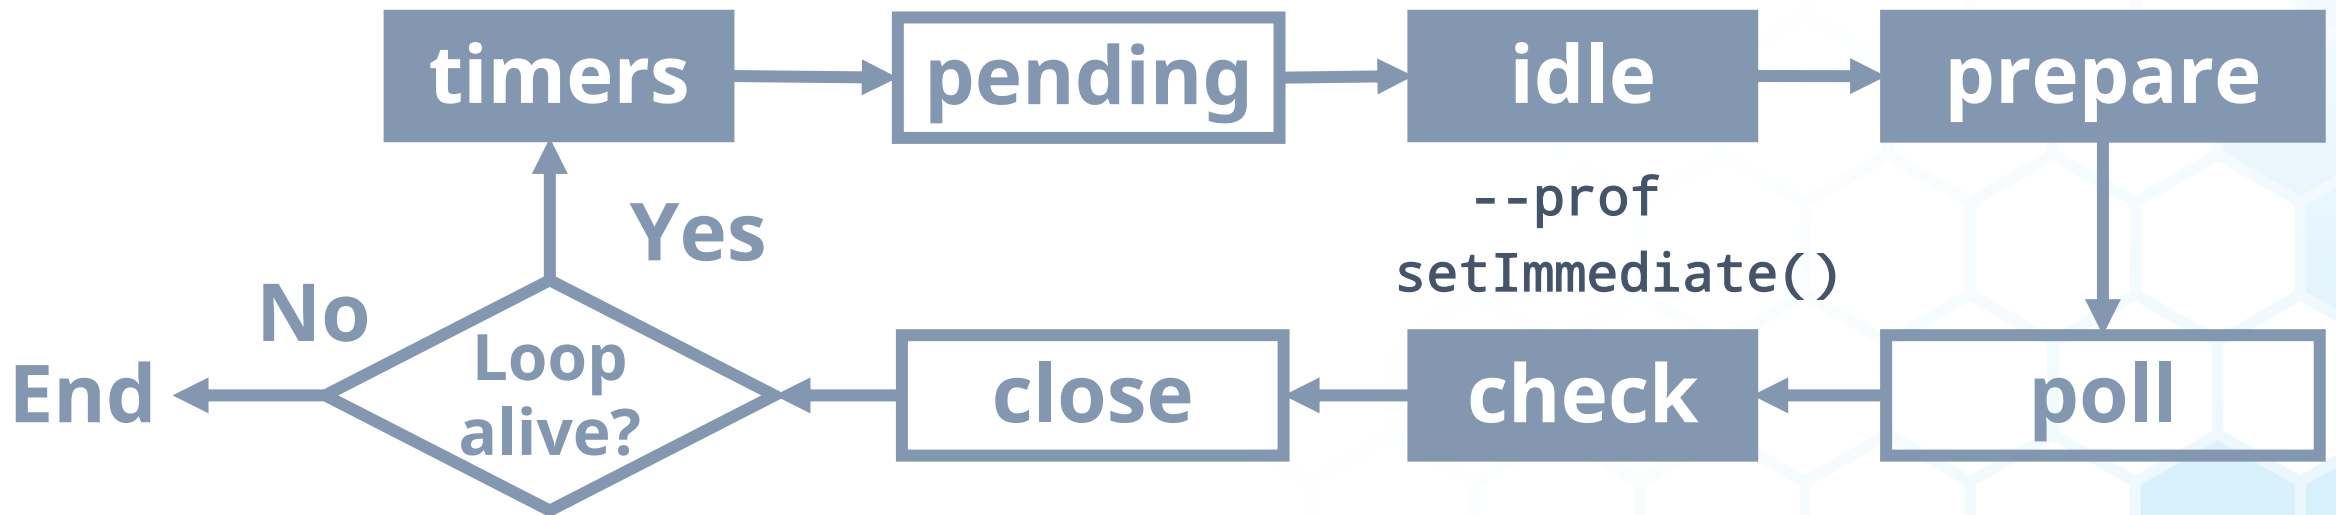
Environment: initialize the event loop



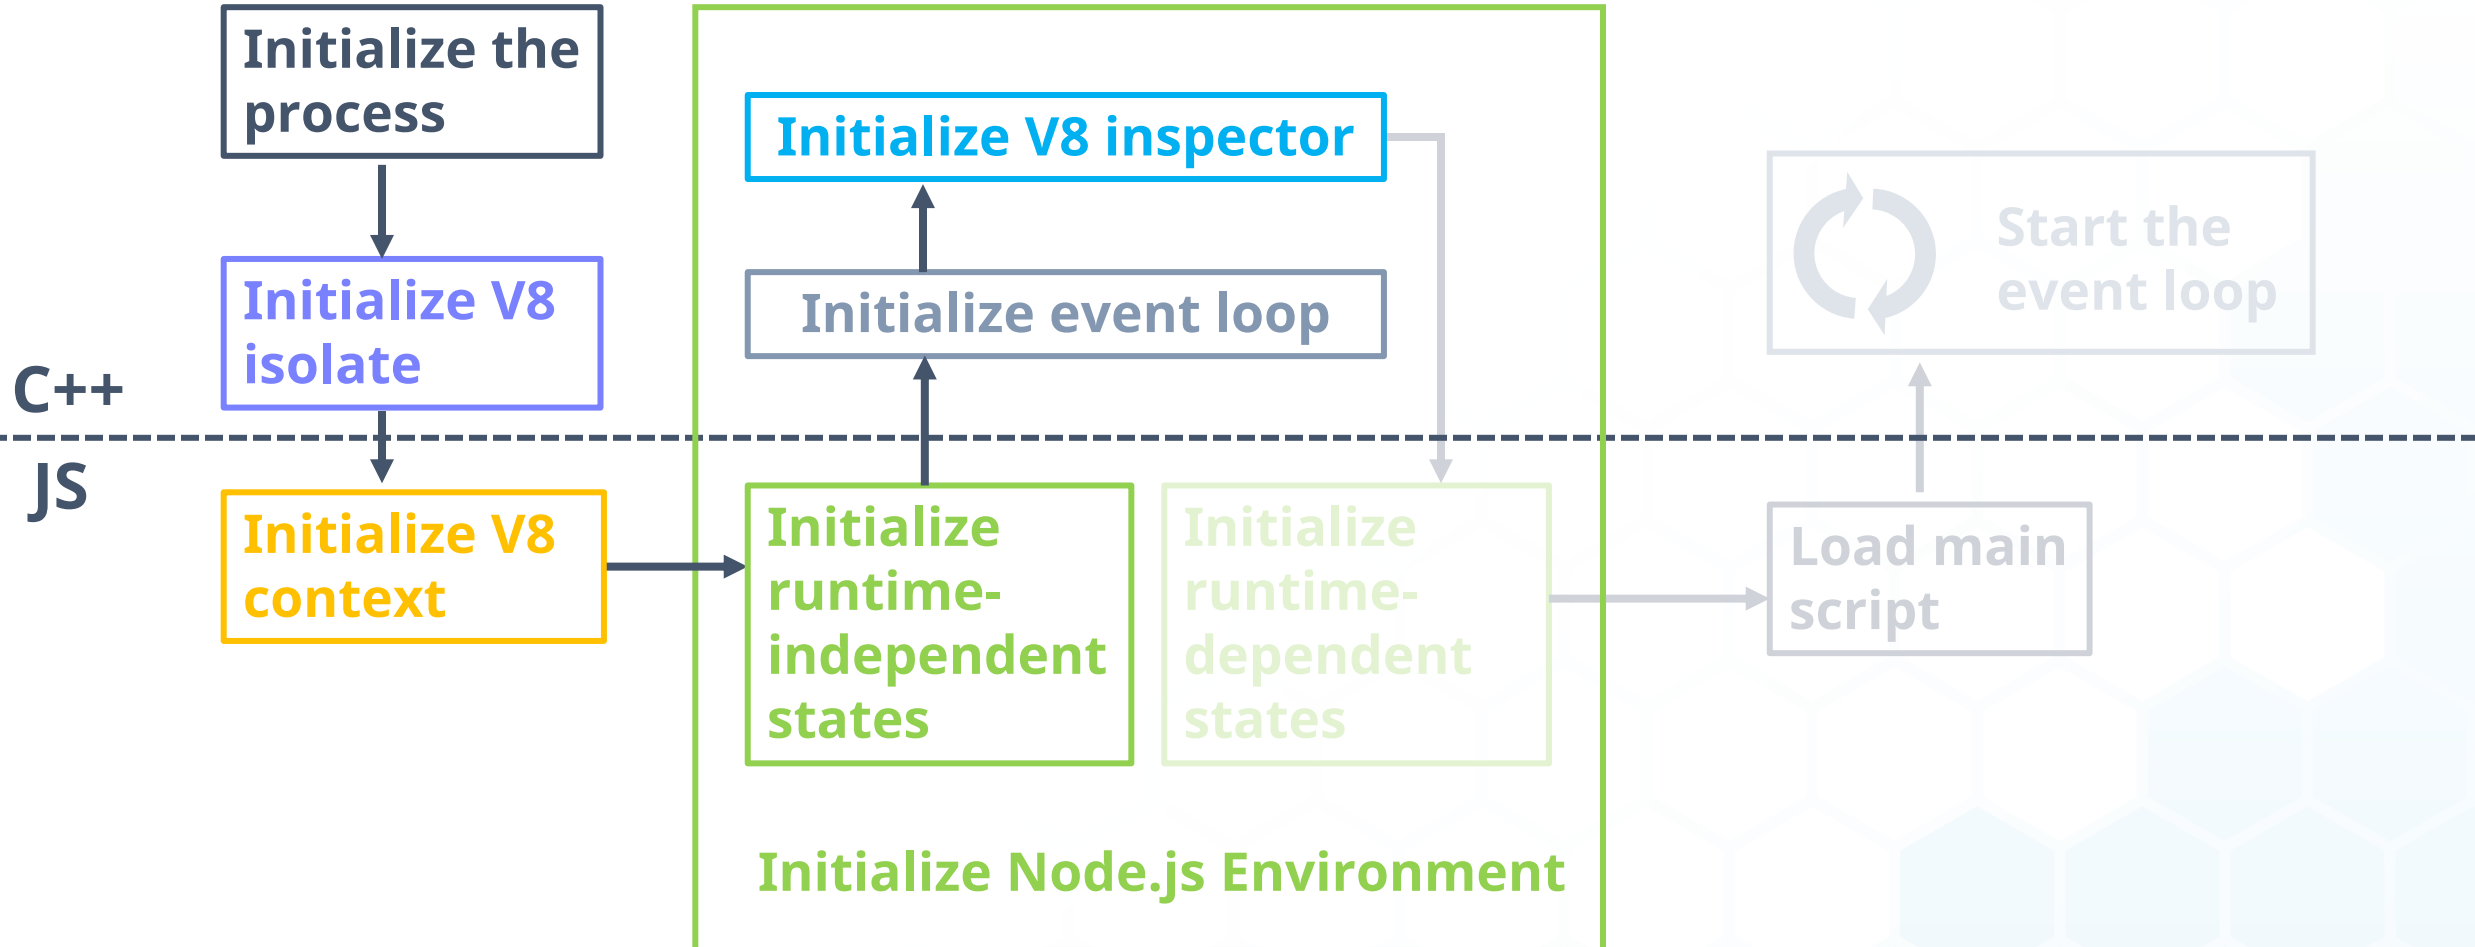
Environment: initialize the event loop



Environment: initialize the event loop

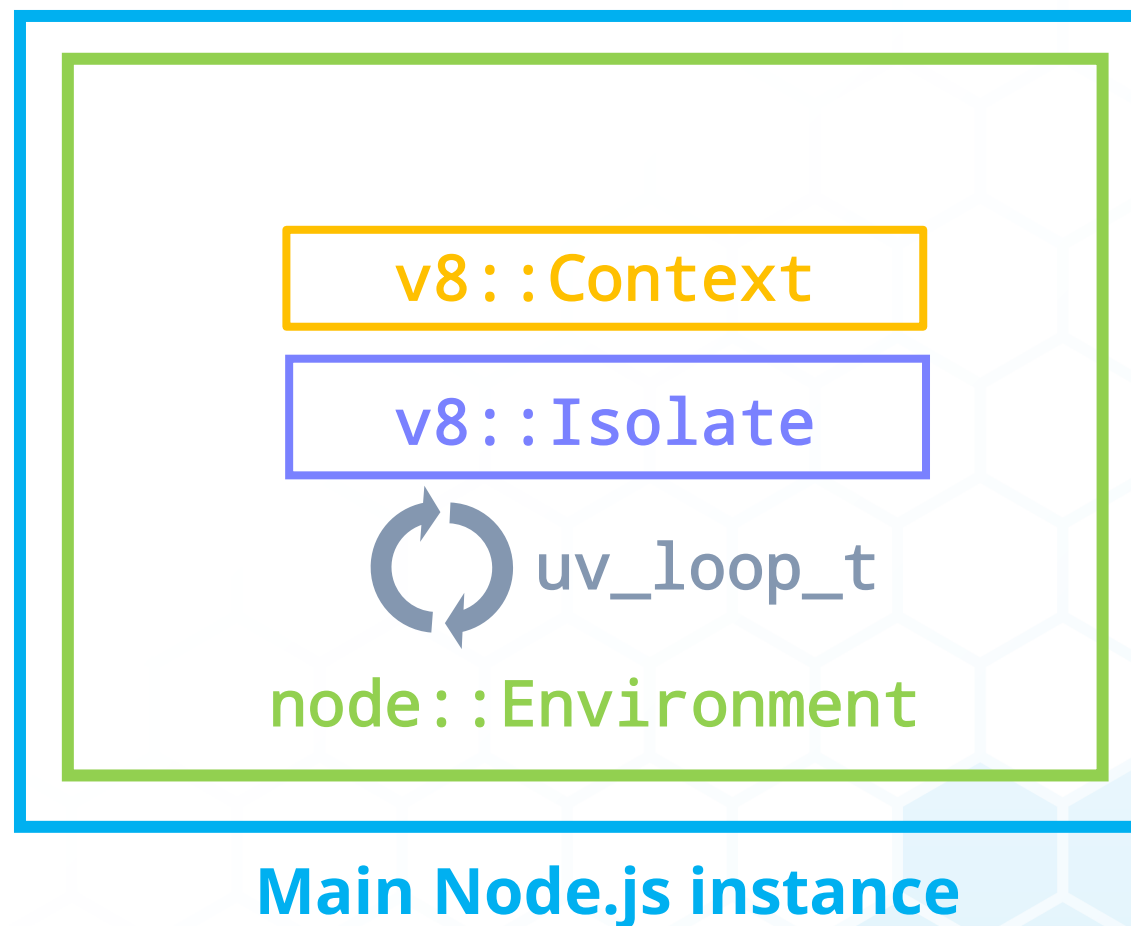


Setting up the Environment



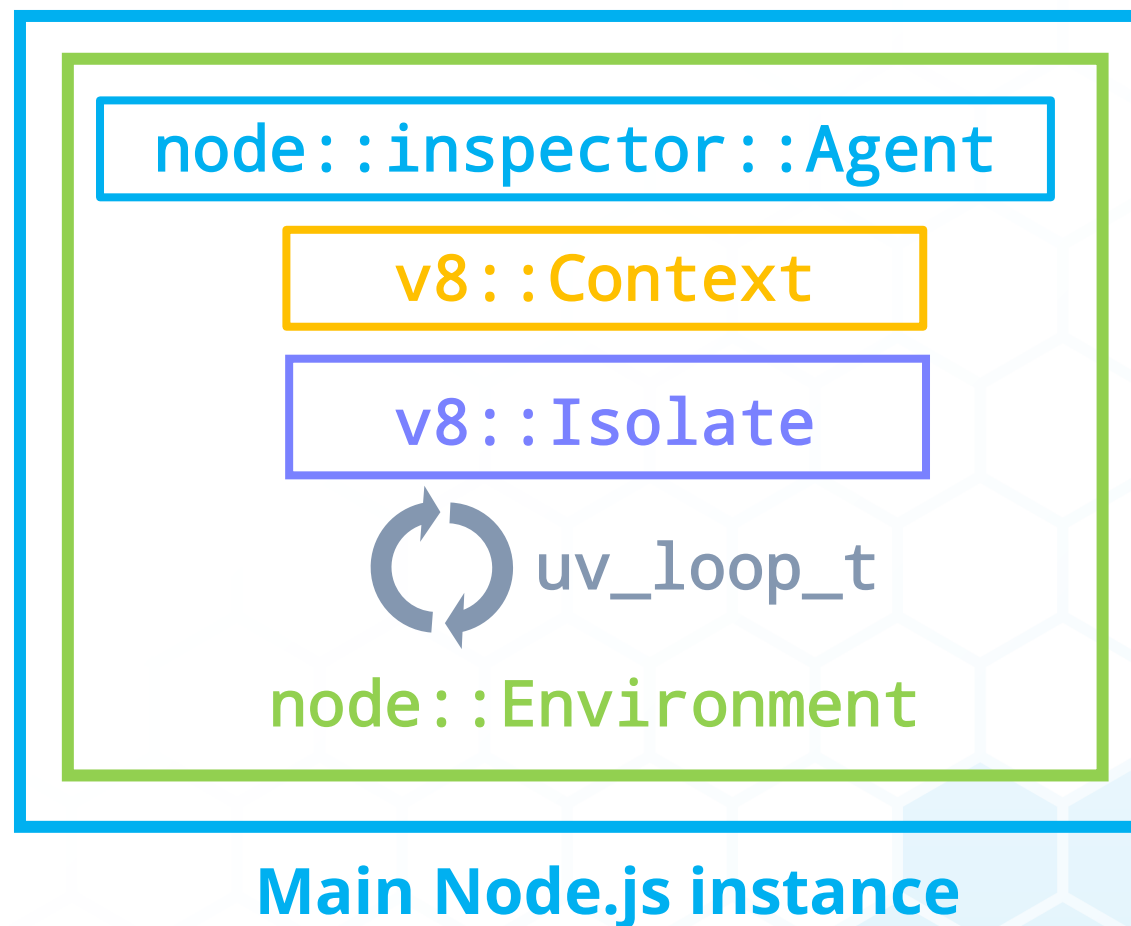
Environment: inspector

- Initialize inspector agent



Environment: inspector

- Initialize inspector agent



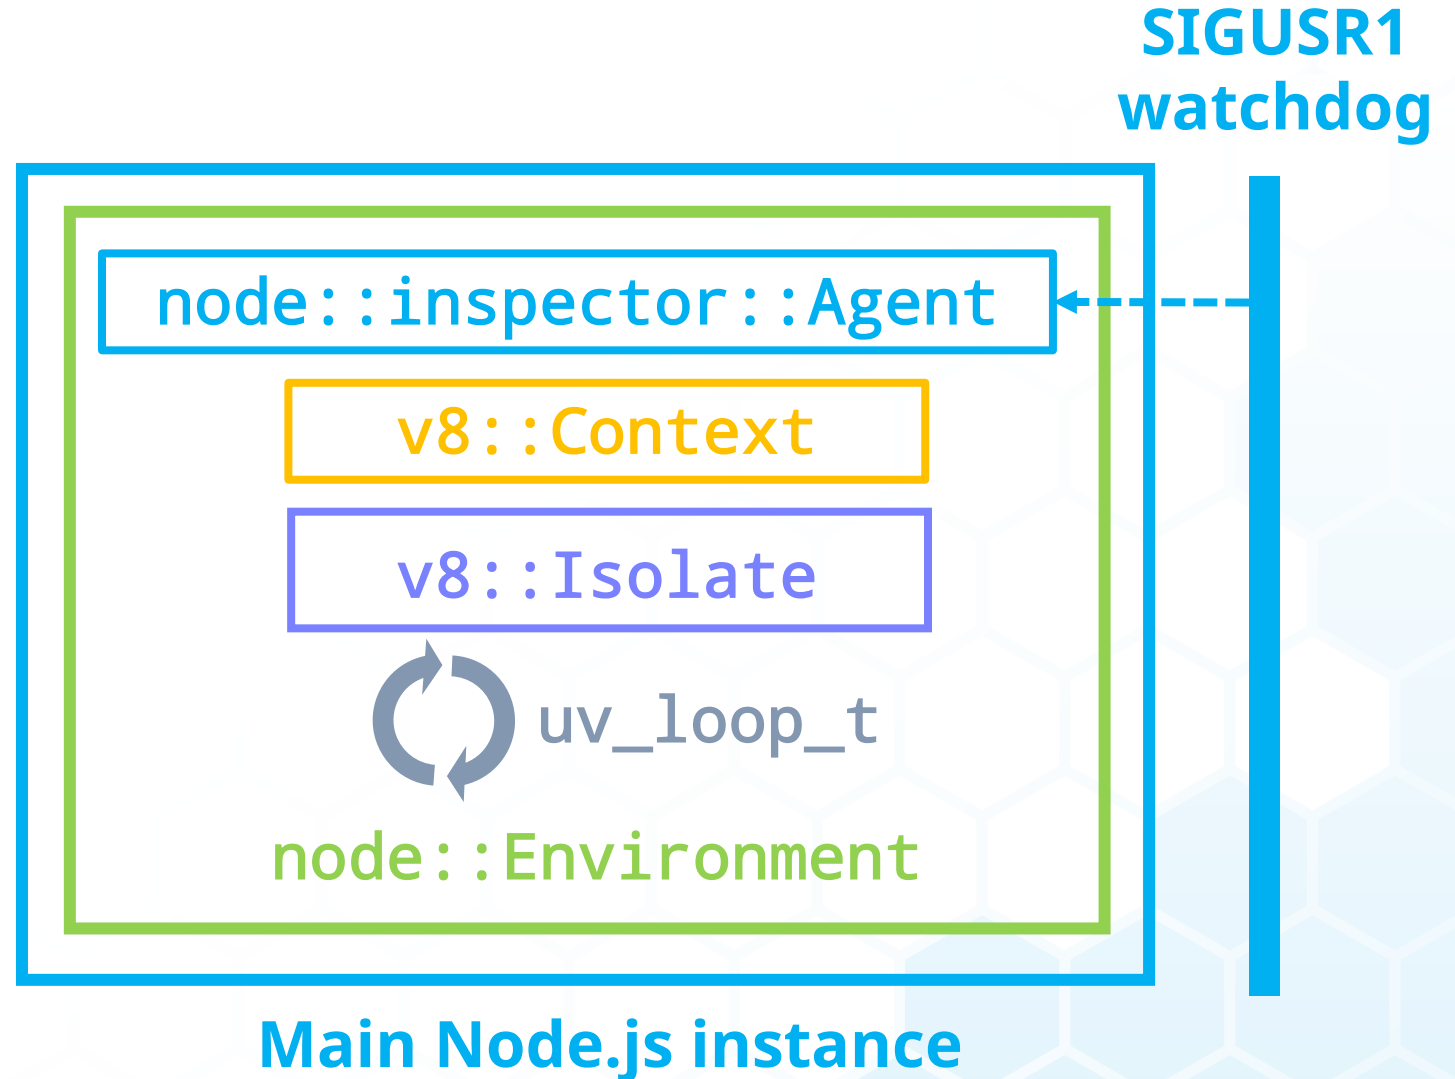


NODE+JS INTERACTIVE

December 11-12, 2019
Montreal, Canada

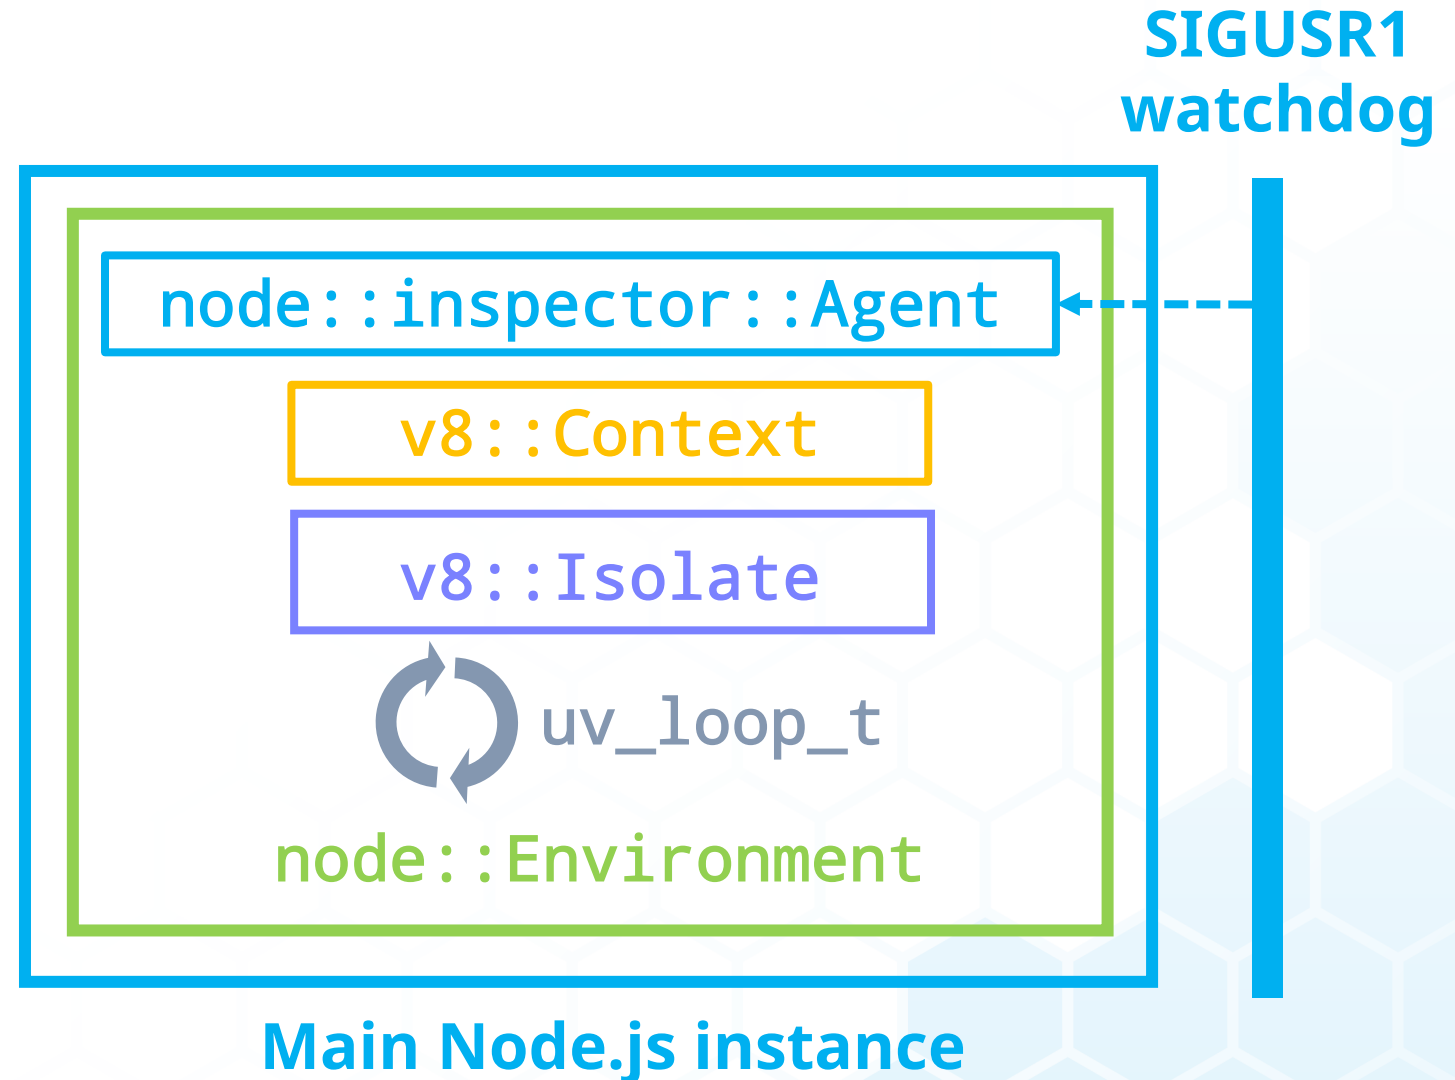
Environment: inspector

- Initialize inspector agent
- Start the SIGUSR1 watchdog thread (main instance only)

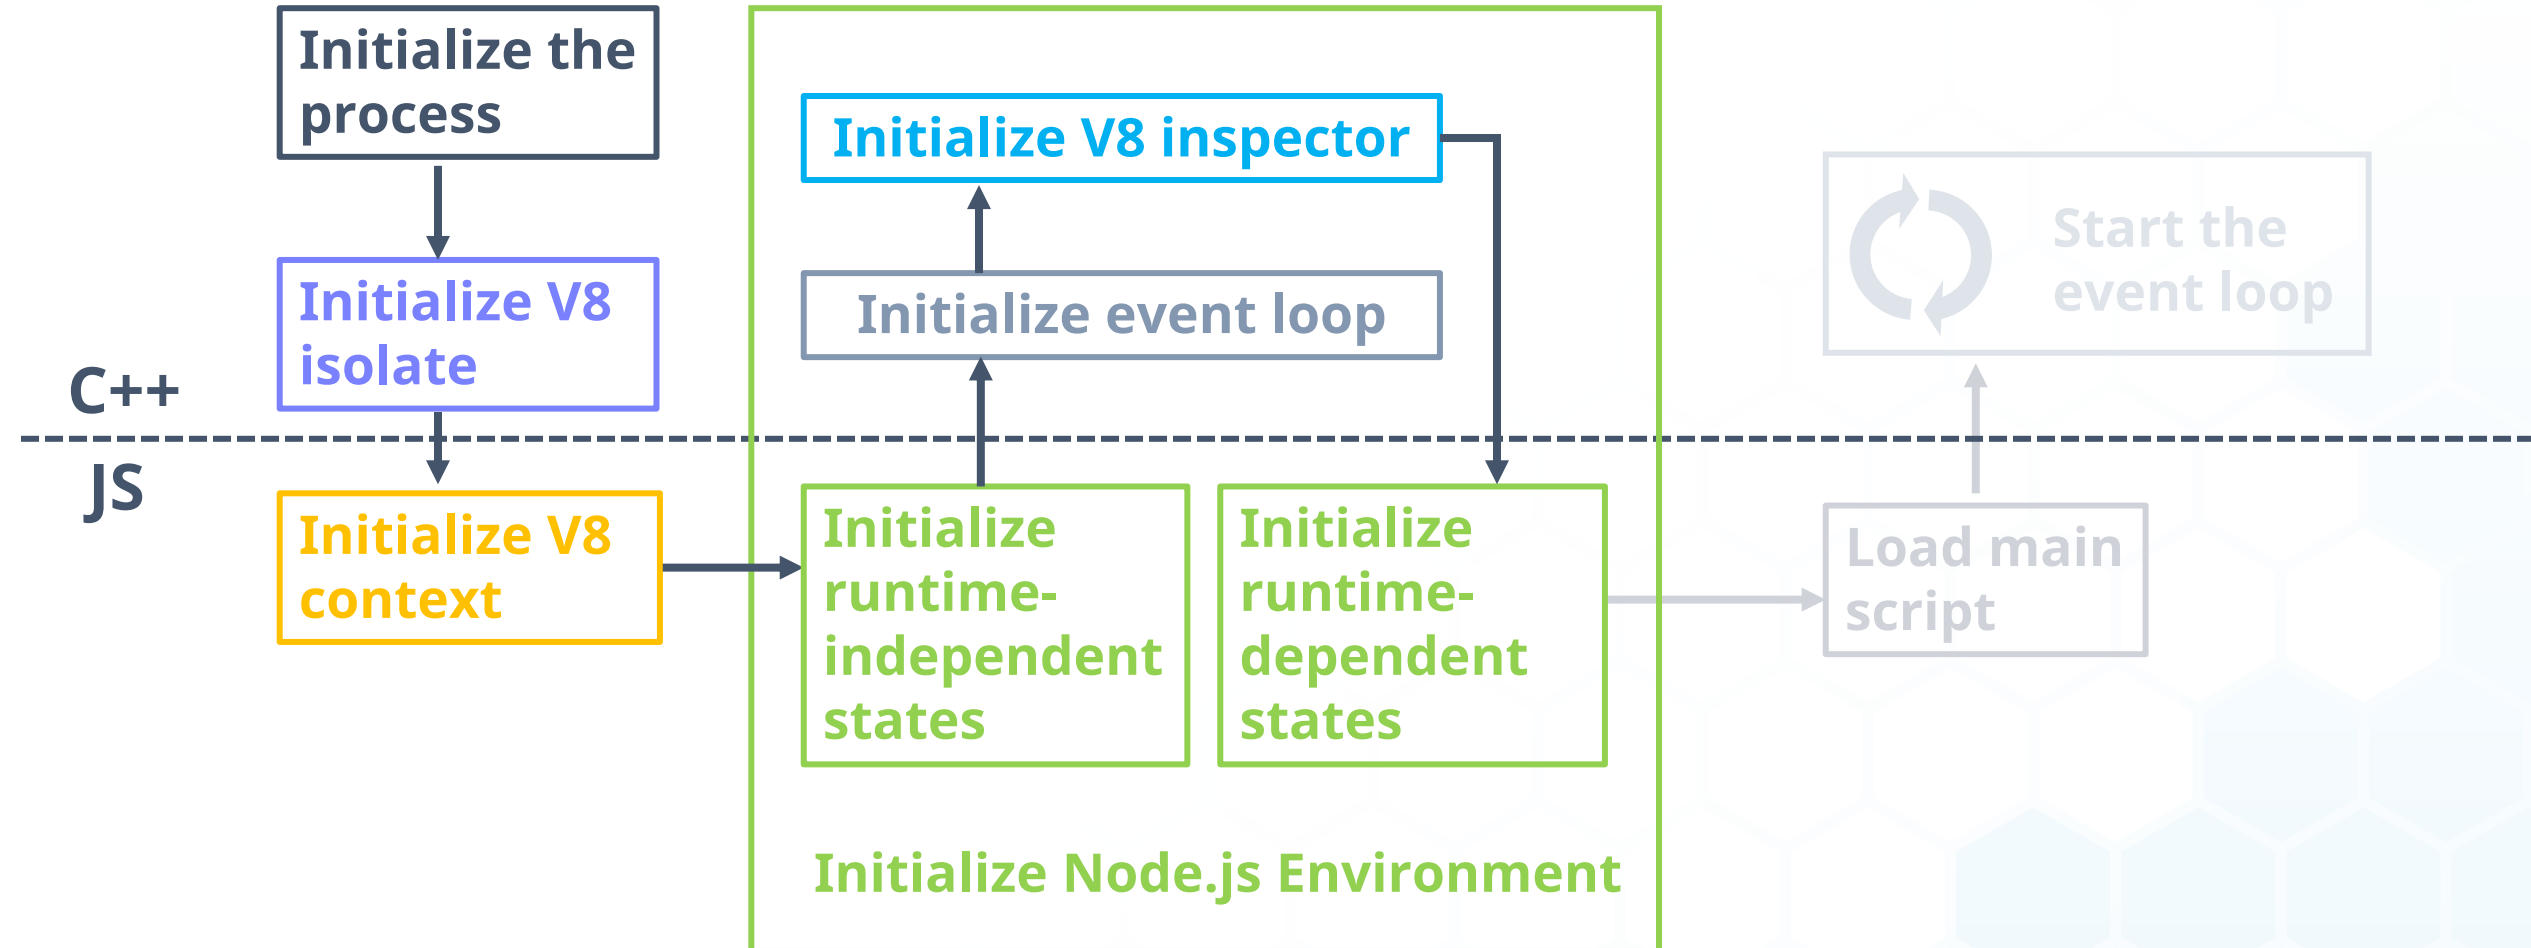


Environment: inspector

- Initialize inspector agent
- Start the SIGUSR1 watchdog thread
- Create more threads for listening on the inspector port and/or profiling, depending on CLI flags
 - `--inspect-brk`, `--cpu-prof`, `--heap-prof`



Setting up the Environment





December 11-12, 2019
Montreal, Canada

Pre-execution

- Handle various runtime configurations
 - **CLI flags:** e.g. `--no-warnings`, `--experimental-policy`, `--experimental-report`
 - **Environment variables:** e.g. `NODE_DEBUG`, `NODE_V8_COVERAGE`

```
const { onWarning } = require('internal/process/warning');  
if (!getOptionValue('--no-warnings') &&  
    process.env.NODE_NO_WARNINGS !== '1') {  
    process.on('warning', onWarning);  
}
```

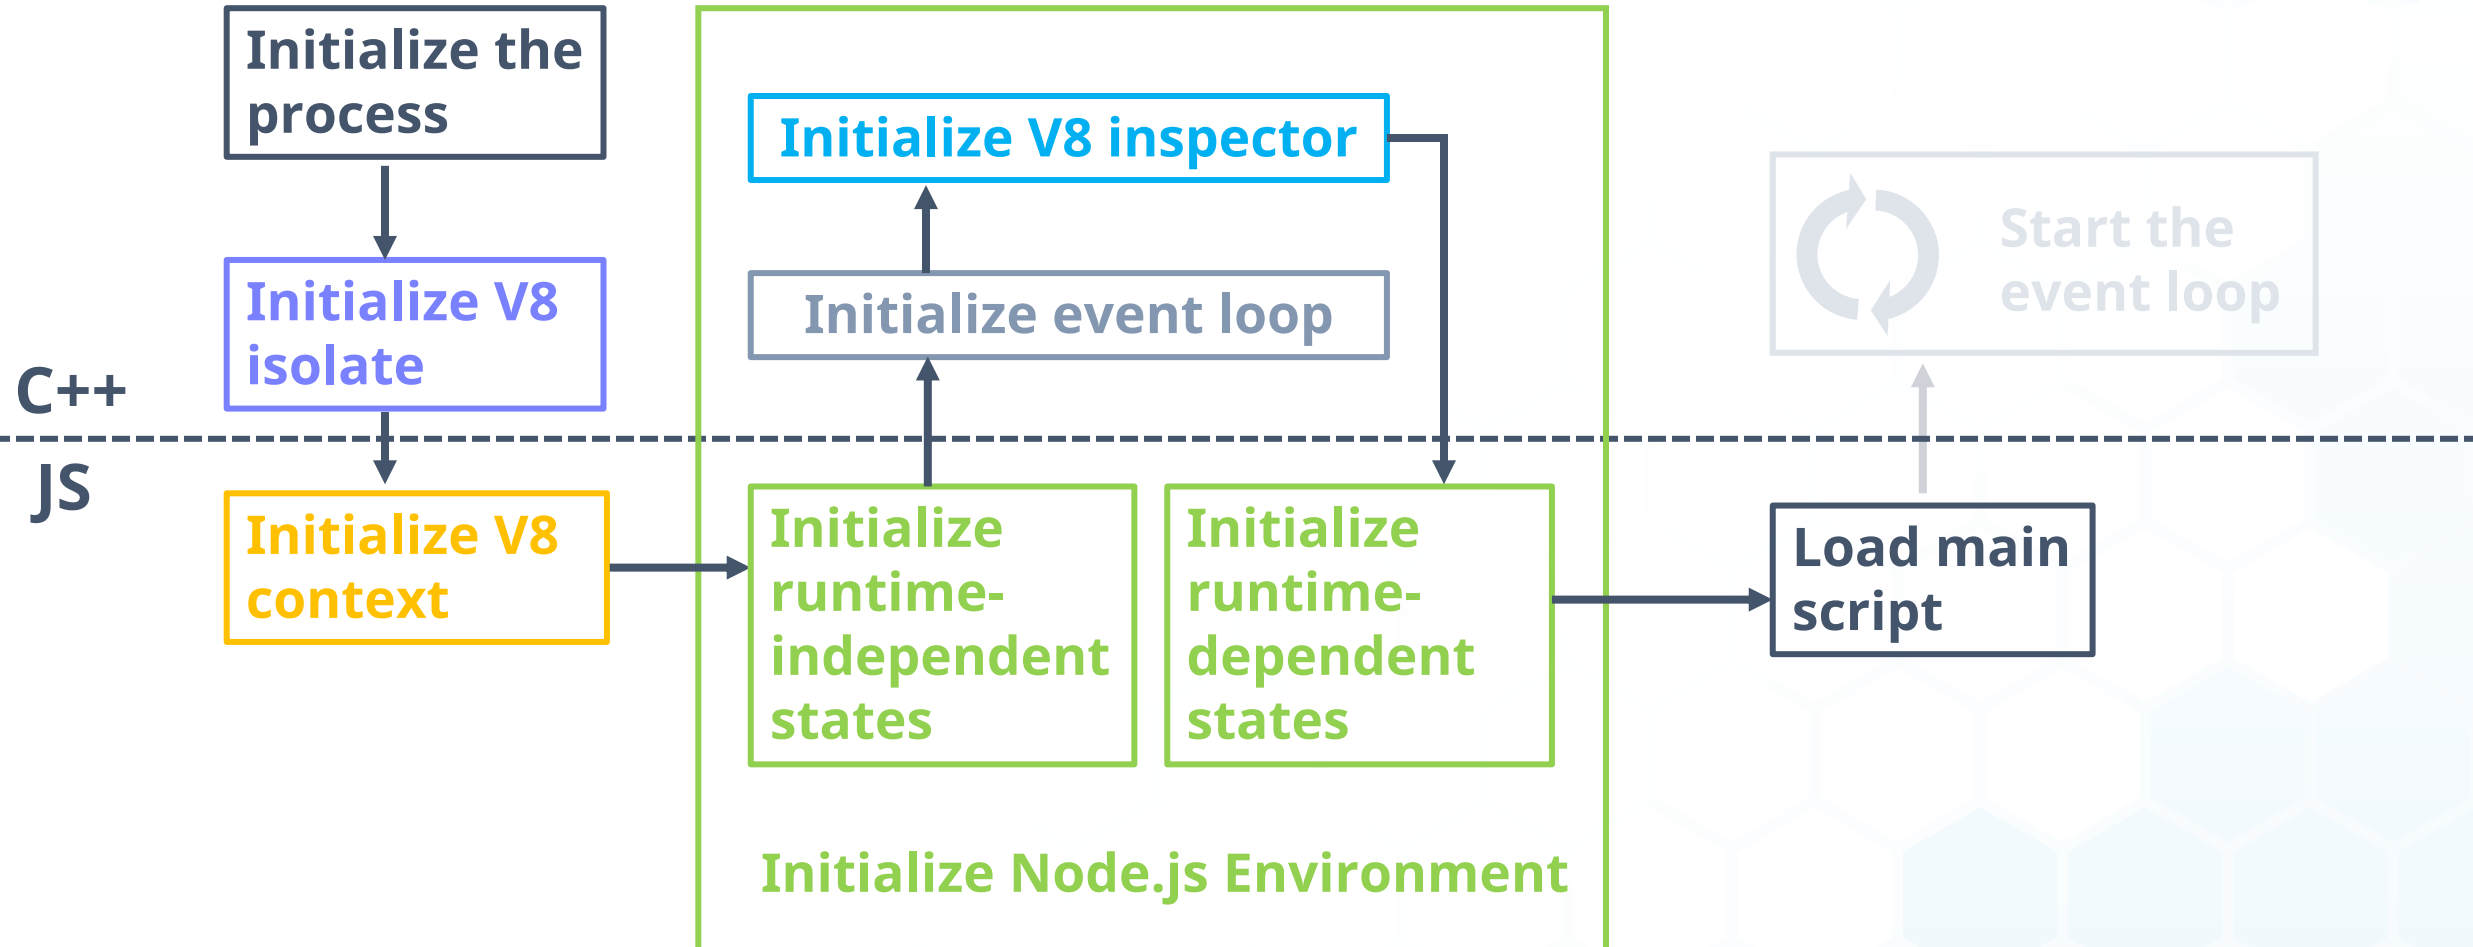



December 11-12, 2019
Montreal, Canada

Pre-execution

- Handle various runtime configurations
- Initialize IPC channel for clusters and child_process
- Initialize user-land module loaders: CJS (`require()`) and ESM (`import`)
- Load `--require` modules

Start execution





December 11-12, 2019
Montreal, Canada

Start execution

- Choose a main script according to the CLI args, etc.
 - `lib/internal/main/*.js`
 - Compiled into the binary at build time (similar to internal JS modules)



December 11-12, 2019
Montreal, Canada

Start execution: from CLI

Create and initialize
Environment

Select a main script

Load `run_main_module.js`

Detect module type

Read and compile
`${cwd}/index.js` as CJS

Start event loop

```
$ node index.js
```



December 11-12, 2019
Montreal, Canada

Start execution: Worker

Create and initialize Environment

↓ Select a main script

Load worker_thread.js

↓ Setup message port
and start listening

Start event loop

↓
Compile and run the script
sent from the port

From user code on the main thread

```
const { Worker } =  
  require('worker_threads');  
const script =  
  `console.log('hello')`;  
new worker_threads  
  .Worker(script, { eval: true });
```

From the worker_thread.js on the worker thread

```
evalScript('[worker eval]', script);
```

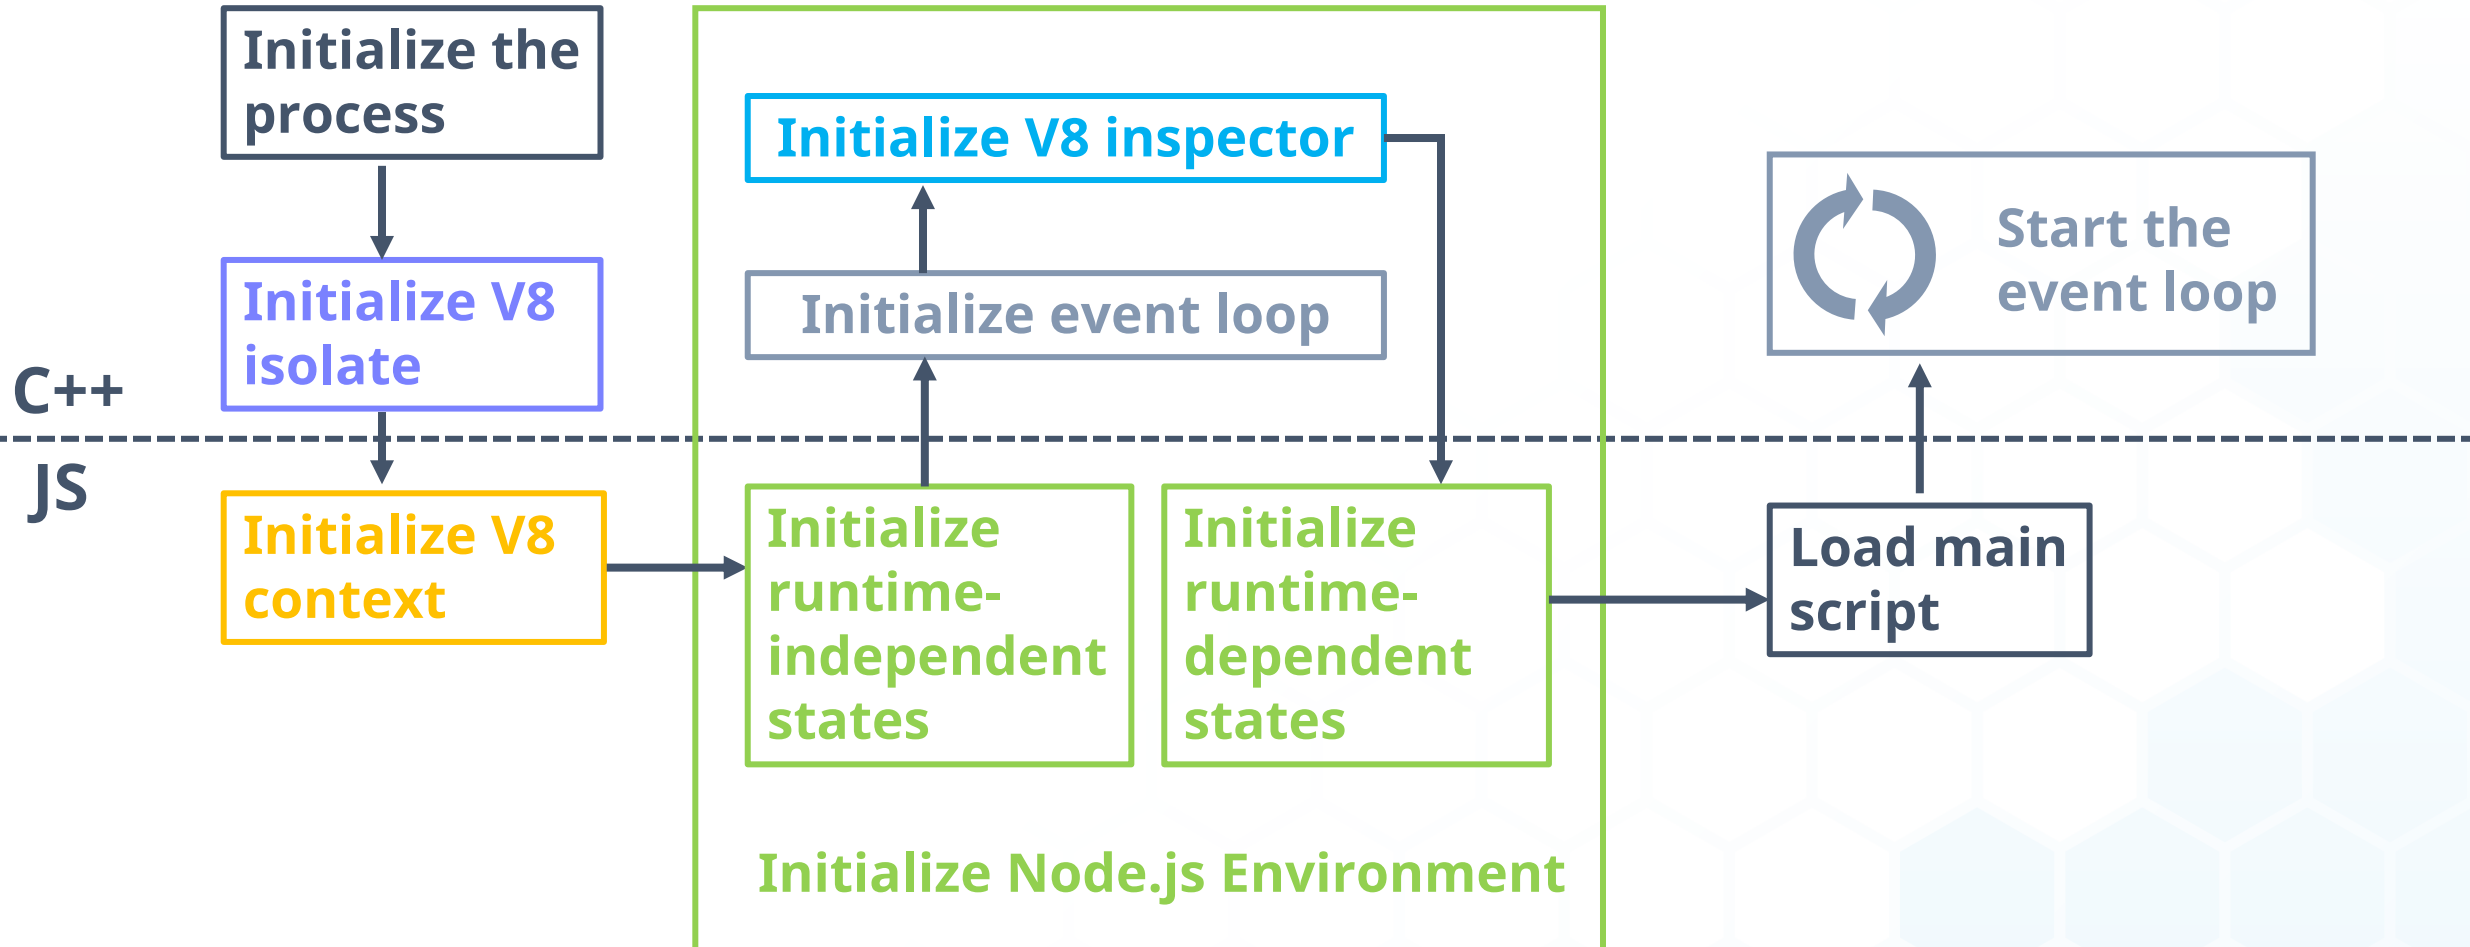


December 11-12, 2019
Montreal, Canada

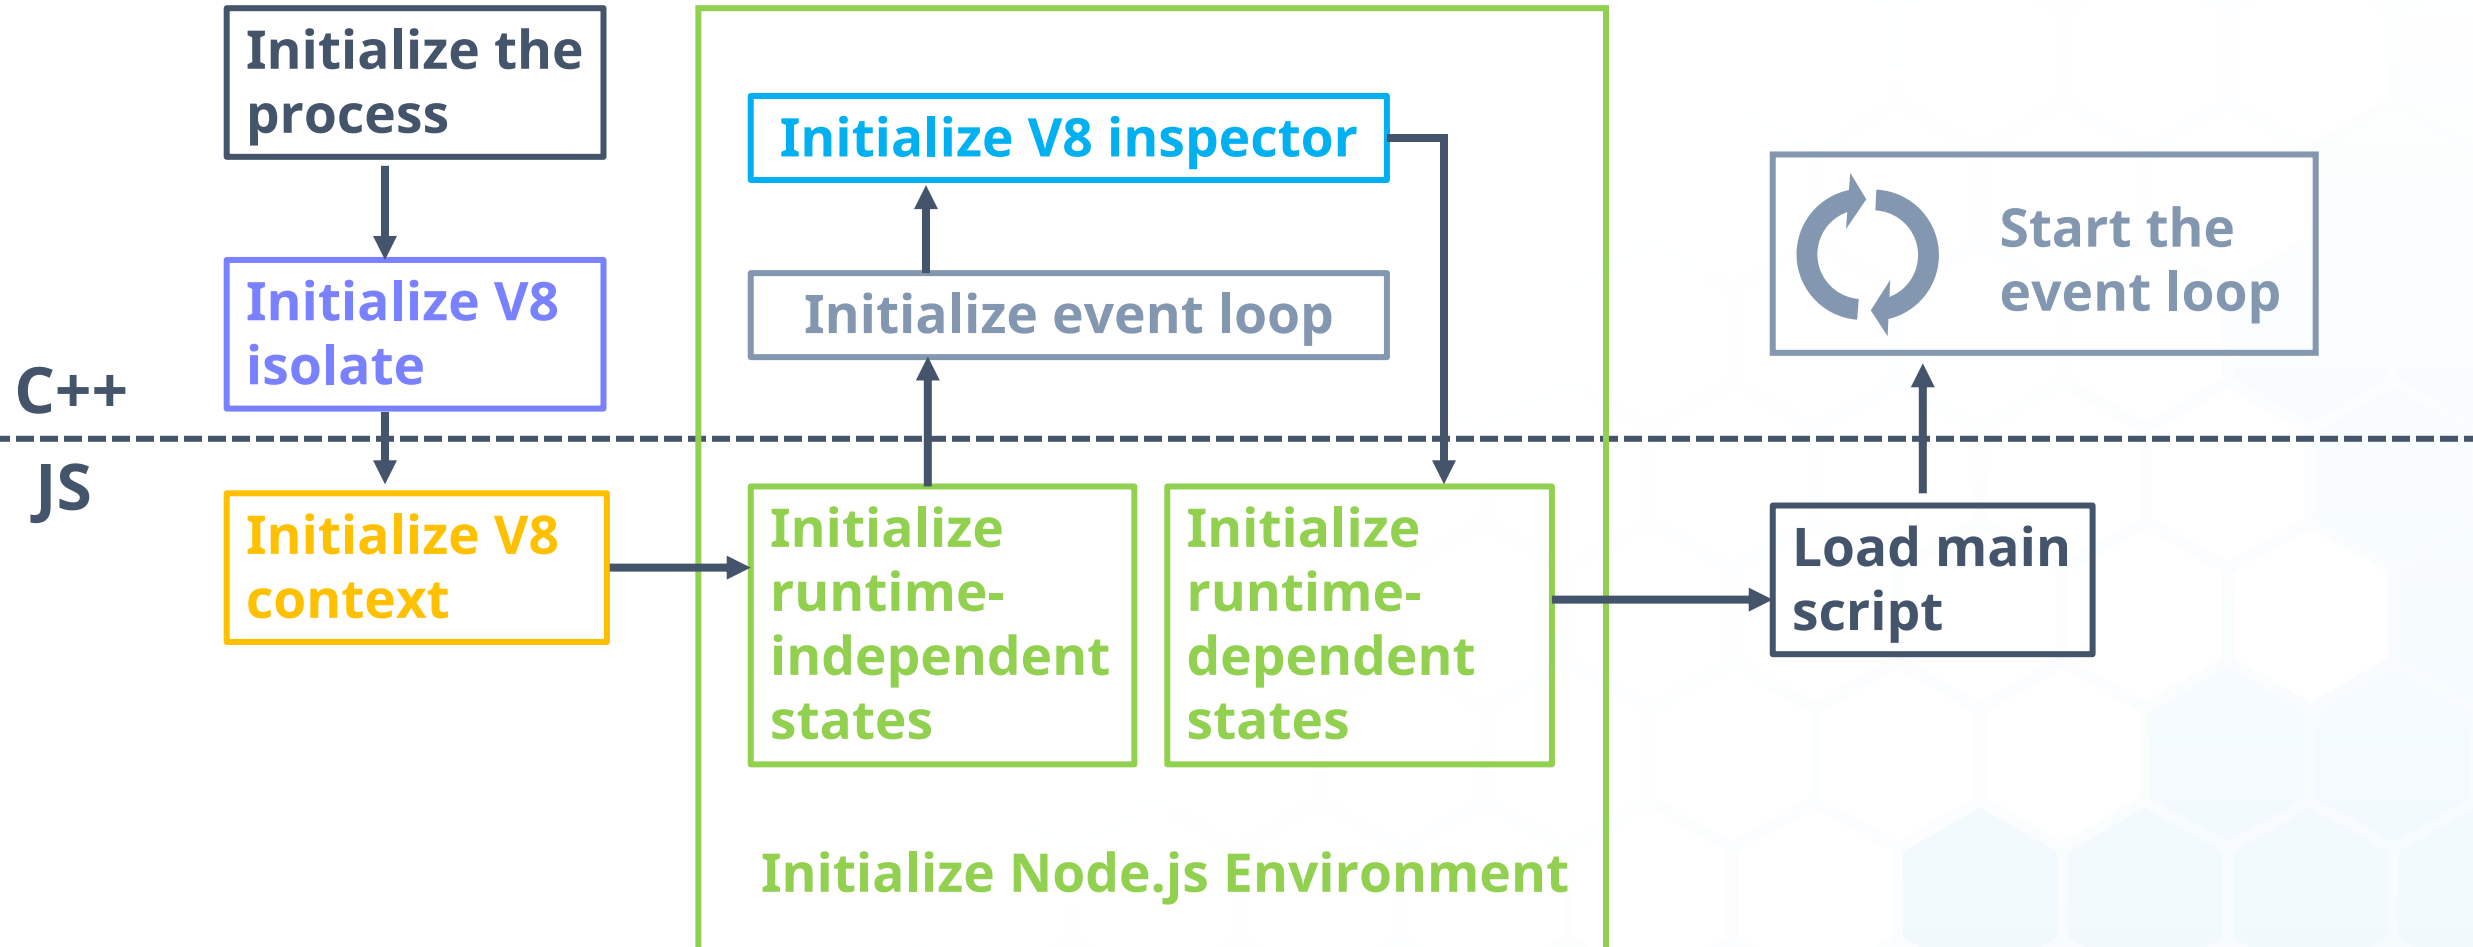
Start execution

- Kick off the event loop and run until nothing keeps it open
 - The libuv thread pool will be created if any asynchronous file system operation is used

Start execution



Summary






December 11-12, 2019
Montreal, Canada

Ongoing work

- V8 startup snapshot integration:
 - Including the runtime-independent part of Environment bootstrap into Nodejs's snapshot
- Startup performance initiative
 - <https://github.com/nodejs/TSC/blob/master/Strategic-Initiatives.md>




December 11-12, 2019
Montreal, Canada

How do hooks typically work?

If it's a per-Isolate hook (set by V8)...


1. Inside the callback, get a pointer to the `v8::Isolate` via e.g. callback arguments



How do hooks typically work?

If it's a per-Isolate hook...


1. Inside the callback, get a pointer to the `v8::Isolate` via e.g. callback arguments
2. Get the `v8::Context` that the isolate enters into



How do hooks typically work?

If it's a per-Isolate hook...

1. Inside the callback, get a pointer to the `v8::Isolate` via e.g. callback arguments
2. Get the `v8::Context` that the isolate enters into
3. Get the pointer to the `node::Environment` embedded in a slot of the context





December 11-12, 2019
Montreal, Canada


How do hooks typically work?

If it's a per-Isolate hook...

1. Inside the callback, get a pointer to the `v8::Isolate` via e.g. callback arguments
2. Get the `v8::Context` that the isolate enters into
3. Get the pointer to the `node::Environment` embedded in a slot of the context

If it's a per-Environment hook (set by node)...

Get access to the `node::Environment` directly via the callback arguments.






NODE+JS INTERACTIVE

December 11-12, 2019
Montreal, Canada

How do hooks typically work?

After gaining access to Environment

1. Get the internal **JavaScript function** stored in the `node::Environment` during bootstrap
2. Invoke the internal **JavaScript function** that calls user-provided callbacks





NODE+JS INTERACTIVE

December 11-12, 2019
Montreal, Canada

How do hooks typically work?

After gaining access to Environment

1. Get the internal JavaScript function stored in the `node::Environment` during bootstrap
2. Invoke the internal JavaScript function that calls user-provided callbacks

How are JS functions stored?

```
internalBinding('errors')  
  .setPrepareStackTraceCallback(  
    require('internal/errors')  
      .prepareStackTrace  
  );
```

