

# JavaScript and C++ in Node.js core

## How do they talk to each other?

Joyee Cheung @ NodeConf EU 2023

# Agenda

- **Overview of layers in Node.js**
- **Code invocations between C++ and JavaScript**
  - JavaScript -> C++
  - V8 Fast API calls
  - C++ -> JavaScript
- **Cross-layer memory management**
  - JavaScript -> C++ references
  - C++ -> JavaScript references
  - Trace-based GC & unified heap

# About me..

- Work on Igalia & Bloomberg collaboration
- Node.js TSC member & V8 committer
- Recent focus
  - Memory infrastructure
  - Startup snapshot
  - Code caching
- Twitter/GitHub: @joyeecheung

# Overview of layers in Node.js

~89K lines	Public JavaScript APIs		e.g. <code>fs.open()</code>
	Internal JavaScript	<b>lib/</b>	Validate & normalize arguments. Some APIs are mostly implemented in JS e.g. streams
	v8	<b>deps/ v8</b>	JavaScript virtual machine
~102K lines	C++ binding & abstractions	<b>src/</b>	Convert JS values <-> C++ values Take care of permission, tracing and hooks, invoke libraries..
	libuv/OpenSSL/zlib/...	<b>deps/</b>	Dependencies
	Operating system		

**Warning: internal property names, especially underscore-prefixed ones, are for demonstration purposes only, don't count on them in the user land**

**...and some names used in the examples are invented pseudo-names**

# JavaScript -> C++ invocation

## JavaScript land - fs.openSync()

```
// lib/fs.js
const binding = internalBinding('fs');

function openSync(path, flags, mode) {
  path = toNamespacedPath(
    getValidatedPath(path));
  flags = stringToFlags(flags);
  mode = parseFileMode(mode, 'mode',
    0o666);
  return binding.open(path, flags, mode);
}
```

- Internal version of **process.binding()** (doc-deprecated)
- Used to organize functionalities C++ expose to JS

## C++ land - binding.open()

```
static void Open(const FunctionCallbackInfo<Value>& args) {
  Environment* env = Environment::GetCurrent(args);

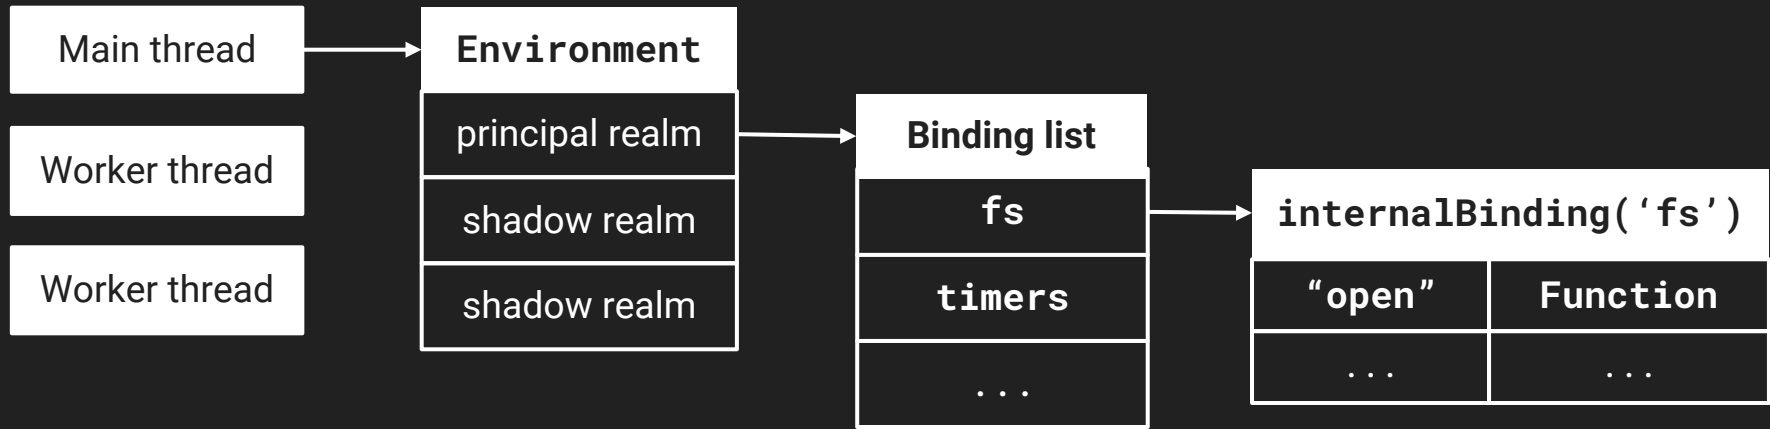
  // Convert JS arguments into native ones
  BufferValue path(env->isolate(), args[0]);
  const int flags = args[1].As<Int32>()->Value();
  const int mode = args[2].As<Int32>()->Value();

  // A bunch of abstractions but essentially...
  int result = uv_fs_open(nullptr, req, *path, flags, mode,
    nullptr /* synchronous */);

  if (!is_uv_error(result)) {
    Local<Integer> ret = v8::Integer::New(isolate, result);
    args.GetReturnValue().Set(ret); return;
  }
}
```

# How the bindings are organized - binding list

- Each Environment contains one V8 VM and one or more realms



Node.js process

# Binding list

ObjectTemplate for the fs binding	
"open"	FunctionTemplate
...	...

Instantiate



internalBinding('fs')	
"open"	Function
...	...

node::fs::Open()  
C++ function

- When **binding.open()** is called, V8 looks up the C++ function following the pointers
- Then V8 sets up some internal states (e.g. for arguments) before calling the C++ function



# JavaScript -> C++ invocation

## JavaScript land - fs.openSync()

```
// lib/fs.js
const binding = internalBinding('fs');
function openSync(path, flags, mode) {
  path = toNamespacedPath(
    getValidatedPath(path));
  flags = stringToFlags(flags);
  mode = parseFileMode(mode, 'mode',
    0o666);
  return binding.open(path, flags, mode);
}
module.exports = { openSync, ... };
```

\* Simplified version of  
**fs.openSync()** - success path  
only

## C++ land - binding.open()

```
static void Open(const FunctionCallbackInfo<Value>& args) {
  Environment* env = Environment::GetCurrent(args);

  // Convert JS arguments into native ones
  BufferValue path(env->isolate(), args[0]);
  const int flags = args[1].As<Int32>()->Value();
  const int mode = args[2].As<Int32>()->Value();

  // A bunch of abstractions but essentially...
  int result = uv_fs_open(nullptr, req, *path, flags, mode,
    nullptr /* synchronous */);

  if (!is_uv_error(result)) {
    Local<Integer> ret = v8::Integer::New(isolate, result);
    args.GetReturnValue().Set(ret); return;
  }
}
```

# JavaScript -> C++ invocation

Yes, today Node.js still uses raw V8 stuff internally instead of using anything like Node-API.

## JavaScript land - `fs.openSync()`

```
// lib/fs.js
const binding = internalBinding('fs');
function openSync(path, flags, mode) {
  path = toNamespacedPath(
    getValidatedPath(path));
  flags = stringToFlags(flags);
  mode = parseFileMode(mode, 'mode',
    0o666);
  return binding.open(path, flags, mode);
}
module.exports = { openSync, ... };
```

\* Simplified version of  
`fs.openSync()` - success path  
only

## C++ land - `binding.open()`

```
static void Open(const FunctionCallbackInfo<Value>& args) {
  Environment* env = Environment::GetCurrent(args);

  // Convert JS arguments into native ones
  BufferValue path(env->isolate(), args[0]);
  const int flags = args[1].As<Int32>()->Value();
  const int mode = args[2].As<Int32>()->Value();

  // A bunch of abstractions but essentially...
  int result = uv_fs_open(nullptr, req, *path, flags, mode,
    nullptr /* synchronous */);

  if (!is_uv_error(result)) {
    Local<Integer> ret = v8::Integer::New(isolate, result);
    args.GetReturnValue().Set(ret); return;
  }
}
```

# JavaScript -> C++ invocation

## JavaScript land - `fs.openSync()`

```
// lib/fs.js
const binding = internalBinding('fs');
function openSync(path, flags, mode) {
  path = toNamespacedPath(
    getValidatedPath(path));
  flags = stringToFlags(flags);
  mode = parseFileMode(mode, 'mode',
    0o666);
  return binding.open(path, flags, mode);
}
module.exports = { openSync, ... };
```

\* Simplified version of  
**`fs.openSync()`** - success path  
only

## C++ land - `binding.open()`

```
static void Open(const FunctionCallbackInfo<Value>& args) {
  Environment* env = Environment::GetCurrent(args);

  // Convert JS arguments into native ones
  BufferValue path(env->isolate(), args[0]);
  const int flags = args[1].As<Int32>()->Value();
  const int mode = args[2].As<Int32>()->Value();

  // A bunch of abstractions but essentially...
  int result = uv_fs_open(nullptr, req, *path, flags, mode,
    nullptr /* synchronous */);

  if (!is_uv_error(result)) {
    Local<Integer> ret = v8::Integer::New(isolate, result);
    args.GetReturnValue().Set(ret); return;
  }
}
```

# JavaScript -> C++ invocation

## JavaScript land - fs.openSync()

```
// lib/fs.js
const binding = internalBinding('fs');
function openSync(path, flags, mode) {
  path = toNamespacedPath(
    getValidatedPath(path));
  flags = stringToFlags(flags);
  mode = parseFileMode(mode, 'mode',
    0o666);
  return binding.open(path, flags, mode);
}
module.exports = { openSync, ... };
```

\* Simplified version of  
**fs.openSync()** - success path  
only

## C++ land - binding.open()

```
static void Open(const FunctionCallbackInfo<Value>& args) {
  Environment* env = Environment::GetCurrent(args);

  // Convert JS arguments into native ones
  BufferValue path(env->isolate(), args[0]);
  const int flags = args[1].As<Int32>()->Value();
  const int mode = args[2].As<Int32>()->Value();

  // A bunch of abstractions but essentially...
  int result = uv_fs_open(nullptr, req, *path, flags, mode,
    nullptr /* synchronous */);

  if (!is_uv_error(result)) {
    Local<Integer> ret = v8::Integer::New(isolate, result);
    args.GetReturnValue().Set(ret); return;
  }
}
```

# JavaScript -> C++ invocation

## JavaScript land - fs.openSync()

```
// lib/fs.js
const binding = internalBinding('fs');
function openSync(path, flags, mode) {
  path = toNamespacedPath(
    getValidatedPath(path));
  flags = stringToFlags(flags);
  mode = parseFileMode(mode, 'mode',
    0o666);
  return binding.open(path, flags, mode);
}
module.exports = { openSync, ... };
```

\* Simplified version of  
**fs.openSync()** - success path  
only

## C++ land - binding.open()

```
static void Open(const FunctionCallbackInfo<Value>& args) {
  Environment* env = Environment::GetCurrent(args);

  // Convert JS arguments into native ones
  BufferValue path(env->isolate(), args[0]);
  const int flags = args[1].As<Int32>()->Value();
  const int mode = args[2].As<Int32>()->Value();

  // A bunch of abstractions but essentially...
  int result = uv_fs_open(nullptr, req, *path, flags, mode,
    nullptr /* synchronous */);

  if (!is_uv_error(result)) {
    Local<Integer> ret = v8::Integer::New(isolate, result);
    args.GetReturnValue().Set(ret); return;
  }
}
```

**That works...but can we make it faster?**

# Yes..with V8 fast API calls

```
void GuessHandleType(const FunctionCallbackInfo<Value>& args) {  
    Isolate* isolate = args.GetIsolate();  
    Local<Context> context = isolate->GetCurrentContext();  
    int fd = args[0]->Int32Value(context).ToChecked();  
    uv_handle_type t = uv_guess_handle(fd);  
    uint32_t result = GetUVHandleTypeCode(t);  
    args.GetReturnValue().Set(Uint32::New(isolate, result));  
}
```

\* An equivalent version of the  
internal guessHandleType()  
binding

# V8 Fast API calls

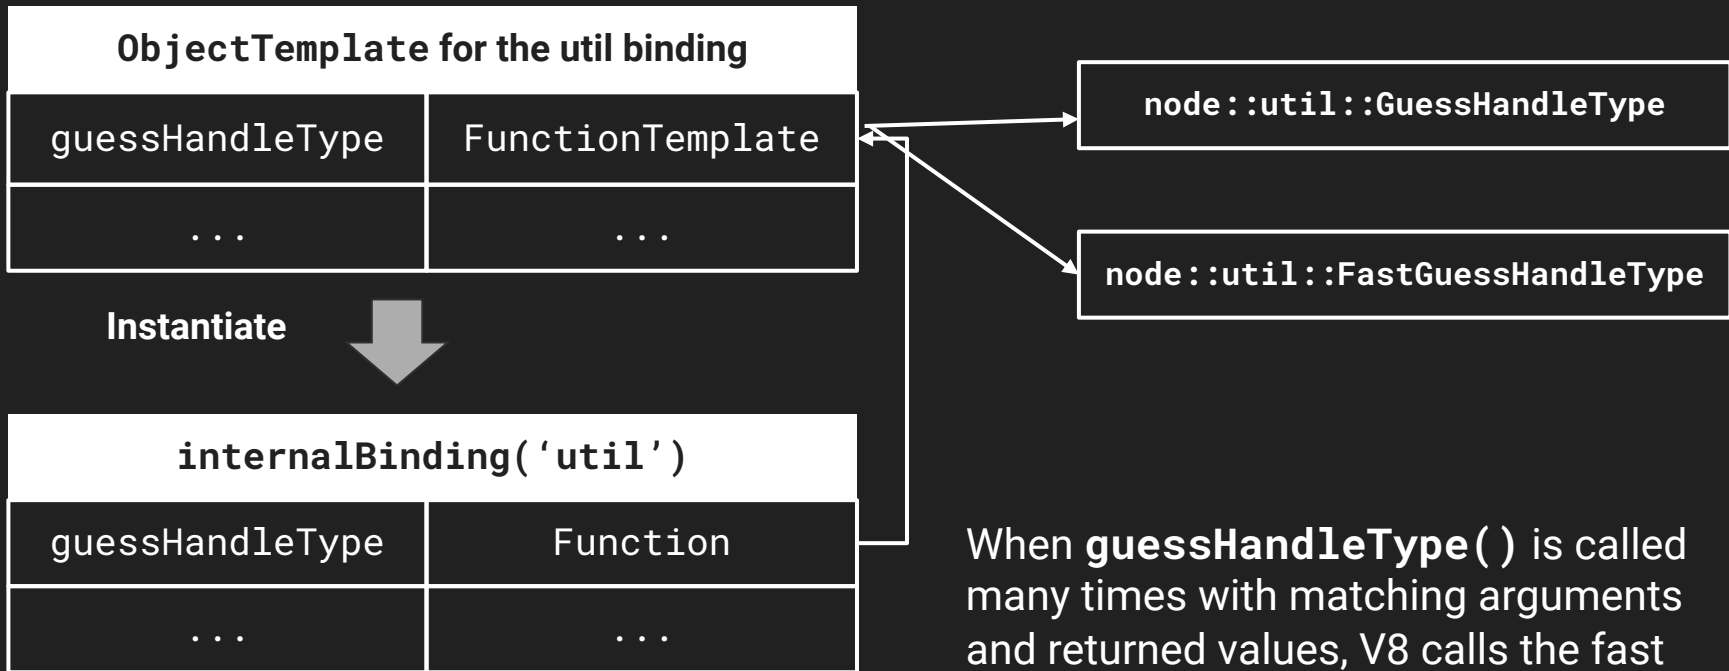
```
void GuessHandleType(const FunctionCallbackInfo<Value>& args) {  
    Isolate* isolate = args.GetIsolate();  
    Local<Context> context = isolate->GetCurrentContext();  
    int fd = args[0]->Int32Value(context).ToChecked();  
    uv_handle_type t = uv_guess_handle(fd);  
    uint32_t result = GetUVHandleTypeCode(t);  
    args.GetReturnValue().Set(Uint32::New(isolate, result));  
}  
  
uint32_t FastGuessHandleType(Local<Value> receiver, const uint32_t fd) {  
    uv_handle_type t = uv_guess_handle(fd);  
    return GetUVHandleTypeCode(t);  
}
```

\* An equivalent version of the internal guessHandleType() binding

- Add a new version whose argument/return value conversions are inlined by v8



# Binding list



When **`guessHandleType()`** is called many times with matching arguments and returned values, V8 calls the fast version for subsequent calls with matching signatures.

# V8 Fast API calls

```
uint32_t FastGuessHandleType(Local<Value> receiver, const uint32_t fd) {  
    uv_handle_type t = uv_guess_handle(fd);  
    return GetUVHandleTypeCode(t);  
}
```

- No allocations, no calling back to JS, no exceptions
- In return, smaller overhead due to not having to set things up for handling them
- ~10% faster for most bindings in Node.js - and they add up
- More bindings in Node.js are being converted to fast API calls if applicable

# C++ -> JavaScript invocation

\* Simplified version of `fs.open()`

## JavaScript land - `fs.openSync()`

```
// lib/fs.js
function open(path, flags, mode, cb) {
  // ...validate and normalize the args
  const req = new FSReqCallback();
  req.oncomplete = cb;
  binding.open(path, flags, mode, req);
}

// user land
fs.open('./f.txt', O_RDONLY, 0o666,
  (err, fd) => { /* ... */ });
```

## C++ land - `binding.open()`

```
static void Open(const FunctionCallbackInfo<Value>& args) {
  // arg handling similar to that of fs.openSync()..
  uv_fs_open(loop, req, *path, flags, mode, [])(uv_fs_t* req) {
    FSReqCallback* native_req = Unwrap(req);
    Local<Object> js_req = native_req->object();
    // req.oncomplete
    Local<Function> callback = GetOnComplete(req);
    // Convert results into callback arguments
    Local<Value> args[] = {
      is_uv_error(req->result)? v8::Null(isolate) : Err(req);
      v8::Integer::New(isolate, req->result)
    };
    callback->Call(context, js_req, 2, args);
  });
}
```

# C++ -> JavaScript invocation

\* Simplified version of `fs.open()`

## JavaScript land - `fs.openSync()`

```
// lib/fs.js
function open(path, flags, mode, cb) {
  // ...validate and normalize the args
  const req = new FSReqCallback();
  req.oncomplete = cb;
  binding.open(path, flags, mode, req);
}

// user land
fs.open('./f.txt', O_RDONLY, 0o666,
  (err, fd) => { /* ... */ });
```

## C++ land - `binding.open()`

```
static void Open(const FunctionCallbackInfo<Value>& args) {
  // arg handling similar to that of fs.openSync()..
  uv_fs_open(loop, req, *path, flags, mode, [])(uv_fs_t* req) {
    FSReqCallback* native_req = Unwrap(req);
    Local<Object> js_req = native_req->object();
    // req.oncomplete
    Local<Function> callback = GetOnComplete(req);
    // Convert results into callback arguments
    Local<Value> args[] = {
      is_uv_error(req->result)? v8::Null(isolate) : Err(req);
      v8::Integer::New(isolate, req->result)
    };
    callback->Call(context, js_req, 2, args);
  });
}
```

# C++ -> JavaScript invocation

\* Simplified version of `fs.open()`

## JavaScript land - `fs.openSync()`

```
// lib/fs.js
function open(path, flags, mode, cb) {
  // ...validate and normalize the args
  const req = new FSReqCallback();
  req.oncomplete = cb;
  binding.open(path, flags, mode, req);
}

// user land
fs.open('./f.txt', O_RDONLY, 0o666,
  (err, fd) => { /* ... */ });
```

## C++ land - `binding.open()`

```
static void Open(const FunctionCallbackInfo<Value>& args) {
  // arg handling similar to that of fs.openSync()..
  uv_fs_open(loop, req, *path, flags, mode, [](uv_fs_t* req) {
    FSReqCallback* native_req = req->data;
    Local<Object> js_req = NativeObject::New(isolate, native_req);
    // req.oncomplete
    Local<Function> callback = GetOnComplete(req);
    // Convert results into callback arguments
    Local<Value> args[] = {
      is_uv_error(req->result)? v8::Null(isolate) : Err(req);
      v8::Integer::New(isolate, req->result)
    };
    callback->Call(context, js_req, 2, args);
  });
}
```

# C++ -> JavaScript invocation

\* Simplified version of `fs.open()`

## JavaScript land - `fs.openSync()`

```
// lib/fs.js
function open(path, flags, mode, cb) {
  // ...validate and normalize the args
  const req = new FSReqCallback();
  req.oncomplete = cb;
  binding.open(path, flags, mode, req);
}

// user land
fs.open('./f.txt', O_RDONLY, 0o666,
  (err, fd) => { /* ... */ });
```

## C++ land - `binding.open()`

```
static void Open(const FunctionCallbackInfo<Value>& args) {
  // arg handling similar to that of fs.openSync()..
  uv_fs_open(loop, req, *path, flags, mode, [])(uv_fs_t* req) {
    FSReqCallback* native_req = Unwrap(req);
    Local<Object> js_req = native_req->object();
    // req.oncomplete
    Local<Function> callback = GetOnComplete(req);
    // Convert results into callback arguments
    Local<Value> args[] = {
      is_uv_error(req->result)? v8::Null(isolate) : Err(req);
      v8::Integer::New(isolate, req->result)
    };
    callback->Call(context, js_req, 2, args);
  });
}
```

# C++ -> JavaScript invocation

\* Simplified version of `fs.open()`

## JavaScript land - `fs.openSync()`

```
// lib/fs.js
function open(path, flags, mode, cb) {
  // ...validate and normalize the args
  const req = new FSReqCallback();
  req.oncomplete = cb;
  binding.open(path, flags, mode, req);
}

// user land
fs.open('./f.txt', O_RDONLY, 0o666,
  (err, fd) => { /* ... */ });
```

## C++ land - `binding.open()`

```
static void Open(const FunctionCallbackInfo<Value>& args) {
  // arg handling similar to that of fs.openSync()..
  uv_fs_open(loop, req, *path, flags, mode, [])(uv_fs_t* req) {
    FSReqCallback* native_req = Unwrap(req);
    Local<Object> js_req = native_req->object();
    // req.oncomplete
    Local<Function> callback = GetOnComplete(req);
    // Convert results into callback arguments
    Local<Value> args[] = {
      is_uv_error(req->result)? v8::Null(isolate) : Err(req),
      v8::Integer::New(isolate, req->result)
    };
    callback->Call(context, js_req, 2, args);
  });
}
```

**That seems *simple* enough...**  
**except that this is a bit too simplified**



# C++ -> JavaScript invocation

\* Simplified version of `fs.open()`

## JavaScript land - `fs.openSync()`

```
// lib/fs.js
function open(path, flags, mode, cb) {
  // ...validate and normalize the args
  const req = new FSReqCallback();
  req.oncomplete = cb;
  binding.open(path, flags, mode, req);
}

// user land
fs.open('./f.txt', O_RDONLY, 0o666,
  (err, fd) => { /* ... */ });
```

## C++ land - `binding.open()`

```
static void Open(const FunctionCallbackInfo<Value>& args) {
  // arg handling similar to that of fs.openSync()..
  uv_fs_open(loop, req, *path, flags, mode, [])(uv_fs_t* req) {
    FSReqCallback* native_req = Unwrap(req);
    Local<Object> js_req = native_req->object();
    // req.oncomplete
    Local<Function> callback = GetOnComplete(req);
    // Convert results into callback arguments
    Local<Value> args[] = {
      is_uv_error(req->result)? v8::Null(isolate) : Err(req);
      v8::Integer::New(isolate, req->result)
    };
    callback->Call(context, js_req, 2, args);
  });
}
```

# C++ -> JavaScript invocation

\* Less simplified version of  
`fs.open()`

## JavaScript land - `fs.openSync()`

```
// lib/fs.js
function open(path, flags, mode, cb) {
  // ...validate and normalize the args
  const req = new FSReqCallback();
  req.oncomplete = cb;
  binding.open(path, flags, mode, req);
}

// user land
fs.open('./f.txt', O_RDONLY, 0o666,
  (err, fd) => { /* ... */ });
```

## C++ land - `binding.open()`

```
static void Open(const FunctionCallbackInfo<Value>& args) {
  // arg handling similar to that of fs.openSync()..
  uv_fs_open(loop, req, *path, flags, mode, [](uv_fs_t* req) {
    FSReqCallback* native_req = Unwrap(req);
    Local<Object> js_req = native_req->object();
    // req.oncomplete
    Local<Function> callback = GetOnComplete(req);
    // Convert results into callback arguments
    Local<Value> args[] = {
      is_uv_error(req->result)? v8::Null(isolate) : Err(req);
      v8::Integer::New(isolate, req->result)
    };
    MakeCallback(... callback, 2, args);
  });
}
```



# What does MakeCallback does?





Besides just invoking the JS function...









- Run microtasks (from promises), if any (common)
- Run callbacks queued by `process.nextTick()`, if any (common)
- Emit AsyncHooks, if any (not as often)

# node::MakeCallback()

- Theoretically, there's no overhead if there are no microtasks, ticks, hooks
- In practice, many of them are in the way in real-world applications
- **MakeCallback()** tends to be slower than “simply calling back to JS in V8” - despite optimizations like callback coalescing
- ...changing it breaks ordering (some assumed by users), how the overhead can be reduced without hurting compatibility is still under investigation

  nodejs / performance

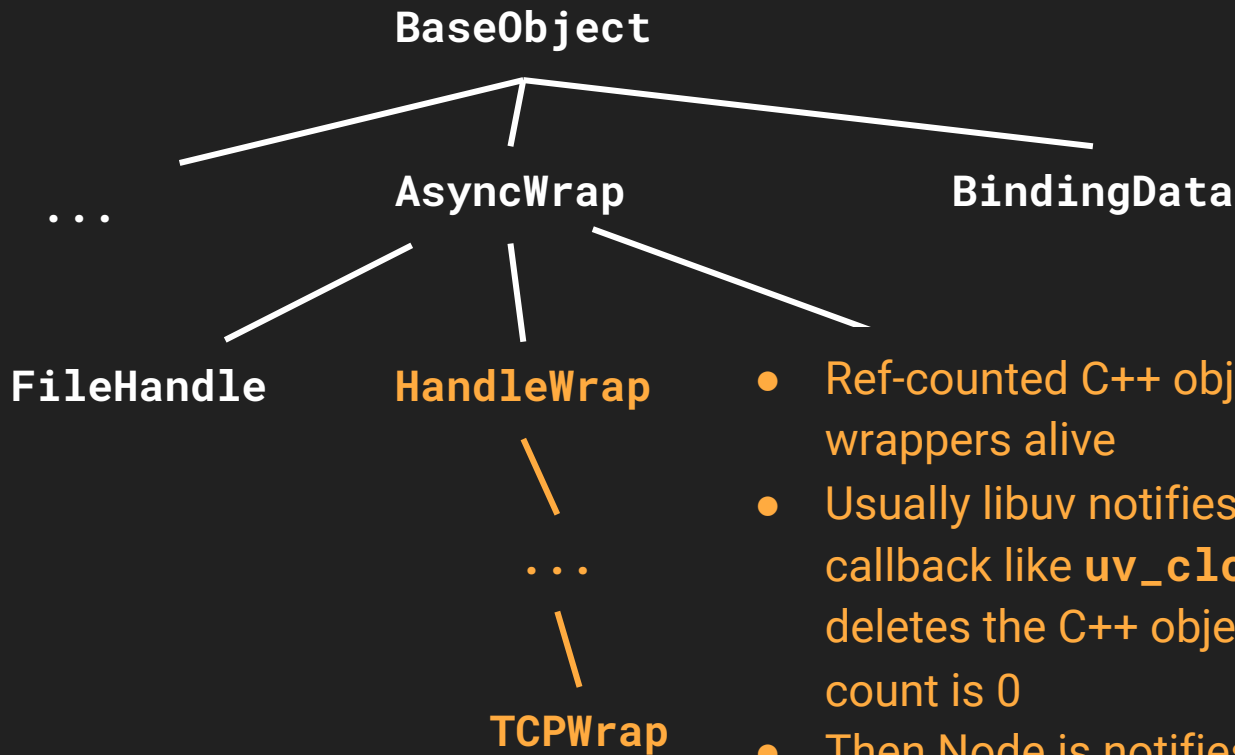
Q Type  to search  |  

 Code  Issues 45  Pull requests  Discussions  Actions  Projects  Wiki  Security

## MakeCallback is very slow #24

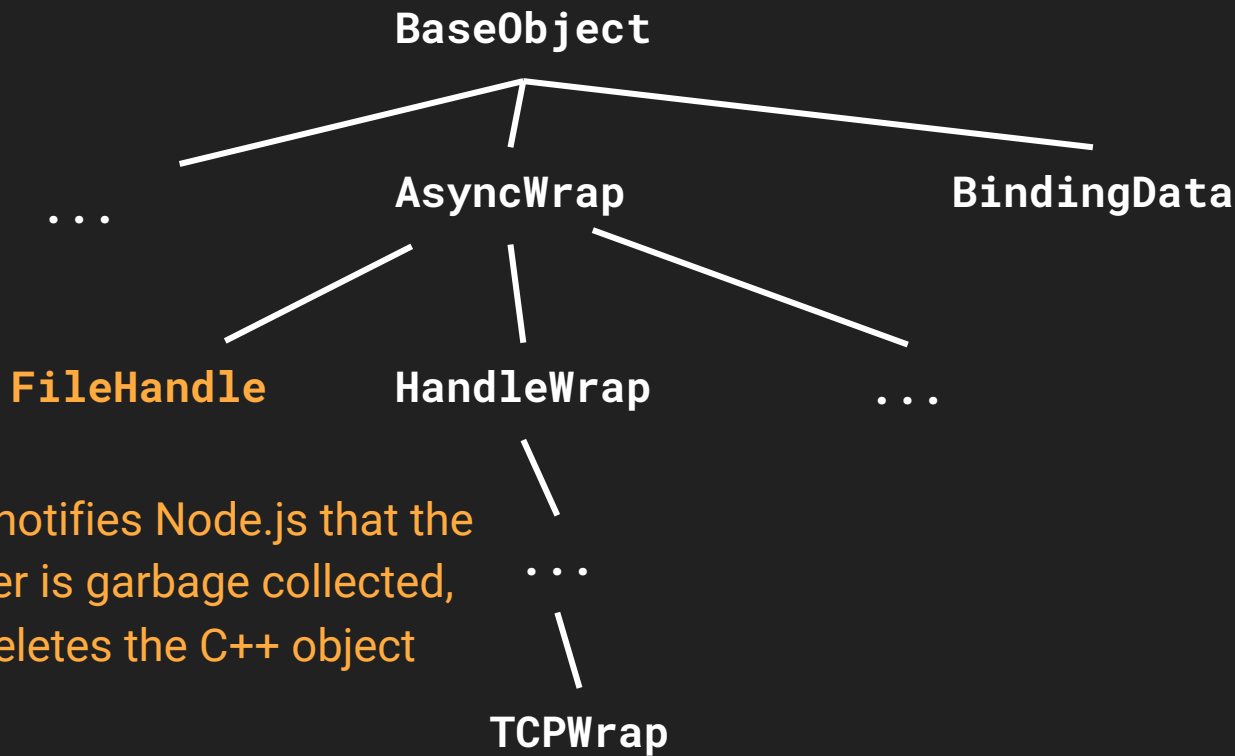
 Open ronag opened this issue on Nov 20, 2022 · 13 comments

# Classes of wrappers for resource management

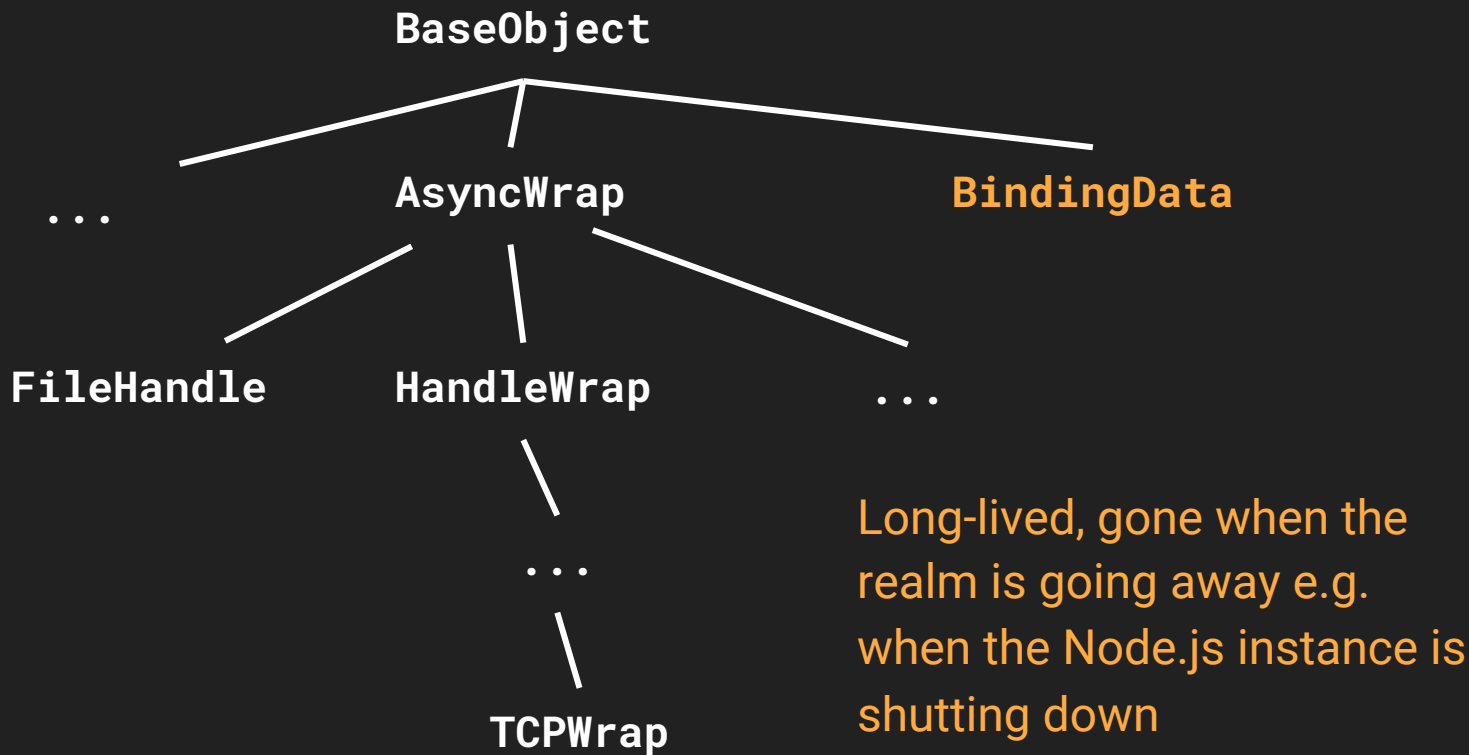


- Ref-counted C++ objects hold JS wrappers alive
- Usually libuv notifies Node.js in a callback like **uv\_close()**, Node.js deletes the C++ object when ref count is 0
- Then Node.js notifies V8 to release the JS wrapper

# Classes of wrappers for resource management



# Classes of wrappers for resource management



# JavaScript -> C++ memory reference

\* Simplified version of `net.Socket`  
- TCP path only

## JavaScript land

```
Socket.prototype.connect = function(...args) {  
  this[kHandle] = new TCP(SOCKET);  
  initSocketHandle(this);  
}
```

```
Socket.prototype.destroy = function(err, cb) {  
  closeSocketHandle(this, ...);  
}
```

```
function closeSocketHandle(self, ...) {  
  self[kHandle].close(() => {  
    self.emit('close', ...);  
    // ...  
  });  
}
```

## C++ land

```
void TCPWrap::New(const FunctionCallbackInfo<Value>& args) {  
  new TCPWrap(env, args.This(), provider);  
}  
  
void HandleWrap::Close(Local<Value> close_callback) {  
  uv_close(handle_, OnClose);  
  object()->Set(context, onclose_symbol, close_callback);  
}  
  
void HandleWrap::OnClose(uv_handle_t* handle) {  
  BaseObjectPtr<HandleWrap> wrap { handle->data };  
  wrap->Detach();  
  wrap->MakeCallback(onclose_symbol, 0, nullptr);  
}
```



# JavaScript -> C++ memory reference

\* Simplified version of `net.Socket`  
- TCP path only

## JavaScript land

```
Socket.prototype.connect = function(...args) {  
  this[kHandle] = new TCP(SOCKET);  
  initSocketHandle(this);  
}
```

```
Socket.prototype.destroy = function(err, cb) {  
  closeSocketHandle(this, ...);  
}
```

```
function closeSocketHandle(self, ...) {  
  self[kHandle].close(() => {  
    self.emit('close', ...);  
    // ...  
  });  
}
```

## C++ land

```
void TCPWrap::New(const FunctionCallbackInfo<Value>& args) {  
  new TCPWrap(env, args.This(), provider);  
}
```

```
void HandleWrap::Close(Local<Value> close_callback) {  
  uv_close(handle_, OnClose);  
  object()->Set(context, onclose_symbol, close_callback);  
}
```

```
void HandleWrap::OnClose(uv_handle_t* handle) {  
  BaseObjectPtr<HandleWrap> wrap { handle->data };  
  wrap->Detach();  
  wrap->MakeCallback(onclose_symbol, 0, nullptr);  
}
```

# JavaScript -> C++ memory reference

\* Simplified version of `net.Socket`  
- TCP path only

## JavaScript land

```
Socket.prototype.connect = function(...args) {  
  this[kHandle] = new TCP(SOCKET);  
  initSocketHandle(this);  
}
```

```
Socket.prototype.destroy = function(err, cb) {  
  closeSocketHandle(this, ...);  
}
```

```
function closeSocketHandle(self, ...) {  
  self[kHandle].close(() => {  
    self.emit('close', ...);  
    // ...  
  });  
}
```

## C++ land

```
void TCPWrap::New(const FunctionCallbackInfo<Value>& args) {  
  new TCPWrap(env, args.This(), provider);  
}  
Holds a strong reference to the JS-land TCP object
```

```
void HandleWrap::Close(Local<Value> close_callback) {  
  uv_close(handle_, OnClose);  
  object()->Set(context, onclose_symbol, close_callback);  
}
```

```
void HandleWrap::OnClose(uv_handle_t* handle) {  
  BaseObjectPtr<HandleWrap> wrap { handle->data };  
  wrap->Detach();  
  wrap->MakeCallback(onclose_symbol, 0, nullptr);  
}
```

# JavaScript -> C++ memory reference

\* Simplified version of `net.Socket`  
- TCP path only

## JavaScript land

```
Socket.prototype.connect = function(...args) {  
  this[kHandle] = new TCP(SOCKET);  
  initSocketHandle(this);  
}
```

```
Socket.prototype.destroy = function(err, cb) {  
  closeSocketHandle(this, ...);  
}
```

```
function closeSocketHandle(self, ...) {  
  self[kHandle].close(() => {  
    self.emit('close', ...);  
    // ...  
  });  
}
```


## C++ land

```
void TCPWrap::New(const FunctionCallbackInfo<Value>& args) {  
  new TCPWrap(env, args.This(), provider);  
}
```

TCPWrap inherits from HandleWrap

```
void HandleWrap::Close(Local<Value> close_callback) {  
  uv_close(handle_, OnClose);  
  object()->Set(context, onclose_symbol, close_callback);  
}
```

```
void HandleWrap::OnClose(uv_handle_t* handle) {  
  BaseObjectPtr<HandleWrap> wrap { handle->data };  
  wrap->Detach();  
  wrap->MakeCallback(onclose_symbol, 0, nullptr);  
}
```



# JavaScript -> C++ memory reference

\* Simplified version of `net.Socket`  
- TCP path only

## JavaScript land

```
Socket.prototype.connect = function(...args) {  
  this[kHandle] = new TCP(SOCKET);  
  initSocketHandle(this);  
}
```

```
Socket.prototype.destroy = function(err, cb) {  
  closeSocketHandle(this, ...);  
}
```

```
function closeSocketHandle(self, ...) {  
  self[kHandle].close(() => {  
    self.emit('close', ...);  
    // ...  
  });  
}
```

## C++ land

```
void TCPWrap::New(const FunctionCallbackInfo<Value>& args) {  
  new TCPWrap(env, args.This(), provider);  
}
```

```
void HandleWrap::Close(Local<Value> close_callback) {  
  uv_close(handle_, OnClose);  
  object()->Set(context, onclose_symbol, close_callback);  
}
```

`tcp._handle[onclose]` assigned to the callback from JS

```
void HandleWrap::OnClose(uv_handle_t* handle) {  
  BaseObjectPtr<HandleWrap> wrap { handle->data };  
  wrap->Detach();  
  wrap->MakeCallback(onclose_symbol, 0, nullptr);  
}
```

# JavaScript -> C++ memory reference

\* Simplified version of `net.Socket`  
- TCP path only

## JavaScript land

```
Socket.prototype.connect = function(...args) {  
  this[kHandle] = new TCP(SOCKET);  
  initSocketHandle(this);  
}
```

```
Socket.prototype.destroy = function(err, cb) {  
  closeSocketHandle(this, ...);  
}
```

```
function closeSocketHandle(self, ...) {  
  self[kHandle].close(() => {  
    self.emit('close', ...);  
    // ...  
  });  
}
```

## C++ land

```
void TCPWrap::New(const FunctionCallbackInfo<Value>& args) {  
  new TCPWrap(env, args.This(), provider);  
}
```

```
void HandleWrap::Close(Local<Value> close_callback) {  
  uv_close(handle_, OnClose);  
  object()->Set(context, onclose_symbol, close_callback);  
}
```

Async callback

```
void HandleWrap::OnClose(uv_handle_t* handle) {  
  BaseObjectPtr<HandleWrap> wrap { handle->data };  
  wrap->Detach();  
  wrap->MakeCallback(onclose_symbol, 0, nullptr);  
}
```

# JavaScript -> C++ memory reference

\* Simplified version of `net.Socket`  
- TCP path only

## JavaScript land

```
Socket.prototype.connect = function(...args) {  
  this[kHandle] = new TCP(SOCKET);  
  initSocketHandle(this);  
}
```

```
Socket.prototype.destroy = function(err, cb) {  
  closeSocketHandle(this, ...);  
}
```

```
function closeSocketHandle(self, ...) {  
  self[kHandle].close(() => {  
    self.emit('close', ...);  
    // ...  
  });  
}
```

## C++ land

```
void TCPWrap::New(const FunctionCallbackInfo<Value>& args) {  
  new TCPWrap(env, args.This(), provider);  
}
```

```
void HandleWrap::Close(Local<Value> close_callback) {  
  uv_close(handle_, OnClose);  
  object()->Set(context, onclose_symbol, close_callback);  
}
```

```
void HandleWrap::OnClose(uv_handle_t* handle) {  
  BaseObjectPtr<HandleWrap> wrap { handle->data };  
  wrap->Detach();  
  wrap->MakeCallback(onclose_symbol, 0, nullptr);  
}
```

Destroy TCPWrap when ref count in C++ is 0

# JavaScript -> C++ memory reference

\* Simplified version of `net.Socket`  
- TCP path only

## JavaScript land

```
Socket.prototype.connect = function(...args) {  
  this[kHandle] = new TCP(SOCKET);  
  initSocketHandle(this);  
}
```

```
Socket.prototype.destroy = function(err, cb) {  
  closeSocketHandle(this, ...);  
}
```

```
function closeSocketHandle(self, ...) {  
  self[kHandle].close(() => {  
    self.emit('close', ...);  
    // ...  
  });  
}
```

## C++ land

```
void TCPWrap::New(const FunctionCallbackInfo<Value>& args) {  
  new TCPWrap(env, args.This(), provider);  
}
```

```
void HandleWrap::Close(Local<Value> close_callback) {  
  uv_close(handle_, OnClose);  
  object()->Set(context, onclose_symbol, close_callback);  
}
```

```
void HandleWrap::OnClose(uv_handle_t* handle) {  
  BaseObjectPtr<HandleWrap> wrap { handle->data };  
  wrap->Detach();  
  wrap->MakeCallback(onclose_symbol, 0, nullptr);  
}
```

..and with it, C++ reference to the JS land TCP wrapper

# JavaScript -> C++ memory reference

\* Simplified version of `net.Socket`  
- TCP path only

## JavaScript land

```
Socket.prototype.connect = function(...args) {  
  this[kHandle] = new TCP(SOCKET);  
  initSocketHandle(this);  
}
```

```
Socket.prototype.destroy = function(err, cb) {  
  closeSocketHandle(this, ...);  
}
```

```
function closeSocketHandle(self, ...) {  
  self[kHandle].close(() => {  
    self.emit('close', ...);  
    // ...  
  });  
}
```

## C++ land

```
void TCPWrap::New(const FunctionCallbackInfo<Value>& args) {  
  new TCPWrap(env, args.This(), provider);  
}
```

```
void HandleWrap::Close(Local<Value> close_callback) {  
  uv_close(handle_, OnClose);  
  object()->Set(context, onclose_symbol, close_callback);  
}
```

```
void HandleWrap::OnClose(uv_handle_t* handle) {  
  BaseObjectPtr<HandleWrap> wrap { handle->data };  
  wrap->Detach();  
  wrap->MakeCallback(onclose_symbol, 0, nullptr)
```

**tcp.\_handle[onclose]**



# JavaScript -> C++ memory reference

\* Simplified version of `net.Socket`  
- TCP path only

## JavaScript land

```
Socket.prototype.connect = function(...args) {  
  this[kHandle] = new TCP(SOCKET);  
  initSocketHandle(this);  
}
```

```
Socket.prototype.destroy = function(err, cb) {  
  closeSocketHandle(this, ...);  
}
```

```
function closeSocketHandle(self, ...) {  
  self[kHandle].close(() => {  
    self.emit('close', ...);  
    // ...  
  });  
}
```

## C++ land

```
void TCPWrap::New(const FunctionCallbackInfo<Value>& args) {  
  new TCPWrap(env, args.This(), provider);  
}
```

```
void HandleWrap::Close(Local<Value> close_callback) {  
  uv_close(handle_, OnClose);  
  object()->Set(context, onclose_symbol, close_callback);  
}
```

```
void HandleWrap::OnClose(uv_handle_t* handle) {  
  BaseObjectPtr<HandleWrap> wrap { handle->data };  
  wrap->Detach();  
  wrap->MakeCallback(onclose_symbol, 0, nullptr);  
}
```

After close, everything will be gone when `Socket` and the `TCP handle` is unreachable from JS

# Deciphering heap snapshots

Summary ▾ socket × All objects ▾				
Constructor	▲	Distance	Shallow Size	Retain
▼ Socket		9	488 0 %	
▼ Socket @77145		9	488 0 %	
▶ <symbol kBuffer> :: system / Oddball @67		4	48 0 %	
▶ <symbol kBufferCb> :: system / Oddball @67		4	48 0 %	
▶ <symbol kBufferGen> :: system / Oddball @67		4	48 0 %	
▶ <symbol kCapture> :: system / Oddball @71		5	48 0 %	
▼ <symbol kHandle> :: TCP @6285		10	56 0 %	
▶ 3 :: onStreamRead() @75257		8	64 0 %	
▶ <symbol owner_symbol> :: Socket @77145		9	488 0 %	
▼ javascript_to_native :: Node / TCPWrap<Socket> @50		11	424 0 %	
native_to_javascript :: TCP @6285		10	56 0 %	
Retainers				
Object		Distance	Shallow Size	Retain
▶ javascript_to_native in TCP @6285		10	56 0 %	
▼ [1] in Node / CleanupQueue @48		–	320 0 %	
▼ cleanup_queue in Node / PrincipalRealm @46		–	752 0 %	
principal_realms in Node / Environment @8		–	2 328 0 %	

# Deciphering heap snapshots

Summary socket All objects

Constructor	Distance	Shallow Size	Retain
▼ Socket	9	488 0 %	
▼ Socket @77145	9	488 0 %	
▶ <symbol kBuffer> :: system / Oddball @67	4	48 0 %	
▶ <symbol kBufferCb> :: system / Oddball @67	4	48 0 %	
▶ <symbol kBufferGen> :: system / Oddball @67	4	48 0 %	
▶ <symbol kCapture> :: system / Oddball @71	5	48 0 %	
▼ <symbol kHandle> :: TCP @6285	10	56 0 %	
▶ 3 :: onStreamRead() @75257	8	64 0 %	
▶ <symbol owner_symbol> :: Socket @77145	9	488 0 %	
▼ javascript_to_native :: Node / TCPWrap<Socket> @50	11	424 0 %	
native_to_javascript :: TCP @6285	10	56 0 %	

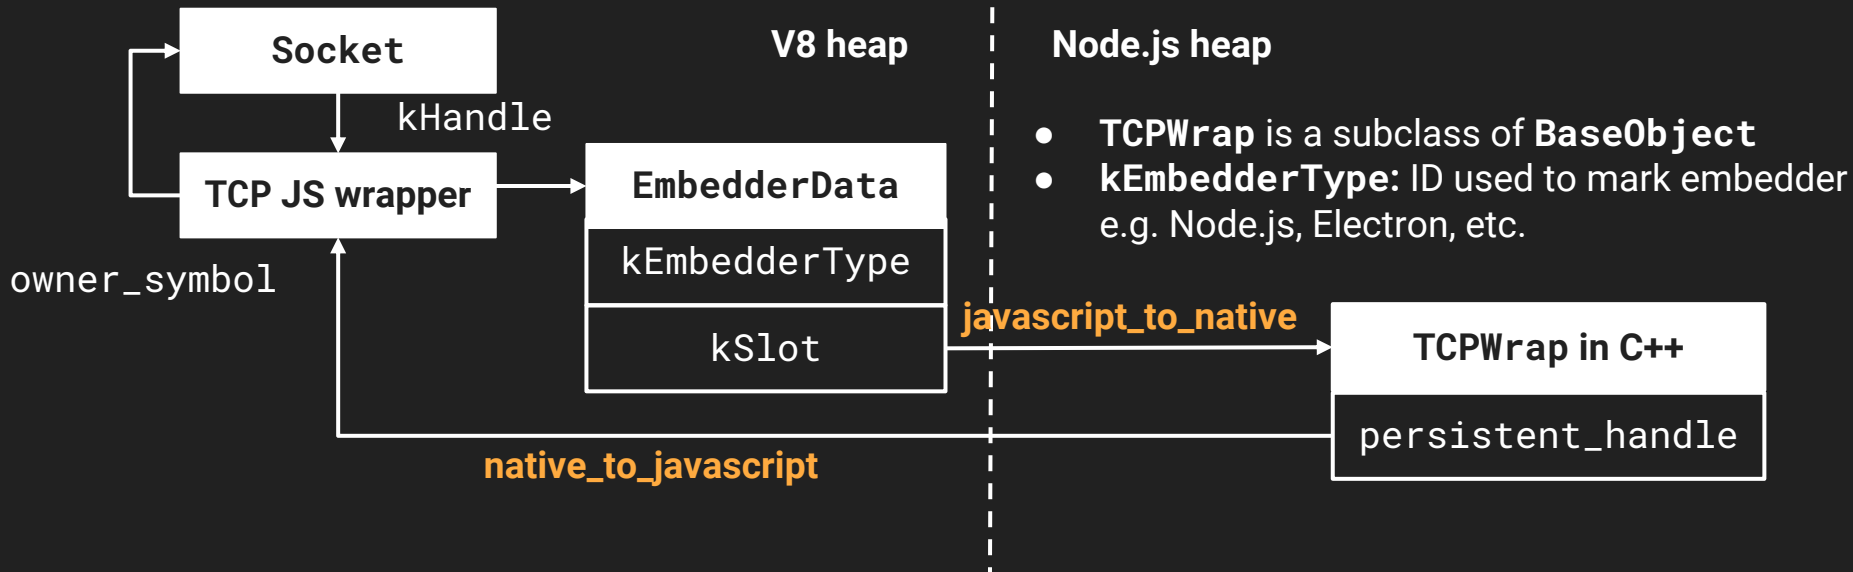
Native nodes are prefixed with "Node /"

Object	Distance	Shallow Size	Retain
▶ javascript_to_native in TCP @6285	10	56 0 %	
▼ [1] in Node / CleanupQueue @48	–	320 0 %	
▼ cleanup_queue in Node / PrincipalRealm @46	–	752 0 %	
principal_realms in Node / Environment @8	–	2 328 0 %	

# Deciphering heap snapshots

Summary	socket	All objects		
Constructor		Distance	Shallow Size	Retain
▼ Socket		9	488	0 %
▼ Socket @77145		9	488	0 %
▶ <symbol kBuffer> :: system / Oddball @67		4	48	0 %
▶ <symbol kBufferCb> :: system / Oddball @67		4	48	0 %
▶ <symbol kBufferGen> :: system / Oddball @67		4	48	0 %
▶ <symbol kCapture> :: system / Oddball @71		5	48	0 %
▼ <symbol kHandle> :: TCP @6285		10	56	0 %
▶ 3 :: onStreamRead() @75257		8	64	0 %
▶ <symbol owner_symbol> :: Socket @77145		9	488	0 %
▼ javascript_to_native :: Node / TCPWrap<Socket> @50		11	424	0 %
native_to_javascript :: TCP @6285		10	56	0 %
Retainers				
Object	Distance	Shallow Size	Retain	
▶ javascript_to_native in TCP @6285	10	56	0 %	
▼ [1] in Node / CleanupQueue @48	–	320	0 %	
▼ cleanup_queue in Node / PrincipalRealm @46	–	752	0 %	
principal_realms in Node / Environment @8	–	2 328	0 %	

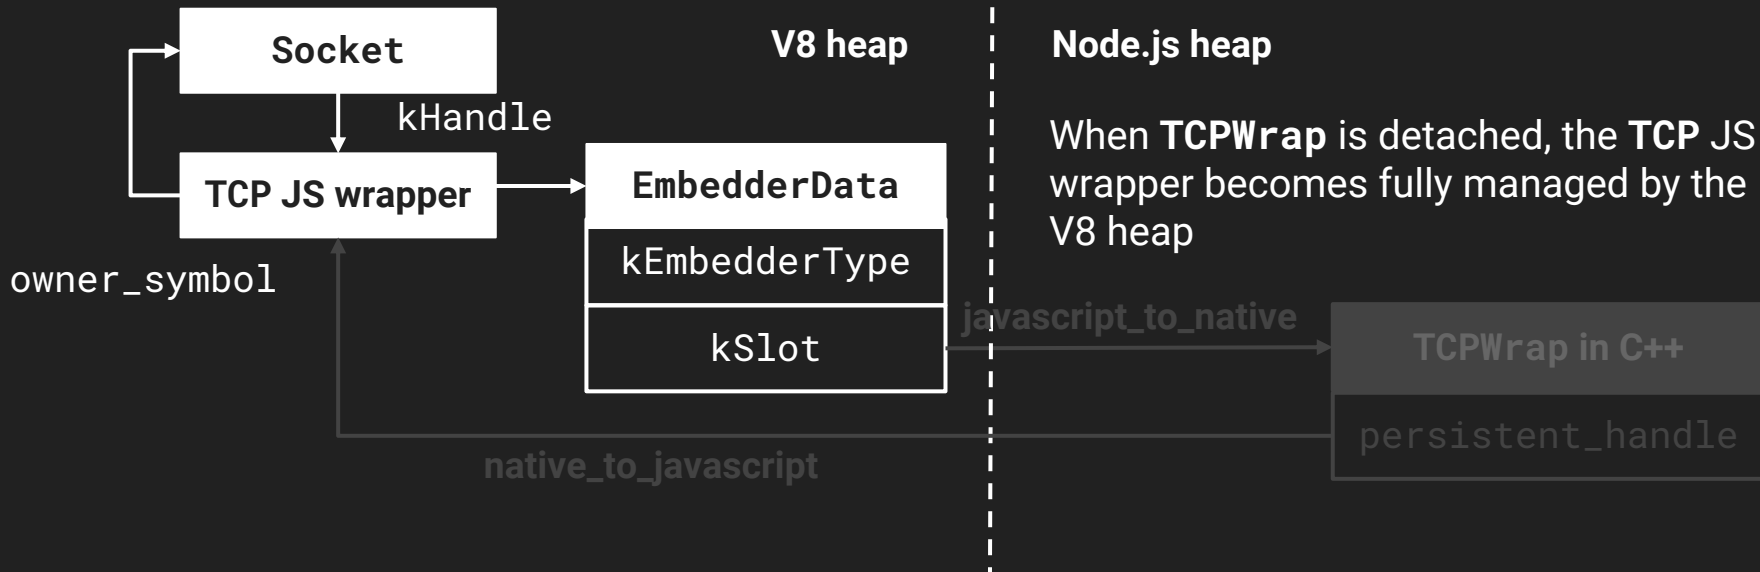
# Deciphering heap snapshots



```
▼ <symbol kHandle> :: TCP @6285
  ► 3 :: onStreamRead() @75257
  ► <symbol owner_symbol> :: Socket @77145
  ▼ javascript_to_native :: Node / TCPWrap<Socket> @50
    native_to_javascript :: TCP @6285
```

10	56	0 %
8	64	0 %
9	488	0 %
11	424	0 %
10	56	0 %

# Deciphering heap snapshots



▼ `<symbol kHandle>` :: TCP @6285

▶ 3 :: `onStreamRead()` @75257

▶ `<symbol owner_symbol>` :: Socket @77145

▼ `javascript_to_native` :: Node / TCPWrap<Socket> @50

`native_to_javascript` :: TCP @6285

10 56 0 %

8 64 0 %

9 488 0 %

11 424 0 %

10 56 0 %

# C++ -> JavaScript memory reference

- Some objects reference the JS wrapper using a weak handle so the C++ object does not hold the JS one alive, it's the other way around
- e.g. **FileHandle** (handle wrapped by **fs.promises.open()**)

```
FileHandle::FileHandle(BindingData* binding_data, Local<Object> obj, int fd) {  
    MakeWeak();  
    // ...  
}  
  
void BaseObject::MakeWeak() {  
    persistent_handle_.SetWeak(this, [] (const WeakCallbackInfo<BaseObject>& data) {  
        BaseObject* obj = data.GetParameter();  
        obj->persistent_handle_.Reset();  
        delete obj;  
    }, WeakCallbackType::kParameter);  
}
```

# C++ -> JavaScript memory reference

- Some objects reference the JS wrapper using a weak handle so the C++ object does not hold the JS one alive, it's the other way around
- e.g. **FileHandle** (handle wrapped by **fs.promises.open()**)

```
FileHandle::FileHandle(BindingData* binding_data, Local<Object> obj, int fd) {  
    MakeWeak();  
    // ...  
}
```

```
void BaseObject::MakeWeak() {  
    persistent_handle_.SetWeak(this, [](const WeakCallbackInfo<BaseObject>& data) {  
        BaseObject* obj = data.GetParameter();  
        obj->persistent_handle_.Reset();  
        delete obj;  
    }, WeakCallbackType::kParameter);  
}
```

When V8 notifies Node.js that the JS wrapper is garbage collected, Node.js delete the C++ object



# C++ -> JavaScript memory reference

The C++ objects being held alive only by JS objects and/or with no other references from C++ are shown as **Detached** in the heap snapshot

Summary	file	All objects
Constructor		
		Distance    Shallow Size
▼ FileHandle x201		4    18 472    0
▼ FileHandle @38127		11    112    0
▶ __proto__ :: EventEmitter @38351		12    24    0
▶ map :: system / Map @38353		12    72    0
▼ <symbol kHandle> :: FileHandle @29549		12    72    0
▶ __proto__ :: AsyncWrap @38339		5    56    0
▼ javascript_to_native :: Detached Node / FileHandle @254		13    160    0
native_to_javascript :: FileHandle @29549		12    72    0
▶ map :: system / Map @38341		5    72    0
▶ _events :: Object @38347		12    24    0
▶ <symbol kClosePromise> :: system / Oddball @67		4    48    0
Retainers		
Object	Distance	Shallow Size

# C++ -> JavaScript memory reference

- Many BaseObjects have additional references to JS values
- Callback-based memory management is also a hack (though it's been working)

```
/**
 * Install a finalization callback on this object.
 * NOTE: There is no guarantee as to when or even if the callback is
 * invoked. The invocation is performed solely on a best effort basis.
 * As always, GC-based finalization should not be relied upon for any
 * critical form of resource management!
 *
 * The callback is supposed to reset the handle. No further V8 API may be
 * called in this callback. In case additional work involving V8 needs to be
 * done, a second callback can be scheduled using
 * WeakCallbackInfo<void>::SetSecondPassCallback.
 */
template <typename P>
V8_INLINE void SetWeak(P* parameter,
                       typename WeakCallbackInfo<P>::Callback callback,
                       WeakCallbackType type);
```

# C++ -> JavaScript memory reference

- When the callbacks are not configured properly, bugs happen
- If you are one of the Jest users getting stuck at 16 because of some memory and/or performance issue from Node.js, please try 21 (fixes will be backported to older LTS soon)

Segmentation fault with `import()` instead of calling `importModuleDynamically` #35889

✓ Closed

🔄 Roadmap for stabilization of vm modul... #37648

nicolo-ribaudo opened this issue on Oct 30, 2020 · 37 comments


Memory leak in `vm.compileFunction` when using `importModuleDynamically` #42080

✓ Closed

sokra opened this issue on Feb 22, 2022 · 11 comments

# WIP: trace-based GC & unified heap

- Oilpan: trace-based GC library in V8 carved out from Blink
- Allows implementing an interface to create C++ -> JS references fully known to V8's garbage collector - and it's not a hack
- Bootstrapped in Node.js now, internal object migration in progress



**Thank you**